

Academic year
2022 - 2023

Software Design: Normalised Systems Theory versus Refactoring

Beau De Clercq

Master's thesis

Master of Science in computer science: software engineering

Supervisors

prof. S. Demeyer, UAntwerpen

prof. H. Mannaert, UAntwerpen



University of Antwerp
| Faculty of Science

Disclaimer Master's thesis

This document is an examination document that has not been corrected for any errors identified. Without prior written permission of both the supervisor(s) and the author(s), any copying, copying, using or realizing this publication or parts thereof is prohibited. For requests for information regarding the copying and/or use and/or realisation of parts of this publication, please contact to the university at which the author is registered.

Prior written permission from the supervisor(s) is also required for the use for industrial or commercial utility of the (original) methods, products, circuits and programs described in this thesis, and for the submission of this publication for participation in scientific prizes or competitions.

This document is in accordance with the master thesis regulations and the Code of Conduct. It has been reviewed by the supervisor and the attendant.

Contents

| | |
|--|-----------|
| 1. Introduction | 3 |
| 1.1. Stability | 3 |
| 1.2. Maintainability | 3 |
| 1.3. Extensibility | 4 |
| 2. Research question | 5 |
| 3. Defining the parameters | 7 |
| 3.1. Software Architectures | 7 |
| 3.1.1. Architectural patterns | 7 |
| 3.1.2. Design patterns | 11 |
| 3.2. Normalised Systems Theory | 14 |
| 3.2.1. Stability | 15 |
| 3.2.2. Maintainability | 17 |
| 3.2.3. Extensibility | 17 |
| 3.2.4. Elements and layers | 17 |
| 3.3. Refactoring: fixing what's broken | 18 |
| 3.4. SonarQube | 18 |
| 3.5. System under observation | 19 |
| 3.5.1. Controller layer | 19 |
| 3.5.2. Entity layer | 19 |
| 3.5.3. Exception layer | 21 |
| 3.5.4. Repository layer | 22 |
| 3.5.5. Service layer | 22 |
| 4. Refactoring into NST | 25 |
| 4.1. μ Radiant | 25 |
| 4.2. The transformation process | 26 |
| 4.2.1. Installation | 26 |
| 4.2.2. The Data Elements | 27 |
| 4.2.3. Creating the project | 27 |
| 4.2.4. Adding Data Elements | 29 |
| 4.3. Future work | 30 |
| 5. Conclusions | 31 |
| 5.1. Short term findings | 31 |
| 5.1.1. Explanations | 31 |
| 5.2. In the long run | 36 |
| 6. Previous research projects | 37 |
| 6.1. Research project 1 | 37 |
| 6.2. Research project 2 | 37 |
| A. Bibliography | 39 |

Abstract

Software Design: Normalised Systems Theory versus Refactoring

Beau De Clercq

1. Introduction

In the current software landscape, properties such as maintainability, extensibility and stability are highly valued and necessary to avoid code rot [Sof23]. In this section we will briefly go over what each of these properties mean, why it is important in a software system and how we can find and solve related problems.

TODO: extend part

1.1. Stability

Software stability refers to the reliability and robustness of a system, which ensures that it performs consistently and predictably under various conditions without crashing or producing unexpected errors. It is a critical aspect of software quality, as unstable software can lead to data loss, system downtime, and user dissatisfaction.

One key factor in achieving software stability is rigorous testing [sta21]. Robust testing processes, including unit testing, integration testing, and regression testing, help identify and fix bugs and vulnerabilities before the software is released to users. Testing also helps uncover edge cases, unusual inputs, and stress conditions that could cause instability.

Other important aspects of software stability are design and architecture. Well-designed software with a clear architecture and modular components is less likely to develop stability issues. Good design practices, such as separation of concerns, loose coupling, and high cohesion, contribute to the stability of software systems. Additionally, adhering to best coding practices, such as error handling, exception handling, and defensive programming, can help prevent software instability caused by unexpected situations.

1.2. Maintainability

The maintainability of a software system refers to the ease with which it can be maintained, i.e., how easy is it to keep the system running and up to date. Take for instance the Agile development process, which has become increasingly popular over the past few years, in which we have a near continuous cycle [Pie] of maintaining, updating and upgrading. If software becomes unmaintainable the entire workflow is at risk of getting disrupted.

As the lifecycle of a software system goes on, due to added functionalities and improvements, the overall code complexity increases and with increased complexity comes an increase in the amount of time needed to maintain the system. Another factor that leads to this increase in maintenance time is technical debt.

Technical debt [tec][sof] is a term that refers to the cost you will pay in the future by trying to re-

alise a quick fix in the present. This can happen for instance when programmers are under pressure to deliver a specific piece of software, resulting in corners being cut or general low code quality. When this happens there often is no plan as to when the debt will be re-paid and debt will silently accumulate as time goes on. The best way to counter technical debt is to keep in mind the long term effects certain changes may have.

A very nice analogy for software and maintenance over time are plants. When you first buy a plant it is compact and well defined, but within a few months this compact plant starts growing in every available direction - the addition of new functionalities and improvements of existing functionalities. And while there are no issues at first, you might notice that in some parts your plant is really getting out of hand - the accumulation of technical debt and general code deterioration. So in an effort to keep your plant maintainable you clip the problematic branches - or in the case of software, you start refactoring.

1.3. Extensibility

Extensibility [Ext21][Bab21] refers to ease with which a system or application can be extended - i.e. updated, customized or upgraded - without needing to make large alterations to the existing codebase. One common way to achieve software extensibility is by working with APIs, i.e. interfaces and protocols that allow external components or plugins to interact with the software.

One of the key benefits of software extensibility is the ability to adapt and evolve software over time. As requirements change or new technologies emerge, software can be extended to accommodate these changes without requiring a complete rewrite or overhaul. This enables software applications to remain relevant and competitive in a fast-paced and ever-changing technology landscape. Additionally, software extensibility promotes collaboration and innovation by allowing third-party developers to contribute new functionalities or integrate with existing software, fostering a rich ecosystem of plugins, extensions, or integrations that can enhance the overall value and utility of the software.

Another advantage of software extensibility is its potential to empower end-users to customize software to their specific needs. With extensible software, users can customize the functionality, appearance, or behavior of the software to align with their unique requirements or preferences. This can result in a more personalized and user-friendly experience, leading to increased user satisfaction. Moreover, software extensibility can enable users to create their own plugins or extensions, fostering a community-driven approach to software development and fostering a sense of ownership and empowerment among users.

2. Research question

The goal of this thesis is to look at the feasibility of refactoring a legacy system so they comply to the NST principles and verify that both the maintainability and the extensibility of the system have improved. To this end we will do the following

- Take a (legacy) system and refactor it using a selected method such that the system complies to the NST principles.
- Add new features to both the old and refactored system and compare the overall ease of implementing these new features as well as the code health determined by SonarQube.

The research question can thus be formulated as follows: If we take a legacy system and transform it so that it complies with the principles and guidelines defined by Normalised Systems Theory, will the promised benefits be visible immediately or will we only be able to tell a difference on the long run?

In what follows we make a few assumptions

- Refactoring a system in such a way that it complies with the principles of NST should be possible in a reasonable amount of time.
- Once refactored, adding new features to the system should be much easier. (= Extensibility)
- Once refactored, updating existing features should be as easy or easier (= Maintainability).
- Once refactored, there should not be a negative impact on the performance of the system.

3. Defining the parameters

In this chapter we will define the parameters of this thesis. These parameters as we like to call them are as follows

- **Normalised Systems Theory:** Normalised Systems Theory provides principles and guidelines which, if followed correctly, result in software which is stable, maintainable and extensible.
- **Software Architectures:** As Normalised Systems Theory defines an architecture, we will first give a brief overview of the different kinds of architectures and patterns.
- **Refactoring:** The line between recreating an application, also known as reengineering, and refactoring is fairly thin. To better distinguish between them we will see how they are related to each other.
- **SonarQube:** As this will be the basis on which we will compare our systems, we introduce what SonarQube is and how we intent to use it.
- **System under observation:** The system with which we would like to verify our research question.

Combined these parameters define this thesis.

3.1. Software Architectures

Just as with the physical architecture, a software architecture is the entirety of all components and their placements needed to build a product. To do so in the most efficient way possible, over the years numerous architectural patterns have been developed to provide clear guidelines on how to solve a certain set of problems. In addition to these architectural patterns, a multitude of design patterns have been developed to provide a clear structure to each of the components of the architectural pattern.

3.1.1. Architectural patterns

Architectural patterns define the foundation on which the software will be build: they provide general solutions to recurring problems. Among the more well known are client-server, multi-layer, microservices and REST.

It is perfectly possible to use architectural patterns in combination with one another. For example, in the microservice pattern a developer could decide to implement one service according to the multi-layer pattern, e.g. an application to rent a car, another one with REST, e.g. the front-end to the rental application, and yet another one as client-server, e.g. the location where all invoices are stored.

3.1.1.1. Client-server

In the client-server pattern, as the name suggests, we have clients which communicate with servers such as a web server or file server. Whether a computer is a client, a server, or both, is determined by the nature of the application that requires the service functions. For example, a single computer can run a web server and file server software at the same time to serve different data to clients making different kinds of requests. The client software can also communicate with server software within the same computer. In order to support frequent/high traffic, the application is often run on multiple servers.

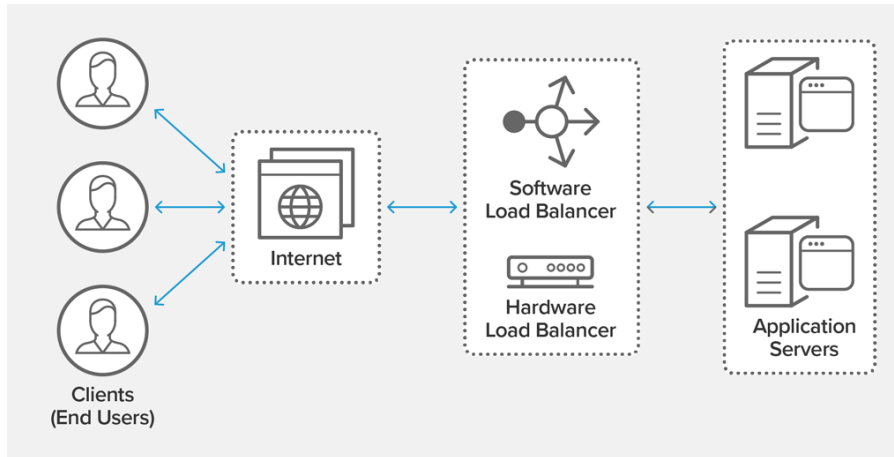


Figure 3.1.: Load balancing diagram

Clients then connect to one of the servers through a load balancer. This load balancer distributes all incoming client traffic among the available servers, thus ensuring that no server receives more than it can handle. An example of how a load balancer can be placed is shown in Figure 3.1 [NGI].

3.1.1.2. Multi-layer

A multi-layer or multitier architecture is a client-server architecture in which functionalities are physically separated over multiple layers. The most common layer division [Ric15] is Presentation-Business-Persistence-Database. An example is shown in Figure 3.2 [Ric15]. One of the powerful features of the layered architecture pattern is separation of concerns among components. Components within a specific layer deal only with logic that belongs to that layer. For example, components in the presentation layer deal only with presentation logic, whereas components in the business layer deal only with business logic.

In this architecture each layer is also closed, meaning that as a request moves from layer to layer, it must go through the layer right below it to get to the next layer below that one. The reason we don't allow each layer to have direct access to all others is due to the concept layers of isolation. This concept means that if changes are made in one of the layers, these changes don't impact or affect the other layers. This concept also means that each layer is independent of the rest, thereby having no knowledge of how the other layers work internally.

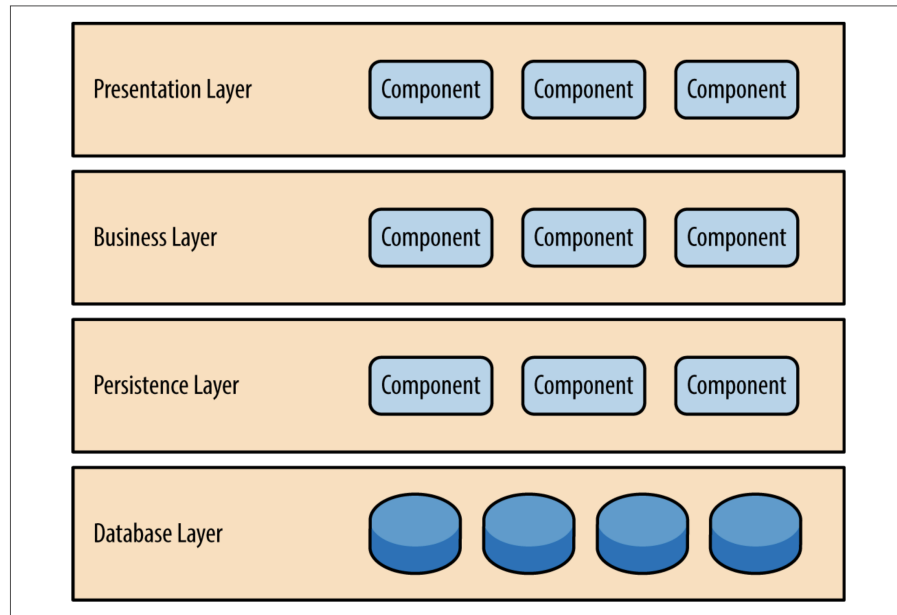


Figure 3.2.: Example of multi-tier architecture

3.1.1.3. Microservices

The microservice architectural style is an approach to developing a single application as a suite of small services where each service runs in its own process and can communicate via lightweight mechanisms.

Microservices are built around business capabilities and are independently deployable by fully automated deployment machinery. They can be written in different programming languages and can use different data storage technologies. When using microservices there is only the bare minimum of centralized management. As an example, consider a webshop as depicted in Figure 3.3. In this example

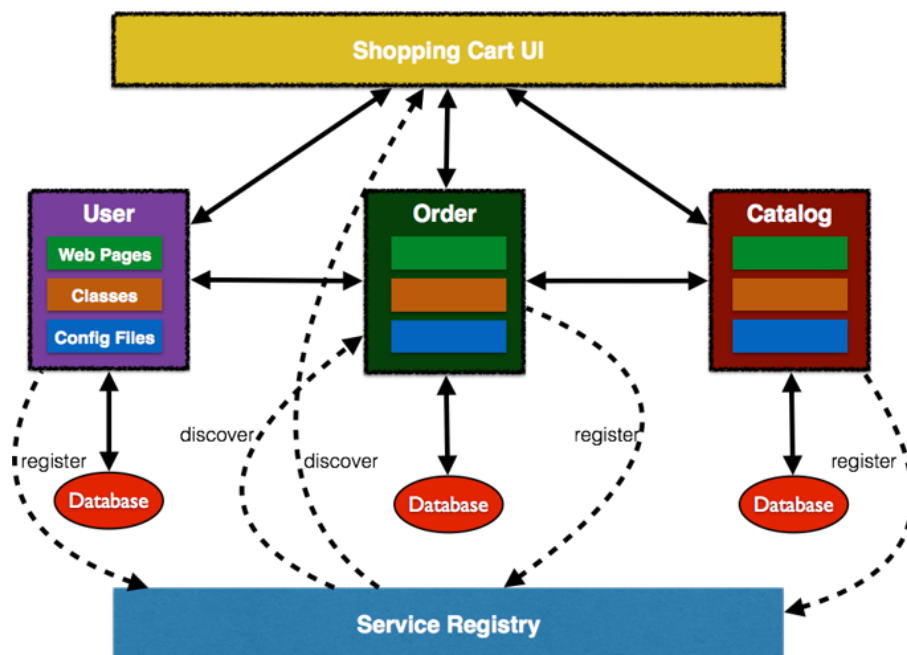


Figure 3.3.: Example of microservice architecture

our application consists of 3 components: Catalog, Order, User. Each component has its own databases so that each microservice can evolve and choose whatever type of datastore – relational, NoSQL, flat file, in-memory or some thing else – is most appropriate. Each component will register with a Service Registry. This is required because multiple stateless instances of each service might be running at a given time and their exact endpoint location will be known only at the runtime. Client interaction for the application is defined in another application, Shopping Cart UI in our case. This application mostly discover the services from Service Registry and composes them together. It should mostly be a dumb proxy where the UI pages of different components are invoked to show the interface.

At time of writing, the most well-known and most used framework for microservices is Docker. Docker works by running software in containers which can then communicate with each other via API calls. Each container can then be developed separately.

3.1.1.4. REST

Representational state transfer [ECPB12] (REST) is a style of software **architecture** for distributed systems such as the World Wide Web.

The term resource is a central concept that denotes any item of interest identified by some identifier. An identifier, also called Uniform Resource Identifier (URI), can be either a URL¹ (the location where you can find the object or the method for finding it) or a URN² (which defines an item identity).

A RESTful API is a Web Service that adheres to the REST style, meaning

- It has an Internet media type for data (XML, JSON, ...)
- It has a base URI
- It uses standard HTTP methods (GET, PUT, POST, DELETE)
- It uses hypertext links to reference state
- It uses hypertext links to reference related resources

The REST style states six constraints which, if adhered to, provide an architecture which has properties such as performance, scalability, simplicity, modifiability, visibility, portability, and reliability. These constraints are as follows

- Client–server architecture: Employing the client-server architecture enforces the principle of separation of concerns by separating UI from data storage, allowing components to evolve independently.
- Statelessness: When communicating with a server, each packet send by the client can be understood separately from the others as it will contain all relevant session data, i.e. no session data is kept by the server.
- Cacheability: Responses received by the client must identify themselves as either cacheable or non-cacheable, this to prevent clients from providing stale or inappropriate data in response to further requests. If used correctly, caching can in part or fully eliminate some client-server interactions.

¹Uniform Resource Locator

²Uniform Resource Name

- Layered system: Clients should not be able to tell whether it is connected to an end-server directly or to a proxy, i.e. the use of load balancers and the addition of security layers won't impact client communications.
- Uniform interface: A uniform interface enables the architecture to be decoupled. It's constraints are the following
 - Resource identification in requests: resources are identified in requests via URIs in RESTful webservices, e.g. the server could send data as XML or JSON.
 - Resource manipulation through representations: when a clients has a representation of a resource, it has enough information to work with the resource's state.
 - Self-descriptive messages: each message includes enough information to describe how to process it.
 - Hypermedia as the engine of application state: starting from a base URI, a client can dynamically access all resources it needs. I.e., starting from a homepage a client can access resources further down the hierarchy.
- Optional - Code on demand: Servers can temporarily alter client functionality by transferring executable code such as JavaScript.

3.1.2. Design patterns

As the name suggests, design patterns provide a pattern for how the actual implementation should look. The difference between an architectural pattern and a design pattern is the level on which it is applied: architectural patterns always end up on top of design patterns.

Design patterns can roughly be divided into three main categories [GHJ⁺94][DSN03], namely behavioral patterns such as the observer, structural patterns such as adapter and proxy and creational patterns such as the abstract factory and singleton.

3.1.2.1. Behavioral patterns

Behavioral patterns focus not only on how objects should be described but also on how they communicate with one another. We can further distinguish between class patterns and object patterns. Class patterns depend on the use of inheritance to distribute behavior among classes, object patterns depend on object composition to define how objects work together in order to perform tasks. In what follows we will briefly go over some example patterns.

| Name | Description |
|----------|---|
| Mediator | The Mediator pattern is a way of defining an object that can be used to coordinate how a set of objects interacts. This promotes loose coupling between objects. It consists of a Mediator class that defines the interface which will be used, Colleague classes which communicate with the Mediator object, and a ConcreteMediator class which implements the coordination of Colleague's. |
| Observer | The Observer pattern defines a one-to-many relationship between objects so that when a change of state occurs all dependents are automatically updated. Using this patterns allows for encapsulating different aspects of an abstraction into their own class, making it possible to let them vary and reuse them independently. The pattern consists of four parts: the Subject, Observer, ConcreteSubject and ConcreteObserver. The Subject knows who its observers are and provides an interface for managing them, the ConcreteSubject stores the state which is relevant for the related ConcreteObserver objects and notifies them when the state changes. The Observer defines the interface used for updating objects when notified by a subject, the ConcreteObserver implements this interface and maintains a link to its ConcreteSubject object to keep its state synced. |
| State | The State pattern allows an object to change its behavior based on the state it is in at run-time. Using this pattern allows each branch of a conditional statement to be encapsulated in its own class, allowing states to be treated as proper objects which can vary independently from other objects. This pattern consists of three parts: a Context, State and ConcreteState subclasses. The Context defines the interface which will be available for clients and maintains an instance of the ConcreteState subclass that represents the current state. The State class defines the interface for encapsulating the behavior associated with a particular Context state. Each of the ConcreteState subclasses implements the behavior associated with that state. |

3.1.2.2. Structural patterns

Structural patterns detail how classes and objects should be composed to form larger structures. Structural class patterns use inheritance to compose interfaces, which can be useful for making libraries work together or conforming interfaces. Structural object patterns on the other hand detail how to compose objects to realize new functionality. In what follows we will briefly go over some example patterns.

| Name | Description |
|-----------|--|
| Adapter | <p>The Adapter takes an interface and transforms it into another interface a client expects, i.e. it resolves compatibility issues.</p> <p>It is both a class and object pattern in which an Adapter takes the interface of an Adaptee and adapts such that it conforms to the interface of the Client.</p> |
| Bridge | <p>The Bridge pattern allows for an abstraction to be decoupled from its implementation so they can vary independently, meaning abstractions can be reused. This pattern consists of 4 classes: Abstraction, RefinedAbstraction, Implementor, ConcreteImplementor. The Abstraction defines an interface which defines higher level operations and maintains a reference to an object of the Implementor type, while the RefinedAbstraction class extends the interface as defined by Abstraction. The Implementor class declares an interface which is implemented by the ConcreteImplementor. This Implementor interface only provides basic operations which serve as basis for those of the Abstract class.</p> |
| Composite | <p>The Composite pattern composes objects into tree structures to represent part-whole hierarchies, allowing clients to treat compositions and individual objects uniformly. This pattern consists of an interface class Component as well as subclasses Leaf and Composite. The Component class implements default behavior for the interface common to all classes. The Leaf class represents leaf nodes, i.e. objects that have no children, and defines the operations which can be executed on base objects. The Composite class defines the behavior for components that can have children and maintains a list to all of its leaves.</p> |
| Facade | <p>A Facade is an object structural pattern which provides a unified interface to a set of interfaces in a subsystem. The Facade knows which part of the subsystem is responsible for handling certain requests and delegates client requests to the correct part of the subsystem. The subsystem classes themselves don't know the Facade exists.</p> |
| Proxy | <p>The Proxy pattern is an object structural pattern which provides proxies or placeholders to provide access control for other objects.</p> <p>The Proxy object maintains a reference to the real object in order to have access. It provides an interface to the outside world which is identical to that of the referred subject. We can further distinguish between remote proxy (handling requests), virtual proxy (cache information about the subject) and protection proxy (performs access control).</p> |

3.1.2.3. Creational patterns

Creational patterns help make a system independent of how its objects are created, composed, and represented. They work by encapsulating the knowledge about which concrete classes the system uses as well as hiding how instances of these classes are created. This means that the system in general only knows what is defined in the interface of the abstract class. In what follows we will briefly go over a few of the creational patterns.

| Name | Description |
|------------------|--|
| Abstract factory | <p>The abstract factory pattern provides an interface for creating families of related or dependent objects without specifying concrete classes.</p> <p>In this pattern we distinguish 5 classes: AbstractFactory, ConcreteFactory, AbstractProduct, ConcreteProduct and Client. The AbstractFactory declares an interface for operations to create abstract product objects, the ConcreteFactory then implements these operations to create concrete product objects. The AbstractProduct class declares an interface for a type of product objects while the ConcreteProduct class implements this interface in order to define a concrete product object to be created by the corresponding factory. The Client only uses the interfaces as declared by AbstractFactory and AbstractProduct.</p> |
| Factory method | <p>A factory provides an interface for creating objects in a superclass, but allows subclasses to alter the type of objects that will be created.</p> <p>This means that we have a class Product, which defines the interface, as well as the ConcreteProduct class which implements this interface. The Creator class declares the factory method, which returns an object of type Product, and may also define a default implementation of this method that returns a default ConcreteProduct object. The final class in the pattern is the ConcreteCreator which overrides the factory method from the Creator class to return a ConcreteProduct object.</p> <p>In this pattern, the Creator class relies on its derived classes to properly alter the factory method such that it returns the correct ConcreteProduct.</p> |
| Singleton | <p>The singleton pattern ensures there is only one instance of the class as well as provide a global access point. This is done by making the constructor a protected method in the class and declare a pointer to an instance as private member. This member is then accessed via the public Instance() method.</p> |

3.2. Normalised Systems Theory

Normalised Systems Theory, NST from here on, is an architecture that aims to provide software that is stable, maintainable and extendable.

In their methodology, the implementation process is viewed as a transformation I of functional requirements R into software primitives S such that $S = I(R)$. The transformation I can be separated into two categories, namely static transformation and dynamic transformation.

The **static transformation** defines how we can transform functional requirements into software primitives. To this end three requirements, as well as the corresponding transformations, are formulated.

- Requirement 1

An information system needs to be able to represent instances of data entities D_m . A data entity consists out of a number of data fields $\{a_i\}$ which may be a basic data field representing a value, or a reference to another data entity.

—→ Transformation implementation

Every data entity D_m is transformed into a data structure $S_m = I(D_m)$. This means that each data entity is instantiated as a software construct for data as provided by the programming language.

- Requirement 2

An information system needs to be able to execute processing actions P_n on instances of data entities. A processing action consists of a number of consecutive tasks $\{t_j\}$. Such a task may be a basic task such as a unit of processing that can change independently or an invocation of another processing action.

—→ Transformation implementation

Every processing action P_n is transformed into a processing function $F_n = I(P_n)$. This means that each processing action is instantiated as a software construct for processing as provided by the programming language.

- Requirement 3

An information system needs to be able to input or output values of instances of data entities through connectors C_l .

—→ Transformation implementation

Every I/O connector C_p of a data structure is transformed into a processing function $F_p = I(C_p)$ of the programming language, in the same way as a processing action, and makes use of the standard I/O functionalities of the language.

A software system thus becomes a set of software primitives S , consisting of a subset of data structures $\{S_m\}$ and a subset of related processing functions $\{F_n\}$.

The **dynamic transformation** ensures that evolving functional requirements can be caught. It does so by extending the static transformation with two additional requirements.

- An existing system representing a set of data entities $\{D_m\}$ needs to be able to represent both a new version of a data entity D_m that corresponds to including an additional data field a_i as well as a completely new data entity.
- An existing system providing a set of processing actions $\{P_n\}$ needs to be able to provide both a new version of a processing task t_j as well as an additional processing task and both a new version of a processing action P_n as well as an additional processing action.

3.2.1. Stability

To ensure stability in the software NST makes use of four design theorems: separation of concerns, data version transparency, action version transparency and separation of states.

3.2.1.1. Separation of concerns

Separation of Concerns is a theorem which is concerned with how tasks are implemented within processing functions. Taking into account that the goal of NST is to deliver software that is evolvable, we identify these tasks based on the concept of change drivers, i.e. a single concern within the application. In order to achieve stability, functions should not address more than one concern. This leads to the following formulation of the theorem.

Theorem 1 *A processing function can only contain a single task in order to achieve stability.*

This theorem can manifest itself in a number of ways.

A first manifestation is the use of an integration bus to manage communication between components

and/or applications. Take for instance a classical application model where each component communicates directly with the others. This would mean that, for N components, there would need to be $\frac{N(N-1)}{2}$ connectors to facilitate communication. This form of communication is in violation with the theorem as it forbids the direct transformation between two protocols as such a transformer would be subject to more than one change driver. If we use a change bus we only need to add one single connector, which does not violate the theorem and is considerably better when taking stability into account. Using a change bus also reduces the amount of work for adding a new connector from $O(N)$ to $O(1)$.

A second manifestation is the use of external workflows as the workflow sequence, i.e. the sequence in which a number of actions are performed, is a separate change driver from the way each of the actions are implemented. Viewing workflows as separate change drivers allows us to modify the workflow, e.g. changing the order in which we want to execute certain steps, without changing the way in which the actions are implemented.

A final manifestation is related to the software architectures. The use of a multi-layer architecture, as described in section 3.1.1.2, separates each change driver into a separate layer.

3.2.1.2. Data version transparency

The data version transparency theorem is concerned with how data entities are passed to processing functions. An entity has data version transparency if it can have multiple versions without affecting related processing functions: it should be possible to upgrade an entity without affecting related functions and methods that process it.

Theorem 2 *A data structure that is passed through the interface of a processing function needs to exhibit version transparency in order to achieve stability.*

This feature is in fact present in nearly every environment as this corresponds to the notion of polymorphism. Other notable examples include web services or microservices as detailed in section 3.1.1.3 and the concept of information hiding and encapsulation.

3.2.1.3. Action version transparency

This theorem is concerned with how functions are called by other functions. It means that functions can have multiple versions without affecting other functions that call it during their own execution, i.e. it should be possible to upgrade functions without affecting the rest of the system.

Theorem 3 *A processing function that is called by another processing function, needs to exhibit version transparency in order to achieve stability.*

In practice this theorem is present in concepts as polymorphism and the use of wrapper functions.

3.2.1.4. Separation of states

p283, 308

3.2.2. Maintainability

In order to ensure maintainability of a software system, NST implements the notion of diagnosability in order to counter the inevitable entropy of the system. This diagnosability dictates that every observable system state is deterministically traceable to the software primitives causing the state.

This is vital as the number of internal states increases during each function invocation, leading to an increased complexity in the state space of the system. In this state space we can distinguish between internal states and external states.

Internal states describe the system during execution of an invocation, e.g. the steps you see when running a debugger in an IDE such as IntelliJ, and consist of the following items.

- The values of the global data structure, e.g. the values which become available when the software starts and can be used during the entire lifetime of the program. They only cease to exist when the program ends.
- The values of the local data structure which are only available during the lifespan of an individual invocation. They may contain intermediate results or store values obtained from external sources or other invocations.

External states on the other hand are far less detailed and contain only

- The external inputs the application may take during its lifetime.
- The external outputs the application may generate during its lifetime.

It goes without saying that the inputs will be taken into the internal state of the system and will reside there until the end of the application's lifetime, and that the outputs generated by the application will go to live on outside of the application and may remain there after the application has stopped.

If we want to use diagnostics to make statements about execution of software we need to define what is considered "properly executed". In NST this has been defined as "the uncertainty related to the identification of the cause of the notion 'has not been executed properly' " where the notion 'has not been executed properly' corresponds to the occurrence of an irregular execution of a task which caused improper execution or failure of the software.

3.2.3. Extendability

p315, 329, 452

3.2.4. Elements and layers

p363

3.3. Refactoring: fixing what's broken

Refactoring is the art of improving the quality of existing code without changing the application's appearance to the outside world. The goal of refactoring is to make code easier to understand, maintain, and modify, while reducing the risk of introducing new bugs and tackling technical debt.

The need for refactoring can arise due to various reasons. For instance, the code may have become difficult to understand or modify due to its complexity or lack of proper documentation. It may also have become outdated or redundant due to changes in business requirements or technological advancements.

Some common refactoring techniques include identifying and isolating code smells, applying design patterns and principles, simplifying code and applying separation of concerns, and using tests to ensure correctness and consistency. Refactoring should be done incrementally and iteratively, focusing on small, manageable chunks of code that can be improved without causing major disruptions or breaking functionality.

Refactoring can be a challenging and time-consuming process, particularly for large and complex codebases. However, it should not be seen as a one-time task or a last-minute solution to fix code problems. Instead, it should be an integral part of the development process, focusing on continuous improvement, collaboration, and code quality.

Refactoring should not be confused with reengineering. Where refactoring is focused on improving the quality of existing code, the aim of reengineering is to alter the fundamental structure of a system or application. Refactoring can thus be used as a tool in reengineering efforts.

3.4. SonarQube

SonarQube is a tool that provides a detailed overview of numerous metrics when run on a codebase. In this thesis we will use SonarQube to obtain results on following metrics:

- Complexity: this value denotes the cyclomatic complexity of the codebase. A higher value indicates more branches in the codebase.
- Duplications: number of duplicated blocks of lines, files and percentage of duplicated lines. The expectancy is that these numbers will be relatively high when run on codebases following NST due to the way μ Radiant generates its code.
- Code smells: the total count of code smell issues found in the codebase.
- Technical debt: provides an estimate for the effort to fix all code smells.
- Maintainability rating: rating given to the project relative to the value of the technical debt ratio. This takes into account the outstanding remediation cost, i.e. the sum of the estimated time to fix all open issues. If this time is
 - $\leq 5\%$ of the time that has already gone into the application, the rating is A
 - $6 \leq time \leq 10$ the rating is B
 - $11 \leq time \leq 20$ the rating is C

- $21 \leq time \leq 50$ the rating is D
 - anything over 50% is an E
- Vulnerabilities: the number of vulnerability issues.
- Security rating:
 - A = 0 Vulnerabilities
 - B = at least 1 Minor Vulnerability
 - C = at least 1 Major Vulnerability
 - D = at least 1 Critical Vulnerability
 - E = at least 1 Blocker Vulnerability
- Security remediation effort: the effort to fix all vulnerability issues.

3.5. System under observation

In order to verify whether NST delivers on the promises it makes we decided to perform a small experiment.

We selected a small legacy system with as criteria that it should be written in Java and it should be a CRUD based system, basically eliminating games for example and music processing software. The system we chose was found on GitHub at https://github.com/VaibhavTyagi010/Online_CabBookingApp.

This Cab booking application was created using a specific kind of **multi-tier** architecture, namely Controller-Service-Repository. In this implementation the developers decided to extend this pattern with two additional layers, resulting in a total of five layers: Controller, Entity, Exception, Repository and Service.

3.5.1. Controller layer

Figure 3.4 provides a diagram of how the Control layer is implemented. As we can see the classes in the Control layer provide the CRUD, Create-Read-Update-Delete, operations needed for the application as well as the necessary getters and setters to allow access to the private members.

Classes in this layer define user-interface according to the REST principles. They relay requests they receive to the corresponding Service layer object.

3.5.2. Entity layer

Classes in the Entity layer define how data is stored in the database. Figure 3.5 provides an overview of the entities which are currently present in the application and how they are connected.

It is worth noting that in the current implementation, the use of inheritance breaks the Separation of

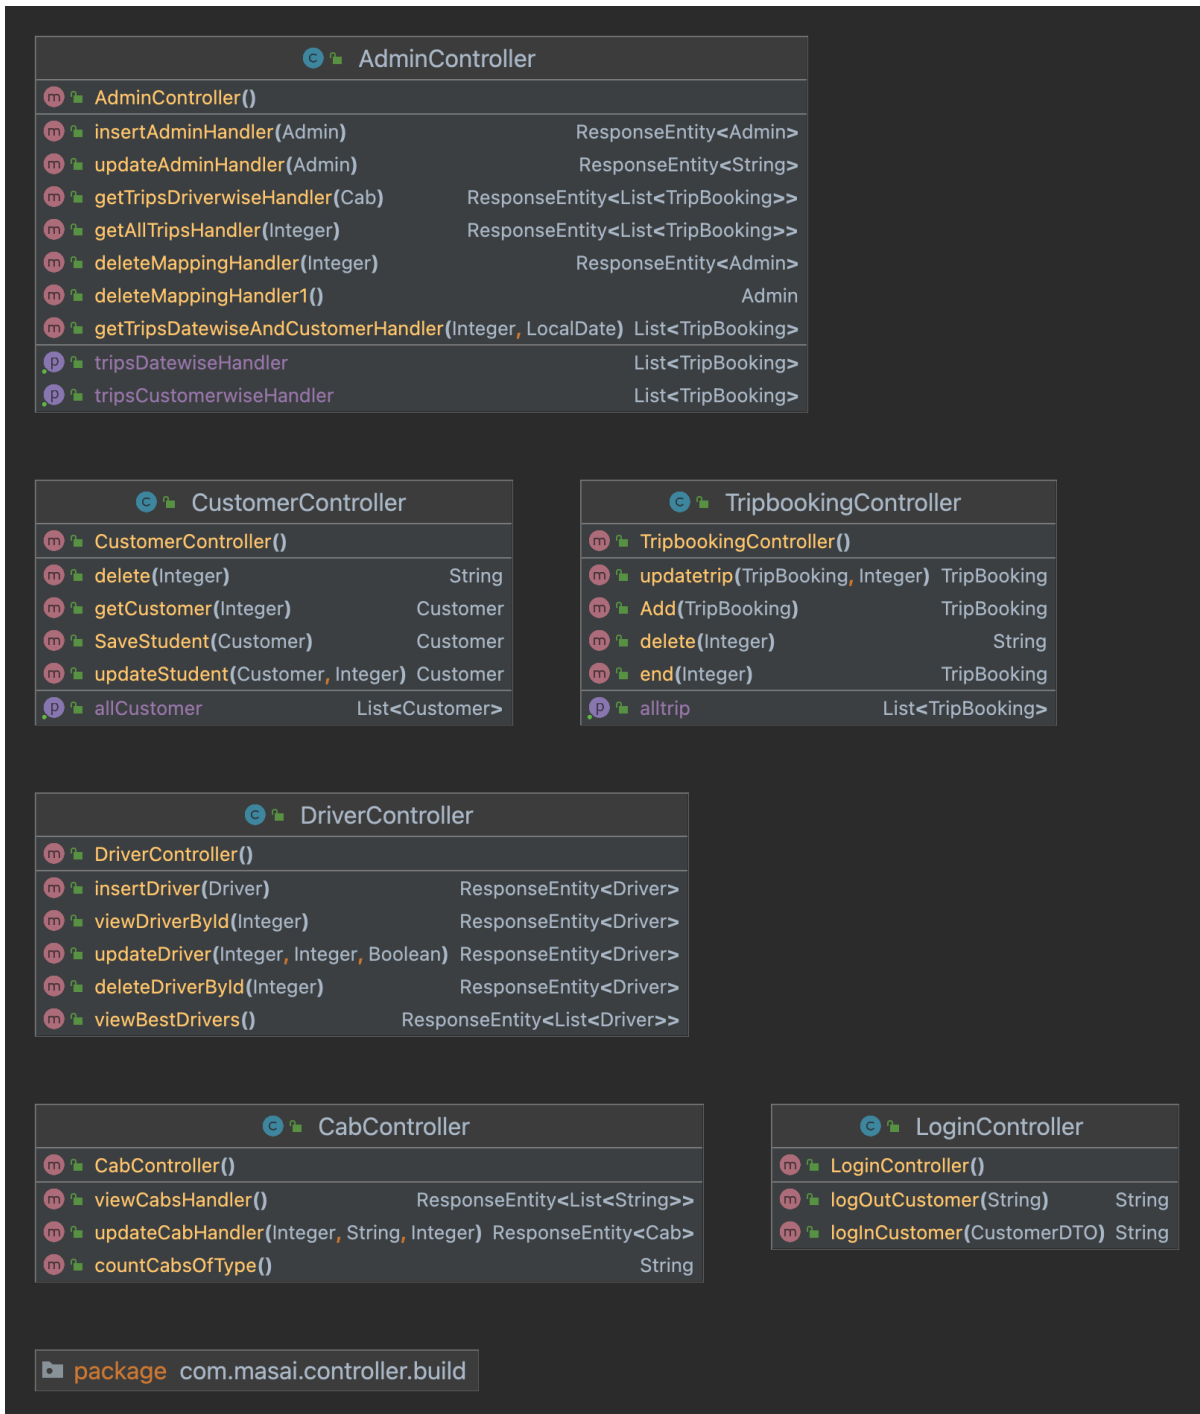


Figure 3.4.: Diagram of the control layer

Concerns principle. This can be circumvented by creating both the superclass and subclass as separate data elements and making the superclass entity a datamember of the subclass entity.

If we apply this to the current system, we would get data elements for Abstractuser, Customer, Admin and Driver where Customer, Admin and Driver would contain a pointer to an Abstractuser.

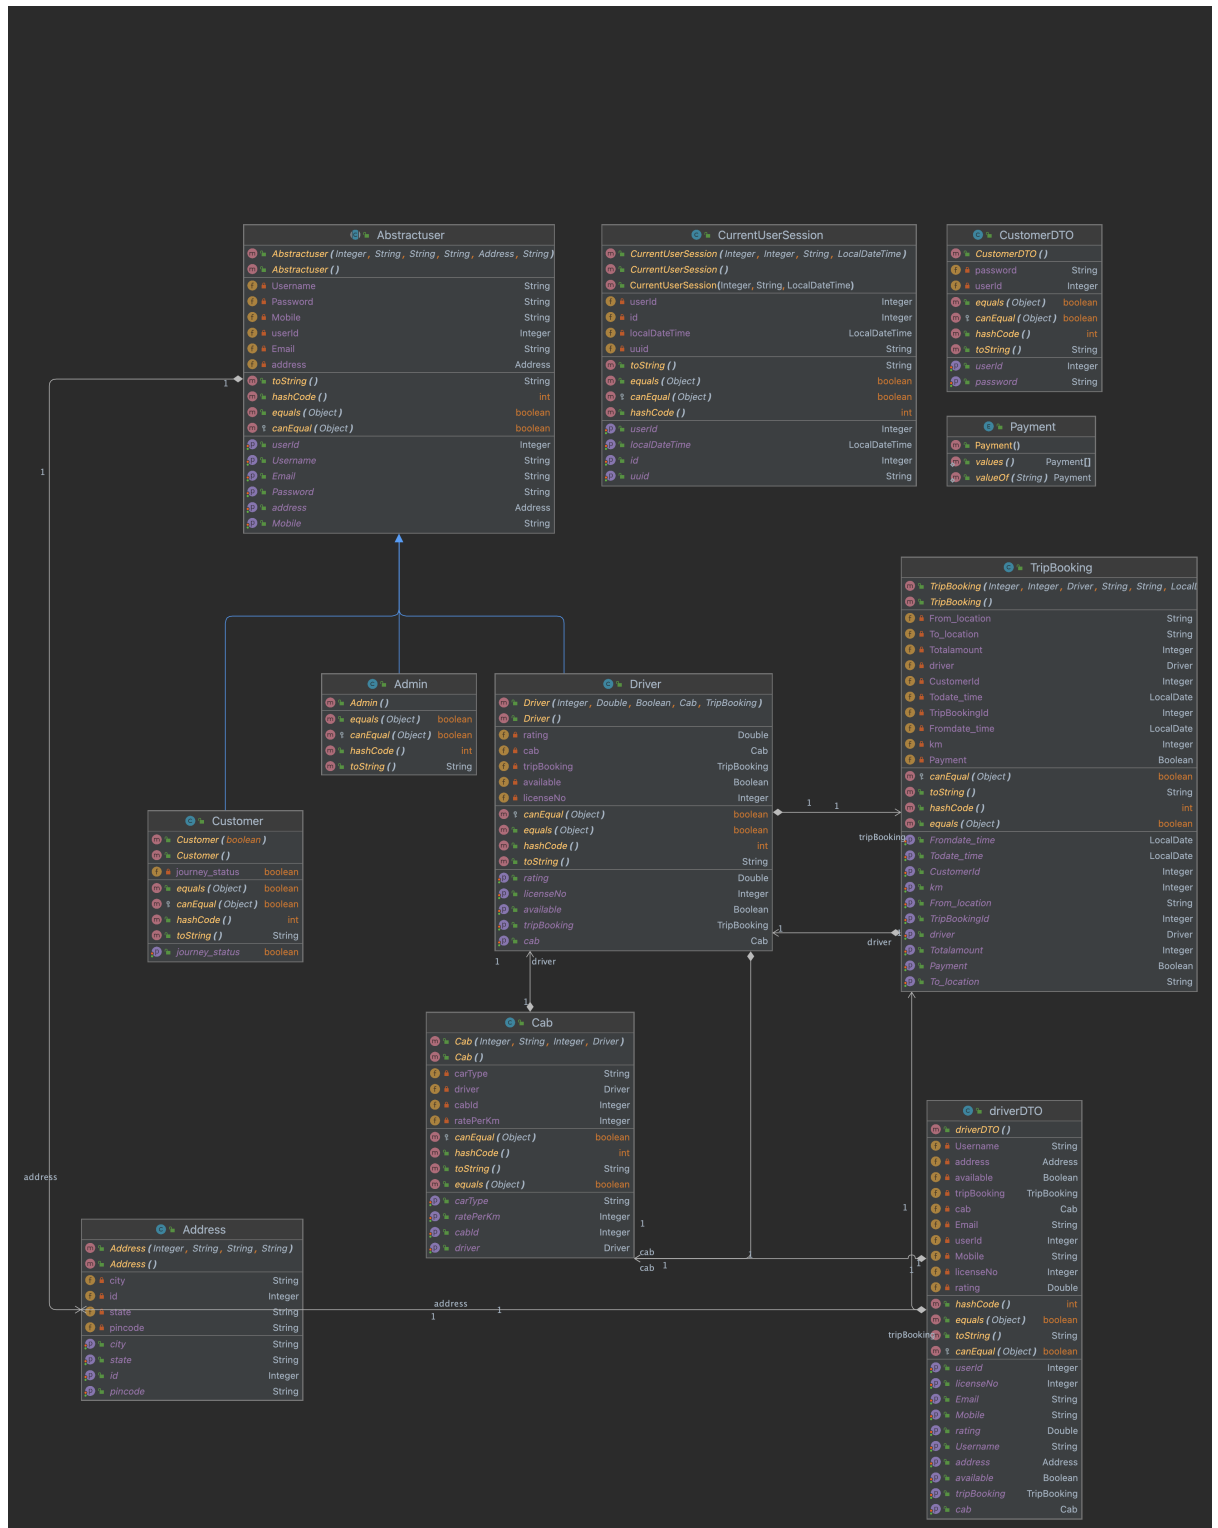


Figure 3.5.: Diagram of the entity layer

3.5.3. Exception layer

In Figure 3.6 we can see that some time has been put towards custom error handling.

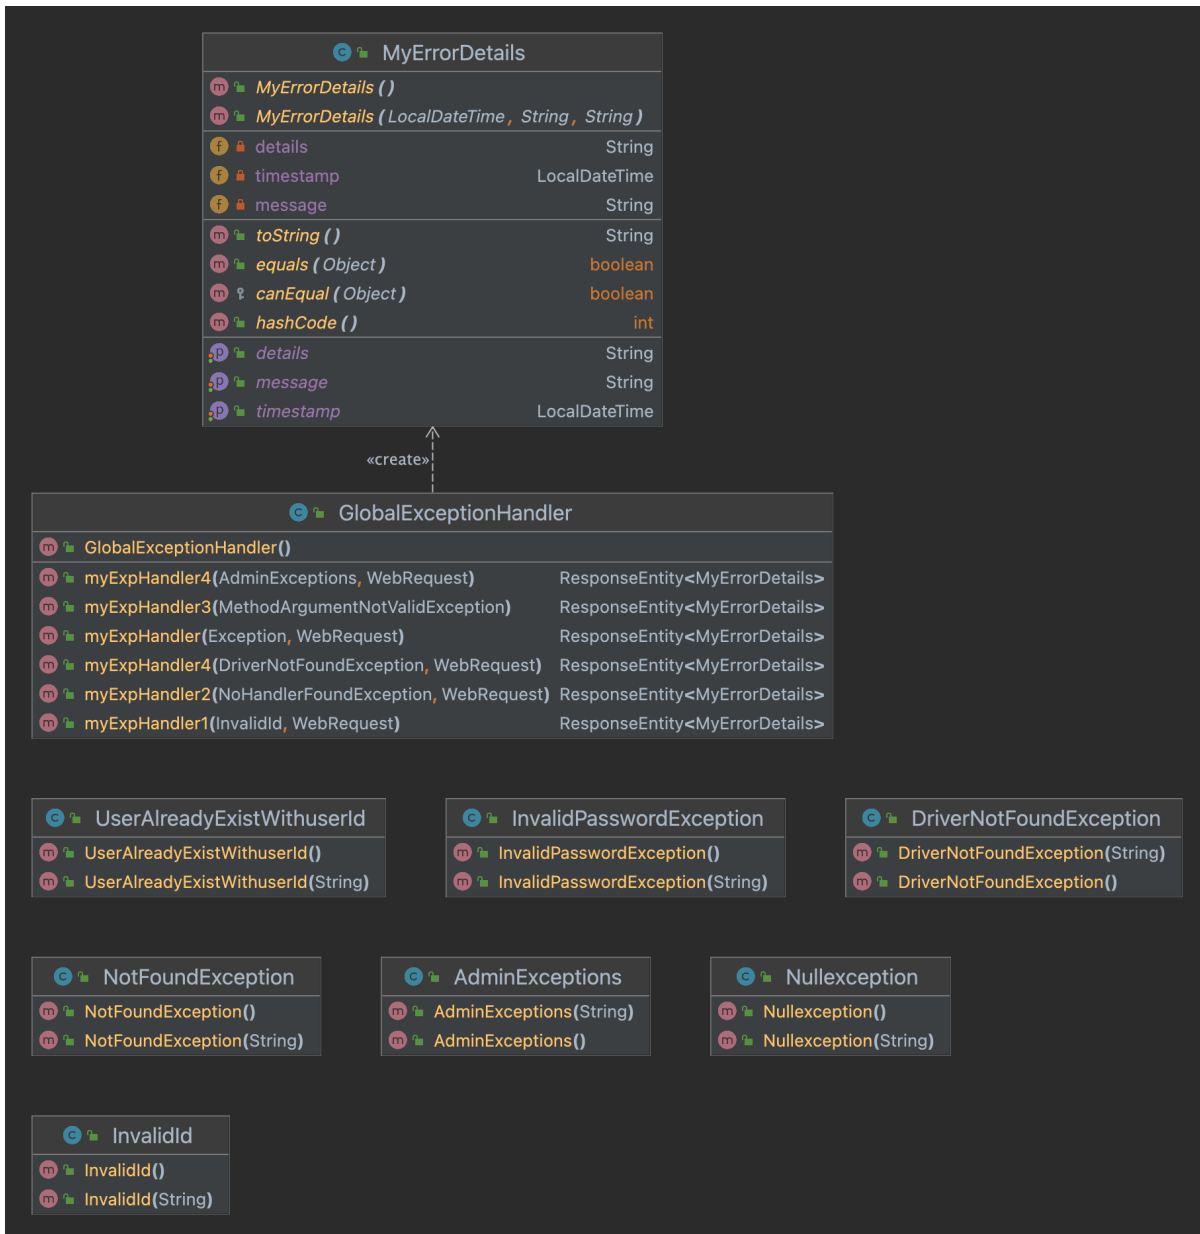


Figure 3.6.: Diagram of the exception layer

3.5.4. Repository layer

The DAOs in this layer, as given in Figure 3.7, provide an API through which the application can access the database layer.

If we transform this system into NST, these finders and views can be generated simply by adding them to the corresponding data element in μ Radiant.

3.5.5. Service layer

The service layer calls the API defined by the Repository layer to allow users to perform CRUD operations in the application. Figure 3.8 provides an overview of the services which are currently provided to

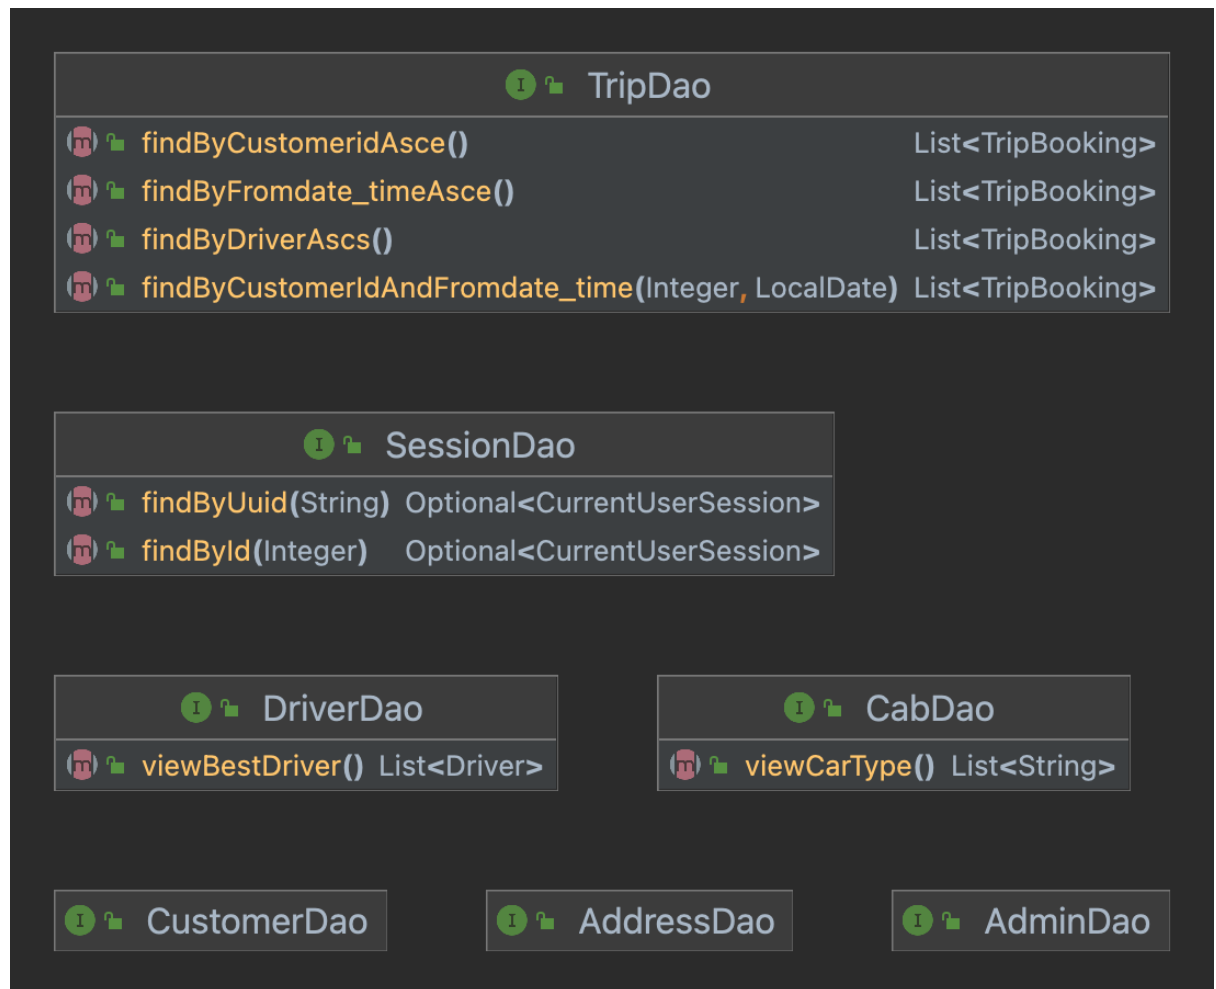


Figure 3.7.: Diagram of the repository layer

the user.

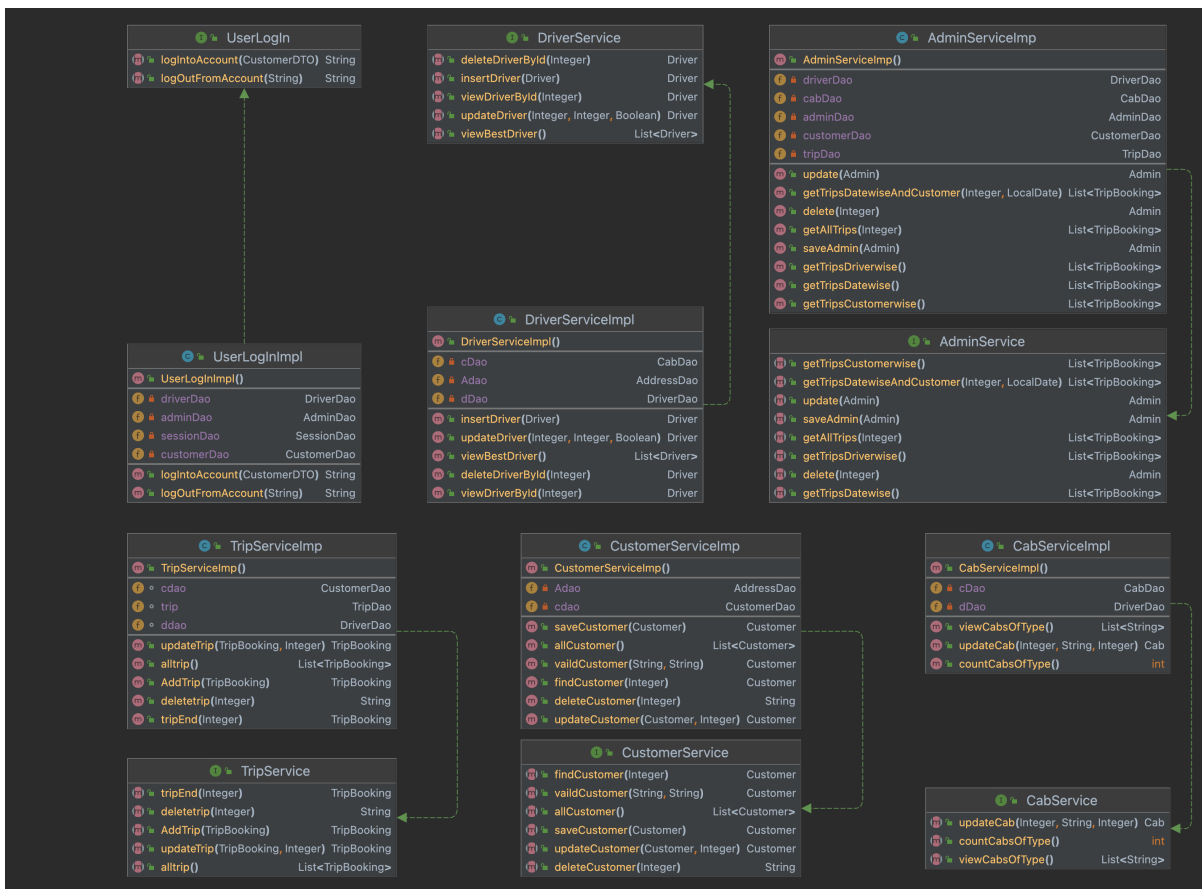


Figure 3.8.: Diagram of the service layer

4. Refactoring into NST

4.1. μ Radiant

μ Radiant is a tool which enables developers to model their NS application. As shown in Figure 4.1, an NS Application has the following structure

- The application is represented by an Application element.
- Each Application can have multiple ApplicationInstances. An ApplicationInstance is a combination of settings to expand the application with.
- Each Application can have multiple Components, and Components can be reused among multiple applications. Components group Data, Task, Flow Elements and ValueFieldTypes.
 - DataElements represent information in the system, which can be stored, queried and updated.
 - TaskElements represent executable pieces of logic. A task has a target DataElement. When executed, the task receives an instance of that DataElement as parameter.
 - Flow Elements are linked to a DataElement with a dedicated status field. The Flow Element describes a number of state transitions. Each transition has a begin-state and a TaskElement to execute if an instance of the DataElement reaches that state.
 - ValueFieldTypes allow the implementation of custom types for ValueFields in the DataElements of the Component.

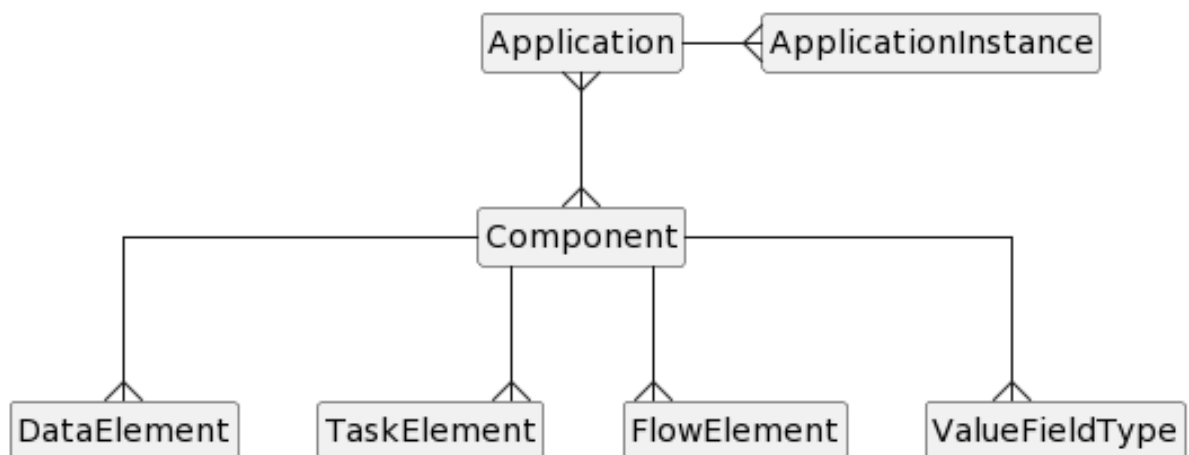


Figure 4.1.: Abstract structure of NS applications

4.2. The transformation process

In order to start our transformation process, we first need to collect a list of all data elements we will need in our application. As our goal is to recreate the system as closely as possible we decided to only take into consideration the data entities which were already present without adding new ones unless absolutely necessary.

4.2.1. Installation

In order to replicate what we will discuss further in this chapter, readers need to obtain a license from <https://foundation.stars-end.net> and proceed with the installation as described on <https://foundation.stars-end.net/docs/tools/micro-radiant/installation>.

After successful installation, μ Radiant can be started by executing following script

```
1 #!/bin/bash
2 cd <path/to/installation/location>/micro-radiant -l/
3 ./start.sh
4 open http://localhost:9050/models
```

This will open the editor in the users default browser and will look similar to Figure4.2. Before we

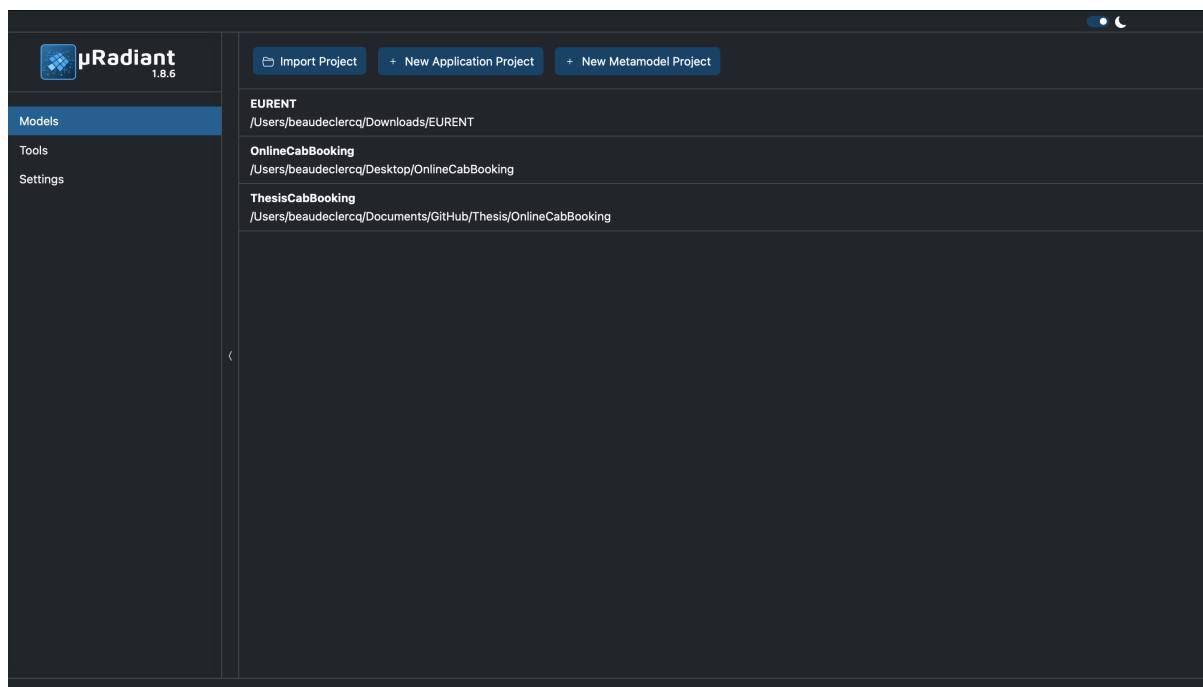


Figure 4.2.: μ Radiant homepage

continue with creating the project, we first need to determine which data elements we will need. After all, good preparation is half of success.

4.2.2. The Data Elements

After going through the different layers of the original application we decided to keep the following list of data elements

- **Address:** this element represents the addresses which will be used in the application. An address consists of a state, city, zipcode, street and house number.
- **Cab:** in accordance to the initial application, a cab has a rate per km, a reference to its driver and a reference to its car type. As the `viewCarType()` method from the `CabDao` simply returns a list of all unique car types currently available as cab, and we opted to create a separate data element for car types to better adhere to the separation of concerns principle, this method can be replaced by using a generic of custom finder in the `CarType` data element.
- **CarType:** currently a car type is nothing more than a string. The reason we decided to create a separate data element for this is that even though currently it only consists of a name, in the future we might feel the need to add more generic information to cabs of a certain type such as a predefined rate per km, a maximum number of occupants etc. By already separating in this stage we create the possibility for easy alteration later on.
- **Customer:** as in the initial system, a customer has a journey status and a link to an `AbstractUser` or `Person` as we renamed it.
- **Driver:** unlike a customer, a driver has a fair number of attributes. In addition to a link to a `Person`, a driver also has a license number, a rating, a link to the cab which it owns, its current booking and a flag to signify whether or not the driver is currently available.
- **Payment:** we decided to take the price of the trip booking out of the `TripBooking` element and put it in the more logical `Payment` element. Along with this total fare the `Payment` element also has a flag to signify whether or not it has been paid.
- **Person:** the `Person` data element represents the `AbstractUser` from the initial system. A person has a username, password, email address, mobile number and a physical address.
- **TripBooking:** a trip booking represents the core of the application. A booking entails a customer who has requested a driver to transport him between two locations at two specified pickup times. Based on the total distance in km and the rate per km from the driver, a payment will be created for the customer to pay.

4.2.3. Creating the project

Now that the preliminary work is completed, it is time to start the fun part: creating the application.

We do this by selecting the “New Application Project” from the home screen as shown in Figure4.2, which will lead us to the project creation screen as depicted in Figure4.3. The first four fields are fairly straightforward. A user selects the folder where the project needs to be saved, provide a name for the application itself and optionally a name for the `applicationInstance` as well as a name for the first component.

The `groupId` is used in the `POM.xml` file which will be generated and should correspond to that of your company. The name for the `dockerImage` is once again self explanatory.

μRadiant 1.8.6

← Home

Import Existing Project

New Application Project

New Metamodel Project

Project Directory: Select undefined

application
The name of the application.

applicationInstance
The name of the application instance. If not provided, the application name is used.

component
The name of the main component.

groupid
The groupid that will be used in the POM files.

dockerImage
The Docker image name.

> Advanced

Create Project

- ☒ **jenkins**
Create required files to build the project on Jenkins.
- ☒ **git**
Initialize a local git repository for the project and create an initial commit.
- ☒ **sonar**
Add SonarQube integration to the project.
- ☒ **docker**
Generate files to package the application as Docker image and deploy locally in Docker with docker-compose.
- ☐ **aws**
Generate Terraform scripts to deploy to Amazon Web Services.
- ☒ **baseComponents**
Add the base components as a dependency for the application.
- ☒ **gitFlow**
Initialize a setup that leverages the Git Flow workflow.
- ☒ **modelValidation**
Generate configuration to run model validations.
- ☒ **scripts**
Generate utility/demo scripts.
- ☐ **examples**
Generate example content for the project.

Figure 4.3.: Creating a new project

On the right hand side a number of boxes are checked by default and it is advised to keep them checked as they make life easier down the road. If wanted for demo purposes it might be useful to check the ‘examples’ checkbox to generate dummy data.

Once everything is filled out correctly, simply click ‘Create Project’ and you’re on your almost ready to do business. If you want to you can go ahead and already take a look at your project by selecting it from the μRadiant homescreen. The project home page will look similar to Figure 4.4.

ThesisCabBooking reload

μRadiant 1.8.6

← Home

Edit

Expansion

Customize

OnlineCabBooking::1.0.0 onlineCabBookingComp settings data

Filter elements...

Right-click to interact

- OnlineCabBooking
 - applicationInstances (1) [elements::App]
 - OnlineCabBooking

edit validations

OnlineCabBooking::1.0.0: OnlineCabBooking ApplicationInstance

name OnlineCabBooking

shortName OnlineCabBooking

version 1.0.0

description

globalOptionSettings 3.0 Modern

presentationSettings 3.1 Future

businessLogicSettings TomEE Postgres Hibernate

technicalInfrastructure Java 17

technologyStacks

options

- ☒ finders.enableTypedFinders ?
- ☒ hibernate.version ? 5
- ☒ struts.security.enforceHttpMethod ?

Figure 4.4.: Overview of an NS application

4.2.4. Adding Data Elements

Once the project has been created, we can start adding the data elements we gathered during our preliminary work to our component. To do so we navigate to our component from the application overview, it will look like Figure 4.5. Adding a data element is easy: right-click on the component in the finder pane,

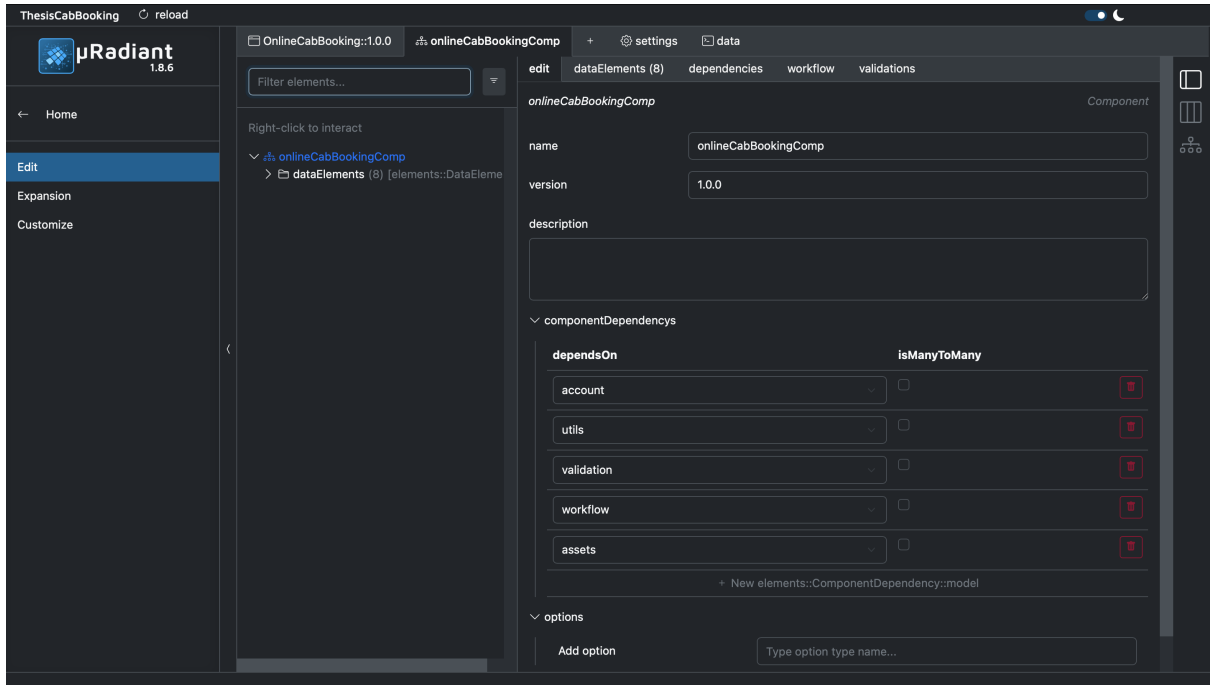


Figure 4.5.: Overview of a component in an application

select “New DataElement” and fill in the required fields as depicted in Figure 4.6.

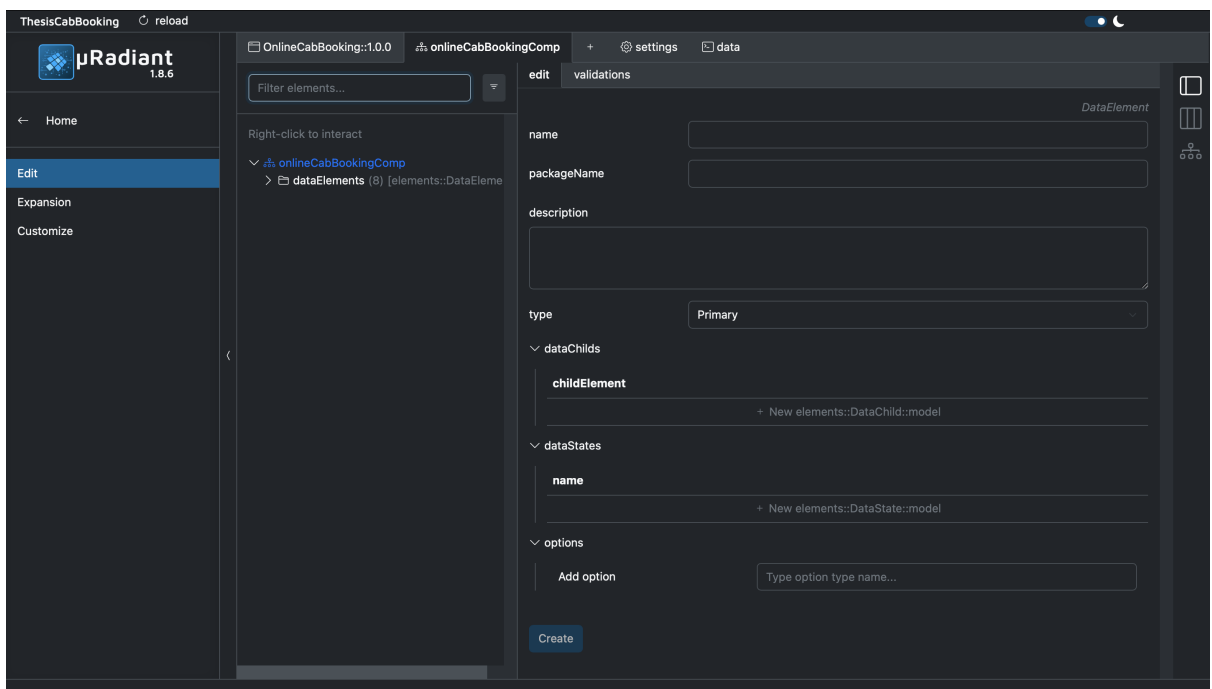


Figure 4.6.: Adding a data element to a component

TODO

Adding task elements and flow elements

Explanation on anchors and custom code

Explanation on finders, projections and custom finders

4.3. Future work

- Testing: at time of writing there is no functionality to automatically generate tests.
- Automatic import: currently there is only functionality to import existing NS projects. It might be useful to provide functionality which allows for automatic importing of data elements as most software architectures already provide some separation of data entities, business logic and visual representation.

5. Conclusions

5.1. Short term findings

| Metric | Original system | NST system |
|--------------------------------|-----------------|------------|
| Bugs | 3 | 0 |
| Reliability rating | C | A |
| Reliability remediation effort | 35m | 0m |
| Cognitive Complexity | 26 | 2104 |
| Code smells | 120 | 3124 |
| Technical debt | 6h26 | 35d |
| Debt ratio | 1.0% | 2.2% |
| Maintainability rating | A | A |
| Vulnerabilities | 8 | 0 |
| Security rating | D | A |
| Security remediation effort | 1h20 | 0 |

5.1.1. Explanations

5.1.1.1. Bugs

According to the SonarQube scan, the original version of the project contains 3 bugs which they classify as major. These bugs are

- In exception/GlobalExceptionHandler, the getFieldError() method could throw a NullPointerException and currently this behavior is not anticipated. Remedying this bug is estimated to take about 10 minutes.
- In service/UserLoginImpl 2 bugs are present, both in the logIntoAccount method. The first one refers to an issue on line 46 where the ID is being retrieved from an object obtained by using the get() method from the base Java library. As this method may throw a NoSuchElementException and this behavior is not anticipated, this is considered a major issue.
The second bug refers to the if-statement on line 50. This statement tries to resolve the before mentioned issue, but in the current implementation it makes little sense to have the check at that location. According to SonarQube remedying these bugs will take 10 and 15 minutes respectively. A possible solution to resolve both bugs would be to refactor the method as follows

```

1 public String logIntoAccount(CustomerDTO userDto) {
2     Optional<Customer> opt_customer = customerDao.findById(userDto.getUserId());
3     // Optional<Driver> opt_driver = driverDao.findById(userDto.getUserId());
4     // Optional<Admin> opt_admin = adminDao.findById(userDto.getUserId());
5     Integer userId = opt_customer.get().getUserId();
6     Optional<CurrentUserSession> currentUserOptional = sessionDao.findById(
7     userId);
8     if (!opt_customer.isPresent()) {
9         throw new AdminExceptions("user not found");
10    }
11    if (currentUserOptional.isPresent()) {
12        throw new UserAlreadyExistWithuserId("User already logged in with this
13        number");
14    }
15    if (opt_customer.get().getPassword().equals(userDto.getPassword())) {
16        String key = RandomString.make(6);
17        CurrentUserSession currentUserSession = new CurrentUserSession(
18        opt_customer.get().getUserId(), key,
19        LocalDateTime.now());
20        sessionDao.save(currentUserSession);
21
22        return currentUserSession.toString();
23    } else {
24        throw new InvalidPasswordException("Please Enter Valid Password");
25    }
26 }

```

```

1 public String logIntoAccount(CustomerDTO userDto) {
2     Optional<Customer> opt_customer = customerDao.findById(userDto.getUserId());
3     // Optional<Driver> opt_driver = driverDao.findById(userDto.getUserId());
4     // Optional<Admin> opt_admin = adminDao.findById(userDto.getUserId());
5     if (opt_customer.isPresent()) {
6         Integer userId = opt_customer.get().getUserId();
7         Optional<CurrentUserSession> currentUserOptional = sessionDao.findById(
8         userId);
9         if (currentUserOptional.isPresent()) {
10            throw new UserAlreadyExistWithuserId("User already logged in with this
11            number");
12        }
13        if (opt_customer.get().getPassword().equals(userDto.getPassword())) {
14            String key = RandomString.make(6);
15            CurrentUserSession currentUserSession = new CurrentUserSession(
16            opt_customer.get().getUserId(), key,
17            LocalDateTime.now());
18            sessionDao.save(currentUserSession);
19
20            return currentUserSession.toString();
21        } else {
22            throw new InvalidPasswordException("Please Enter Valid Password");
23        }
24    } else {
25        throw new AdminExceptions("user not found");
26    }
27 }

```

The NST version of the same project currently does not contain any bugs.

5.1.1.2. Reliability rating and remediation

Due to the bugs mentioned in the previous point, the original project has received a reliability rating of C. The total time to resolve all bugs and go from C to A is estimated to be 35 minutes, which is still an acceptable amount of time given the nature of the bugs.

The NST version of the project already has a rating of A as there are no bugs found.

Conclusion: when it comes down to bugs, using NST limits the possibility for introducing them due to the use of expanders during code generation/rejuvenation.

5.1.1.3. Cognitive complexity

The cognitive complexity of an application refers to how difficult it is to understand it. According to the white paper which can be retrieved at <https://www.sonarsource.com/resources/cognitive-complexity/>, the cognitive complexity is increased when encountering loop structures, conditionals, catches, nested structures, switches and recursion. For example, the use of a switch statement with multiple cases will result in a lower cognitive complexity then when using a if-else if structure.

The original system has a cognitive complexity of 26, all of which is located in the service layer of the application, while the NST version has a cognitive complexity of 2104 with following distribution

| | | | |
|---------|-----|-------|-----|
| Control | 468 | Data | 856 |
| Logic | 152 | Proxy | 416 |
| Shared | 212 | View | 0 |

This high complexity is due to the way in which code is currently generated by the expanders. For example, in the create method from the DriverCruds class try-catch structures are used in combination with nested conditionals. This use of if-else structures can be justified by the need for anchor points between which developers can add custom code.

5.1.1.4. Code smells

Our scan of the original project revealed the presence of 120 code smells in the code base. Some of the revealed tags are

| Code smell | Original system | NST system |
|--|-----------------|------------|
| Unused: these code smells refer to deletion of commented code blocks and unused imports. | 38 | 2k |
| Clumsy: clumsy code smells indicate situations that could have been implemented in more efficient ways, such as specifying types in constructors via diamond constructors, immediately returning expressions in stead of storing them in temporary variables and using appropriate methods to check for emptiness. | 38 | 146 |
| Convention: reports violations of coding conventions such as formatting and naming. | 29 | 283 |
| Bad practice: these issues will work as intended, but the way in which they are implemented are acknowledged to be bad practice. | 3 | 101 |
| CERT: reports violations of CERT standard rules as can be found at https://wiki.sei.cmu.edu/ . | 5 | 226 |
| CWE: relates to rules in the Common Weakness Enumeration | 16 | 258 |
| Design: reveals questionable design choices, such as having duplicates of literals in stead of using a variable for it. | 4 | 80 |

Fixing all code smells in the original system would take an estimated 6h 26 minutes while the NST system would require 35 days of work.

It is actually quite surprising that the system following the NST principles has such a high number of code smells, so let's take a closer look at some of the categories.

- **Convention:** almost all issues with this tag refer to the renaming of the package from cabBooking-Core to cabbookingcore. This can easily be done in the μ Radiant editor and can be fixed in less than 5 minutes. The estimated time to fix all these issues is estimated to be 5 days and 5 hours.
- **Unused:** SonarQube reports over 2k code smells of unused code. The main issue here is that, while these are indeed correctly identified as currently unused, not all of them can be removed due to the way in which the code is generated by the expanders. For example, imports are included on a "these are most commonly used in this context" basis. Therefore it's not possible to resolve these smells as they are created by design. This means that of the estimated 35 days needed to resolve the code smells, 12 will always be there.
- **CWE:** the majority of these issues fall in one of three categories. A first category is one where the different agents created to provide access to the different objects should be declared as either transient or serializable. A next category would be one removal of deprecated methods. A final category is avoiding the use of generic exceptions such as RuntimeException. Code smells that fall in either of the first two categories can be resolved by updating the expanders. Code smells that fall in the third category are harder to resolve as they would require the automatic creation of per-class exceptions which can then be further customized by the developer. This would lead to an exponential increase in number of classes which is both impractical and undesirable.
In total it would take 9 days and 7 hours to resolve all smells with the CWE tag. If we take into account that about a third of the workload cannot be resolved by design, there will always be an outstanding cost of at least 3 days.

- Bad practice: all code smells with this tag can be resolved by either adding a `@deprecated` tag or `@Override` annotation in the expanders. SonarQube gives an estimated cost of 1 day.
- Design: all except one of the design code smells refer to the use of duplicate string literals. The estimated cost for fixing these smells is 1 day 6 hours.

5.1.1.5. Technical debt

As stated in the previous paragraph, it would take 35 days to fix all code smells from our NST system. We can solve almost 6 days of work in less than 5 minutes using μ Radiant and an additional 9 days of work can be taken care of by tweaking the expanders. Of the 35 days needed to fix all code smells at least 15 days can not be resolved due to the design of NST.

In the previous discussion we have not mentioned CERT, CWE and clumsy code smells. The reason for this is that both CERT and CWE smells are also tagged with other tags which have already been mentioned and their remediation cost has thus already been counted. The clumsy smells require a case-by-case inspection: some could be fixed quite easily in the expanders, such as merging nested conditional statements, while others cannot be resolved due to the presence of anchors as is the case in for instance the `getCustomDisplayName` method from the `CarTypeCruds` class. As these smells only account for a workload of 7 hours to remedy and overall they are classified as minor issues, it might not be feasible to spend time at tweaking the expanders just for the sake of reducing code smells. The remaining 5 days of work are distributed among various other tags which we will not discuss here for sake of interest.

Our initial system on the other hand can be released of its code smells in less than 7 hours.

We already know that the technical debt of software will increase over time if a developer does not consciously tries and avoid it. The main benefit of NST is that, since the majority of the code is generated/rejuvenated by μ Radiant, the technical debt will almost never increase as fast as otherwise.

5.1.1.6. Maintainability rating

The maintainability rating is closely related to the technical debt ratio. As both versions have a debt ratio below 5%, both versions get a maintainability rating of A.

5.1.1.7. Vulnerabilities

In our initial system, SonarQube found 8 vulnerability issues, all of which entail replacing a persistent entity with a POJO or DTO object. Our NST version comes with zero vulnerability issues.

5.1.1.8. Security rating and remediation

As all of the vulnerabilities found in the initial system are deemed to be critical issues, this version has received a security rating of D which is the second last lowest rating. Luckily the estimated remediation effort is 1 hour 20 minutes, so we can easily improve this rating from D to A.

Since our NST version has no detected vulnerabilities, it received a rating of A.

5.1.1.9. Conclusion

While at first glance the NST version of the system looks to be way worse than our initial version - according to SonarQube we would need well over a month to remedy all issues! - things aren't as bad at all.

Indeed, if we take a closer look at the two areas in which NST seems to perform poorly - namely cognitive complexity and code smells - it becomes clear that many of the issues can be resolved by tweaking the expanders by for instance switching from try-catch statements to switch statements in order to lower the cognitive complexity or resolve 8 days of remediation cost from code smells, thereby providing cleaner code.

The main question here would be the following: is it worth recreating a system using the μ Radiant tool? The answer is: it depends.

If an architectural blueprint is present - which really should be the case in a professional setting - it should not take long to transform an application using μ Radiant. The advantage obtained from doing so, i.e. obtaining a system build according to the four design principles given in section 3.2.1 which can be mathematically proven [VM03] to deliver software with the desired properties, outweigh the initial high cognitive complexity and high amount of code smells.

5.2. In the long run

As the long term benefits of transforming a system according to NST versus traditional refactoring cannot be measured during this thesis, this part is not covered at time of writing.

6. Previous research projects

6.1. Research project 1

Title/subject: Sketch recognition

Promotor: Hans Vangheluwe

Related to this thesis: No.

6.2. Research project 2

Title/subject: Visualizing large TCR networks

Promotor: Pieter Meysman and Sofie Giellis

Related to this thesis: No.

A. Bibliography

Bibliography

- [Bab21] Shibu Babuchandran. Software release life cycle (srhc): Understand the 6 main stages, 2021. [Online; publication date Dec 6 2021]. 4
- [DSN03] Serge Demeyer, Ducasse Stephane, and Oscar Nierstrasz. *Object-oriented reengineering patterns*. Morgan Kaufman Pub., 2003. 11
- [ECPB12] Thomas Erl, Benjamin Carlyle, Cesare Pautasso, and Raj Balasubramanian. *SOA with REST Principles, Patterns Constraints for Building Enterprise Solutions with REST*. Pearson, 2012. 10
- [Ext21] Extensibility. Extensibility — Wikipedia, the free encyclopedia, 2021. [Online]. 4
- [GHJ⁺94] Erich Gamma Gamma, Richard Helm, Ralph Johnson, John Vlissides, and Grady Booch. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley Professional, 1994. 11
- [NGI] NGINX Glossary. What is load balancing? [Online]. 8
- [Pie] Michael Pierce. Software release life cycle (srhc): Understand the 6 main stages. [Online; no publication date]. 3
- [Ric15] Mark Richards. *Software Architecture Patterns Understanding Common Architecture Patterns and When to Use Them*. O'Reilly, 2015. 8
- [sof] What is software rot, and why does it happen? [Online; no publication date or author]. 3
- [Sof23] Software rot. Software rot — Wikipedia, the free encyclopedia, 2023. [Online]. 3
- [sta21] The importance of stability and reliability testing in software development, 2021. [Online]. 3
- [tec] What is technical debt and what should you do about it? [Online; no publication date or author]. 3
- [VM03] Jan Verelst and Herwig and Mannaert. *Normalised Systems Theory*. Morgan Kaufman Pub., 2003. 36