

Academic year  
2022 - 2023

# Software Design: Normalised Systems Theory versus Refactoring

Beau De Clercq

Master's thesis  
**Master of Science in computer science: software engineering**

Supervisors  
**prof. S. Demeyer, UAntwerpen**  
**prof. H. Mannaert, UAntwerpen**



University of Antwerp  
Faculty of Science

#### **Disclaimer Master's thesis**

This document is an examination document that has not been corrected for any errors identified. Without prior written permission of both the supervisor(s) and the author(s), any copying, copying, using or realizing this publication or parts thereof is prohibited. For requests for information regarding the copying and/or use and/or realisation of parts of this publication, please contact to the university at which the author is registered.

Prior written permission from the supervisor(s) is also required for the use for industrial or commercial utility of the (original) methods, products, circuits and programs described in this thesis, and for the submission of this publication for participation in scientific prizes or competitions.

This document is in accordance with the master thesis regulations and the Code of Conduct. It has been reviewed by the supervisor and the attendant.

# Contents

<b>1. Introduction</b>	<b>3</b>
1.1. Stability . . . . .	3
1.2. Maintainability . . . . .	3
1.3. Extensibility . . . . .	4
<b>2. Research question</b>	<b>5</b>
<b>3. Defining the parameters</b>	<b>7</b>
3.1. Software Architectures . . . . .	7
3.1.1. Architectural patterns . . . . .	7
3.1.2. Design patterns . . . . .	11
3.2. Normalised Systems Theory . . . . .	14
3.2.1. Stability . . . . .	15
3.2.2. Maintainability . . . . .	17
3.2.3. Extendibility . . . . .	18
3.2.4. Elements and layers . . . . .	20
3.3. Refactoring: fixing what's broken . . . . .	20
3.4. SonarQube . . . . .	22
3.4.1. Other tools . . . . .	23
3.5. System under observation . . . . .	25
3.5.1. Controller layer . . . . .	26
3.5.2. Entity layer . . . . .	26
3.5.3. Exception layer . . . . .	26
3.5.4. Repository layer . . . . .	26
3.5.5. Service layer . . . . .	26
<b>4. Refactoring into NST</b>	<b>33</b>
4.1. $\mu$ Radianit . . . . .	33
4.2. The transformation process . . . . .	34
4.2.1. Installation . . . . .	34
4.2.2. The Data Elements . . . . .	35
4.2.3. Creating the project . . . . .	35
4.2.4. Adding Data Elements . . . . .	37
4.2.5. Adding Flow and Task Elements . . . . .	38
4.2.6. Creating the workflow . . . . .	41
4.2.7. On the $\mu$ Radianit Flow . . . . .	42
4.2.8. On Anchors and Custom Code . . . . .	43
4.2.9. On Finders, Projections and Custom Finders . . . . .	43
4.3. Future work . . . . .	45
<b>5. Conclusions</b>	<b>47</b>
5.1. Short term findings . . . . .	47
5.1.1. Explanations . . . . .	47
5.2. In the long run . . . . .	54

<b>6. Previous research projects</b>	<b>55</b>
6.1. Research project 1 . . . . .	55
6.2. Research project 2 . . . . .	55
<b>A. Bibliography</b>	<b>57</b>

## **Abstract**

**Software Design: Normalised Systems Theory versus Refactoring**  
Beau De Clercq



# 1. Introduction

The idea for this thesis came after many years of frustrations, because for some reason software has always been a pain to work with: you could have one version that's been working for years but when you finally decide to upgrade to a newer version it turns out that you need to install a bunch of extra dependencies; that one application you've been using at your job just keeps growing and being extended, but that also means it starts getting slower and starts behaving more buggy with every update; ...

In the current software landscape, properties such as maintainability, extensibility and stability are highly valued and necessary to avoid code rot [Sof23]. In this chapter we will define the concepts of stability, maintainability and extensibility, cover why they are important in software and provide ways on how to detect them.

**TODO:** extend part

## 1.1. Stability

Software stability refers to the reliability and robustness of a system, which ensures that it performs consistently and predictably under various conditions without crashing or producing unexpected errors. It is a critical aspect of software quality, as unstable software can lead to data loss, system downtime, and user dissatisfaction.

One key factor in achieving software stability is rigorous testing [sta21]. Robust testing processes, including unit testing, integration testing, and regression testing, help identify and fix bugs and vulnerabilities before the software is released to users. Testing also helps uncover edge cases, unusual inputs, and stress conditions that could cause instability.

Other important aspects of software stability are design and architecture. Well-designed software with a clear architecture and modular components is less likely to develop stability issues. Good design practices, such as separation of concerns, loose coupling, and high cohesion, contribute to the stability of software systems. Additionally, adhering to best coding practices, such as error handling, exception handling, and defensive programming, can help prevent software instability caused by unexpected situations.

## 1.2. Maintainability

The maintainability of a software system refers to the ease with which it can be maintained, i.e., how easy is it to keep the system running and up to date. Take for instance the Agile development process, which has become increasingly popular over the past few years, in which we have a near continuous cycle [Pie] of maintaining, updating and upgrading. If software becomes unmaintainable the entire workflow is at risk of getting disrupted.

As the lifecycle of a software system goes on, due to added functionalities and improvements, the overall code complexity increases and with increased complexity comes an increase in the amount of time needed to maintain the system. Another factor that leads to this increase in maintenance time is technical debt.

Technical debt [tec][sof] is a term that refers to the cost you will pay in the future by trying to realise a quick fix in the present. This can happen for instance when programmers are under pressure to deliver a specific piece of software, resulting in corners being cut or general low code quality. When this happens there often is no plan as to when the debt will be re-paid and debt will silently accumulate as time goes on. The best way to counter technical debt is to keep in mind the long term effects certain changes may have.

A very nice analogy for software and maintenance over time are plants. When you first buy a plant it is compact and well defined, but within a few months this compact plant starts growing in every available direction - the addition of new functionalities and improvements of existing functionalities. And while there are no issues at first, you might notice that in some parts your plant is really getting out of hand - the accumulation of technical debt and general code deterioration. So in an effort to keep your plant maintainable you clip the problematic branches - or in the case of software, you start refactoring.

### 1.3. Extensibility

Extensibility [Ext21][Bab21] refers to ease with which a system or application can be extended - i.e. updated, customized or upgraded - without needing to make large alterations to the existing codebase. One common way to achieve software extensibility is by working with APIs, i.e. interfaces and protocols that allow external components or plugins to interact with the software.

One of the key benefits of software extensibility is the ability to adapt and evolve software over time. As requirements change or new technologies emerge, software can be extended to accommodate these changes without requiring a complete rewrite or overhaul. This enables software applications to remain relevant and competitive in a fast-paced and ever-changing technology landscape. Additionally, software extensibility promotes collaboration and innovation by allowing third-party developers to contribute new functionalities or integrate with existing software, fostering a rich ecosystem of plugins, extensions, or integrations that can enhance the overall value and utility of the software.

Another advantage of software extensibility is its potential to empower end-users to customize software to their specific needs. With extensible software, users can customize the functionality, appearance, or behavior of the software to align with their unique requirements or preferences. This can result in a more personalized and user-friendly experience, leading to increased user satisfaction. Moreover, software extensibility can enable users to create their own plugins or extensions, fostering a community-driven approach to software development and fostering a sense of ownership and empowerment among users.

## 2. Research question

The goal of this thesis is to look at the feasibility of refactoring a legacy system so they comply to the NST principles and verify that both the maintainability and the extensibility of the system have improved. To this end we will do the following

- Take a (legacy) system and refactor it using a selected method such that the system complies to the NST principles.
- Add new features to both the old and refactored system and compare the overall ease of implementing these new features as well as the code health determined by SonarQube.

The research question can thus be formulated as follows: If we take a legacy system and transform it so that it complies with the principles and guidelines defined by Normalised Systems Theory, will the promised benefits be visible immediately or will we only be able to tell a difference on the long run?

In what follows we make a few assumptions

- Refactoring a system in such a way that it complies with the principles of NST should be possible in a reasonable amount of time.
- Once refactored, adding new features to the system should be much easier. (= Extensibility)
- Once refactored, updating existing features should be as easy or easier (= Mainainability).
- Once refactored, there should not be a negative impact on the performance of the system.



# 3. Defining the parameters

In this chapter we will define the parameters of this thesis. These parameters as we like to call them are as follows

- Normalised Systems Theory: Normalised Systems Theory provides principles and guidelines which, if followed correctly, result in software which is stable, maintainable and extensible.
- Software Architectures: As Normalised Systems Theory defines an architecture, we will first give a brief overview of the different kinds of architectures and patterns.
- Refactoring: The line between recreating an application, also known as reengineering, and refactoring is fairly thin. To better distinguish between them we will see how they are related to each other.
- SonarQube: As this will be the basis on which we will compare our systems, we introduce what SonarQube is and how we intent to use it.
- System under observation: The system with which we would like to verify our research question.

Combined these parameters define this thesis.

## 3.1. Software Architectures

Just as with the physical architecture, a software architecture is the entirety of all components and their placements needed to build a product. To do so in the most efficient way possible, over the years numerous architectural patterns have been developed to provide clear guidelines on how to solve a certain set of problems. In addition to these architectural patterns, a multitude of design patterns have been developed to provide a clear structure to each of the components of the architectural pattern.

### 3.1.1. Architectural patterns

Architectural patterns define the foundation on which the software will be build: they provide general solutions to recurring problems. Among the more well known are client-server, multi-layer, microservices and REST.

It is perfectly possible to use architectural patterns in combination with one another. For example, in the microservice pattern a developer could decide to implement one service according to the multi-layer pattern, e.g. an application to rent a car, another one with REST, e.g. the front-end to the rental application, and yet another one as client-server, e.g. the location where all invoices are stored.

### 3.1.1.1. Client-server

In the client-server pattern, as the name suggests, we have clients which communicate with servers such as a web server or file server. Whether a computer is a client, a server, or both, is determined by the nature of the application that requires the service functions. For example, a single computer can run a web server and file server software at the same time to serve different data to clients making different kinds of requests. The client software can also communicate with server software within the same computer. In order to support frequent/high traffic, the application is often run on multiple servers.

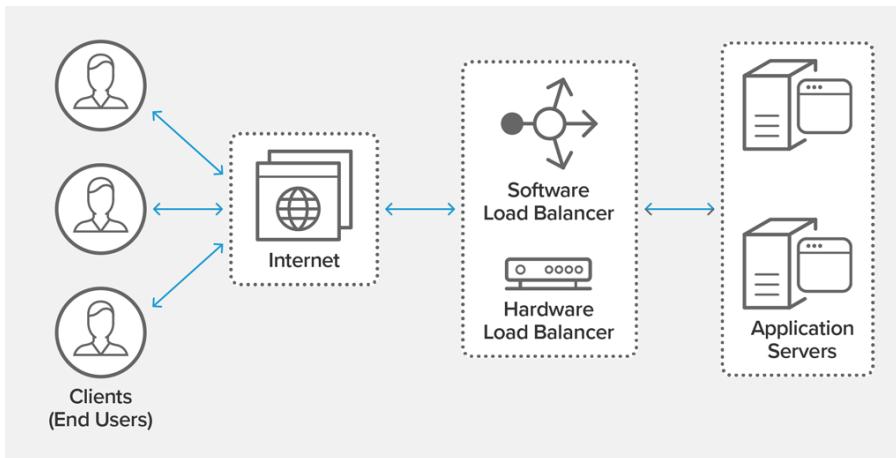


Figure 3.1.: Load balancing diagram

Clients then connect to one of the servers through a load balancer. This load balancer distributes all incoming client traffic among the available servers, thus ensuring that no server receives more than it can handle. An example of how a load balancer can be placed is shown in Figure 3.1 [NGI].

### 3.1.1.2. Multi-layer

A multi-layer or multitier architecture is a client-server architecture in which functionalities are physically separated over multiple layers. The most common layer division [Ric15] is Presentation-Business-Persistence-Database. An example is shown in Figure 3.2 [Ric15]. One of the powerful features of the layered architecture pattern is separation of concerns among components. Components within a specific layer deal only with logic that belongs to that layer. For example, components in the presentation layer deal only with presentation logic, whereas components in the business layer deal only with business logic.

In this architecture each layer is also closed, meaning that as a request moves from layer to layer, it must go through the layer right below it to get to the next layer below that one. The reason we don't allow each layer to have direct access to all others is due to the concept layers of isolation. This concept means that if changes are made in one of the layers, these changes don't impact or affect the other layers. This concept also means that each layer is independent of the rest, thereby having no knowledge of how the other layers work internally.

### 3.1. SOFTWARE ARCHITECTURES

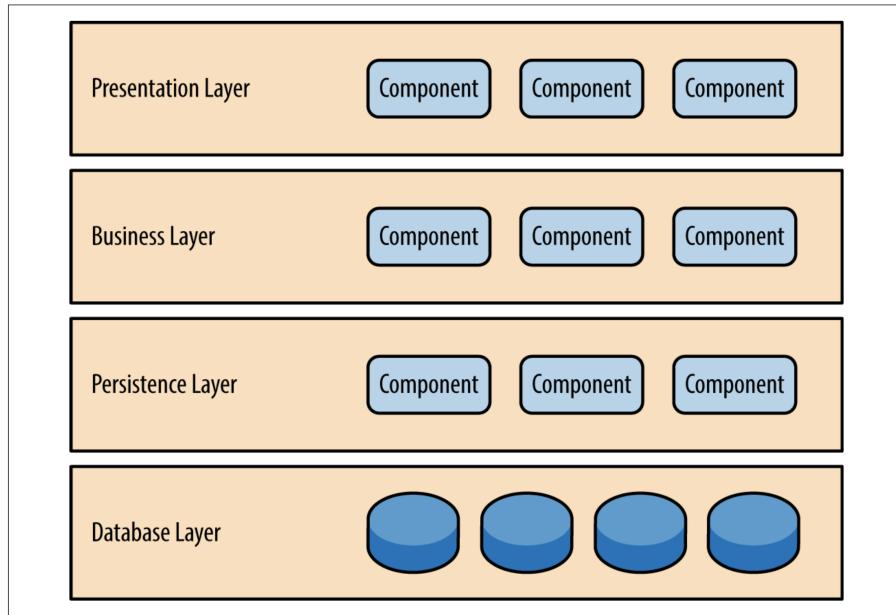


Figure 3.2.: Example of multi-tier architecture

#### 3.1.1.3. Microservices

The microservice architectural style is an approach to developing a single application as a suite of small services where each service runs in its own process and can communicate via lightweight mechanisms.

Microservices are built around business capabilities and are independently deployable by fully automated deployment machinery. They can be written in different programming languages and can use different data storage technologies. When using microservices there is only the bare minimum of centralized management. As an example, consider a webshop as depicted in Figure 3.3. In this example

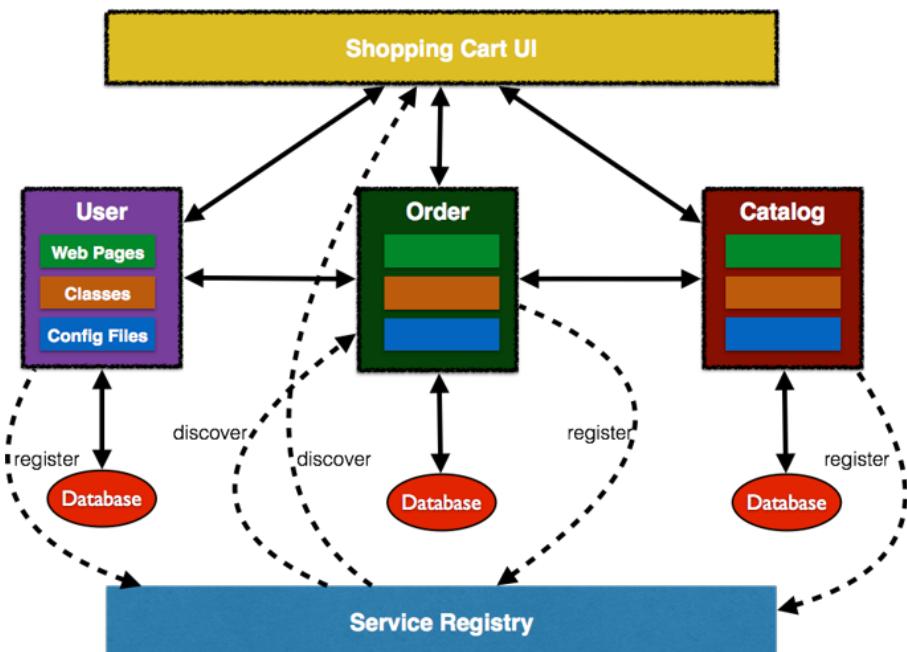


Figure 3.3.: Example of microservice architecture

our application consists of 3 components: Catalog, Order, User. Each component has its own databases so that each microservice can evolve and choose whatever type of datastore – relational, NoSQL, flat file, in-memory or some thing else – is most appropriate. Each component will register with a Service Registry. This is required because multiple stateless instances of each service might be running at a given time and their exact endpoint location will be known only at the runtime. Client interaction for the application is defined in another application, Shopping Cart UI in our case. This application mostly discover the services from Service Registry and composes them together. It should mostly be a dumb proxy where the UI pages of different components are invoked to show the interface.

At time of writing, the most well-known and most used framework for microservices is Docker. Docker works by running software in containers which can then communicate with each other via API calls. Each container can then be developed separately.

### 3.1.1.4. REST

Representational state transfer [ECPB12] (REST) is a style of software **architecture** for distributed systems such as the World Wide Web.

The term resource is a central concept that denotes any item of interest identified by some identifier. An identifier, also called Uniform Resource Identifier (URI), can be either a URL<sup>1</sup> (the location where you can find the object or the method for finding it) or a URN<sup>2</sup> (which defines an item identity).

A RESTful API is a Web Service that adheres to the REST style, meaning

- It has an Internet media type for data (XML, JSON, ...)
- It has a base URI
- It uses hypertext links to reference state
- It uses standard HTTP methods (GET, PUT, POST, DELETE)
- It uses hypertext links to reference related resources

The REST style states six constraints which, if adhered to, provide an architecture which has properties such as performance, scalability, simplicity, modifiability, visibility, portability, and reliability. These constraints are as follows

- Client–server architecture: Employing the client–server architecture enforces the principle of separation of concerns by separating UI from data storage, allowing components to evolve independently.
- Statelessness: When communicating with a server, each packet sent by the client can be understood separately from the others as it will contain all relevant session data, i.e. no session data is kept by the server.
- Cacheability: Responses received by the client must identify themselves as either cacheable or non-cacheable, this to prevent clients from providing stale or inappropriate data in response to further requests. If used correctly, caching can in part or fully eliminate some client–server interactions.

---

<sup>1</sup>Uniform Resource Locator

<sup>2</sup>Uniform Resource Name

### 3.1. SOFTWARE ARCHITECTURES

- Layered system: Clients should not be able to tell whether it is connected to an end-server directly or to a proxy, i.e. the use of load balancers and the addition of security layers won't impact client communications.
- Uniform interface: A uniform interface enables the architecture to be decoupled. Its constraints are the following
  - Resource identification in requests: resources are identified in requests via URIs in RESTful webservices, e.g. the server could send data as XML or JSON.
  - Resource manipulation through representations: when a client has a representation of a resource, it has enough information to work with the resource's state.
  - Self-descriptive messages: each message includes enough information to describe how to process it.
  - Hypermedia as the engine of application state: starting from a base URI, a client can dynamically access all resources it needs. I.e., starting from a homepage a client can access resources further down the hierarchy.
- Optional - Code on demand: Servers can temporarily alter client functionality by transferring executable code such as JavaScript.

#### 3.1.2. Design patterns

As the name suggests, design patterns provide a pattern for how the actual implementation should look. The difference between an architectural pattern and a design pattern is the level on which it is applied: architectural patterns always end up on top of design patterns.

Design patterns can roughly be divided into three main categories [GHJ<sup>+</sup>94][DSN03], namely behavioral patterns such as the observer, structural patterns such as adapter and proxy and creational patterns such as the abstract factory and singleton.

##### 3.1.2.1. Behavioral patterns

Behavioral patterns focus not only on how objects should be described but also on how they communicate with one another. We can further distinguish between class patterns and object patterns. Class patterns depend on the use of inheritance to distribute behavior among classes, object patterns depend on object composition to define how objects work together in order to perform tasks. In what follows we will briefly go over some example patterns.

Name	Description
Mediator	The Mediator pattern is a way of defining an object that can be used to coordinate how a set of objects interacts. This promotes loose coupling between objects. It consists of a Mediator class that defines the interface which will be used, Colleague classes which communicate with the Mediator object, and a ConcreteMediator class which implements the coordination of Colleague's.
Observer	The Observer pattern defines a one-to-many relationship between objects so that when a change of state occurs all dependents are automatically updated. Using this patterns allows for encapsulating different aspects of an abstraction into their own class, making it possible to let them vary and reuse them independently. The pattern consists of four parts: the Subject, Observer, ConcreteSubject and ConcreteObserver. The Subject knows who its observers are and provides an interface for managing them, the ConcreteSubject stores the state which is relevant for the related ConcreteObserver objects and notifies them when the state changes. The Observer defines the interface used for updating objects when notified by a subject, the ConcreteObserver implements this interface and maintains a link to its ConcreteSubject object to keep its state synced.
State	The State pattern allows an object to change its behavior based on the state it is in at run-time. Using this pattern allows each branch of a conditional statement to be encapsulated in its own class, allowing states to be treated as proper objects which can vary independently from other objects. This pattern consists of three parts: a Context, State and ConcreteState subclasses. The Context defines the interface which will be available for clients and maintains an instance of the ConcreteState subclass that represents the current state. The State class defines the interface for encapsulating the behavior associated with a particular Context state. Each of the ConcreteState subclasses implements the behavior associated with that state.

### 3.1.2.2. Structural patterns

Structural patterns detail how classes and objects should be composed to form larger structures. Structural class patterns use inheritance to compose interfaces, which can be useful for making libraries work together or conforming interfaces. Structural object patterns on the other hand detail how to compose objects to realize new functionality. In what follows we will briefly go over some example patterns.

### 3.1. SOFTWARE ARCHITECTURES

Name	Description
Adapter	The Adapter takes an interface and transforms it into another interface a client expects, i.e. it resolves compatibility issues. It is both a class and object pattern in which an Adapter takes the interface of an Adaptee and adapts such that it conforms to the interface of the Client.
Bridge	The Bridge pattern allows for an abstraction to be decoupled from its implementation so they can vary independently, meaning abstractions can be reused. This pattern consists of 4 classes: Abstraction, RefinedAbstraction, Implementor, ConcreteImplementor. The Abstraction defines an interface which defines higher level operations and maintains a reference to an object of the Implementor type, while the RefinedAbstraction class extends the interface as defined by Abstraction. The Implementor class declares an interface which is implemented by the ConcreteImplementor. This Implementor interface only provides basic operations which serve as basis for those of the Abstract class.
Composite	The Composite pattern composes objects into tree structures to represent part-whole hierarchies, allowing clients to treat compositions and individual objects uniformly. This pattern consists of an interface class Component as well as subclasses Leaf and Composite. The Component class implements default behavior for the interface common to all classes. The Leaf class represents leaf nodes, i.e. objects that have no children, and defines the operations which can be executed on base objects. The Composite class defines the behavior for components that can have children and maintains a list to all of its leaves.
Facade	A Facade is an object structural pattern which provides a unified interface to a set of interfaces in a subsystem. The Facade knows which part of the subsystem is responsible for handling certain requests and delegates client requests to the correct part of the subsystem. The subsystem classes themselves don't know the Facade exists.
Proxy	The Proxy pattern is an object structural pattern which provides proxies or placeholders to provide access control for other objects. The Proxy object maintains a reference to the real object in order to have access. It provides an interface to the outside world which is identical to that of the referred subject. We can further distinguish between remote proxy (handling requests), virtual proxy (cache information about the subject) and protection proxy (performs access control).

#### 3.1.2.3. Creational patterns

Creational patterns help make a system independent of how its objects are created, composed, and represented. They work by encapsulating the knowledge about which concrete classes the system uses as well as hiding how instances of these classes are created. This means that the system in general only knows what is defined in the interface of the abstract class. In what follows we will briefly go over a few of the creational patterns.

Name	Description
Abstract factory	<p>The abstract factory pattern provides an interface for creating families of related or dependent objects without specifying concrete classes.</p> <p>In this pattern we distinguish 5 classes: AbstractFactory, ConcreteFactory, AbstractProduct, ConcreteProduct and Client. The AbstractFactory declares an interface for operations to create abstract product objects, the ConcreteFactory then implements these operations to create concrete product objects. The AbstractProduct class declares an interface for a type of product objects while the ConcreteProduct class implements this interface in order to define a concrete product object to be created by the corresponding factory. The Client only uses the interfaces as declared by AbstractFactory and AbstractProduct.</p>
Factory method	<p>A factory provides an interface for creating objects in a superclass, but allows subclasses to alter the type of objects that will be created.</p> <p>This means that we have a class Product, which defines the interface, as well as the ConcreteProduct class which implements this interface. The Creator class declares the factory method, which returns an object of type Product, and may also define a default implementation of this method that returns a default ConcreteProduct object. The final class in the pattern is the ConcreteCreator which overrides the factory method from the Creator class to return a ConcreteProduct object.</p> <p>In this pattern, the Creator class relies on its derived classes to properly alter the factory method such that it returns the correct ConcreteProduct.</p>
Singleton	<p>The singleton pattern ensures there is only one instance of the class as well as provide a global access point. This is done by making the constructor a protected method in the class and declare a pointer to an instance as private member. This member is then accessed via the public Instance() method.</p>

## 3.2. Normalised Systems Theory

Normalised Systems Theory, NST from here on, is an architecture that aims to provide software that is stable, maintainable and extendable.

In their methodology, the implementation process is viewed as a transformation  $I$  of functional requirements  $R$  into software primitives  $S$  such that  $S = I(R)$ . The transformation  $I$  can be separated into two categories, namely static transformation and dynamic transformation.

The **static transformation** defines how we can transform functional requirements into software primitives. To this end three requirements, as well as the corresponding transformations, are formulated.

- Requirement 1

An information system needs to be able to represent instances of data entities  $D_m$ . A data entity consists out of a number of data fields  $\{a_i\}$  which may be a basic data field representing a value, or a reference to another data entity.

→ Transformation implementation

Every data entity  $D_m$  is transformed into a data structure  $S_m = I(D_m)$ . This means that each data entity is instantiated as a software construct for data as provided by the programming language.

### 3.2. NORMALISED SYSTEMS THEORY

- Requirement 2

An information system needs to be able to execute processing actions  $P_n$  on instances of data entities. A processing action consists of a number of consecutive tasks  $\{t_j\}$ . Such a task may be a basic task such as a unit of processing that can change independently or an invocation of another processing action.

→ Transformation implementation

Every processing action  $P_n$  is transformed into a processing function  $F_n = I(P_n)$ . This means that each processing action is instantiated as a software construct for processing as provided by the programming language.

- Requirement 3

An information system needs to be able to input or output values of instances of data entities through connectors  $C_l$ .

→ Transformation implementation

Every I/O connector  $C_p$  of a data structure is transformed into a processing function  $F_p = I(C_p)$  of the programming language, in the same way as a processing action, and makes use of the standard I/O functionalities of the language.

A software system thus becomes a set of software primitives  $S$ , consisting of a subset of data structures  $\{S_m\}$  and a subset of related processing functions  $\{F_n\}$ .

The **dynamic transformation** ensures that evolving functional requirements can be caught. It does so by extending the static transformation with two additional requirements.

- An existing system representing a set of data entities  $\{D_m\}$  needs to be able to represent both a new version of a data entity  $D_m$  that corresponds to including an additional data field  $a_i$  as well as a completely new data entity.
- An existing system providing a set of processing actions  $\{P_n\}$  needs to be able to provide both a new version of a processing task  $t_j$  as well as an additional processing task and both a new version of a processing action  $P_n$  as well as an additional processing action.

#### 3.2.1. Stability

To ensure stability in the software NST makes use of four design theorems: separation of concerns, data version transparency, action version transparency and separation of states.

##### 3.2.1.1. Separation of concerns

Separation of Concerns is a theorem which is concerned with how tasks are implemented within processing functions. Taking into account that the goal of NST is to deliver software that is evolvable, we identify these tasks based on the concept of change drivers, i.e. a single concern within the application. In order to achieve stability, functions should not address more than one concern. This leads to the following formulation of the theorem.

**Theorem 1** *A processing function can only contain a single task in order to achieve stability.*

This theorem can manifest itself in a number of ways.

A first manifestation is the use of an integration bus to manage communication between components

and/or applications. Take for instance a classical application model where each component communicates directly with the others. This would mean that, for  $N$  components, there would need to be  $\frac{N(N-1)}{2}$  connectors to facilitate communication. This form of communication is in violation with the theorem as it forbids the direct transformation between two protocols as such a transformer would be subject to more than one change driver. If we use a change bus we only need to add one single connector, which does not violate the theorem and is considerably better when taking stability into account. Using a change bus also reduces the amount of work for adding a new connector from  $O(N)$  to  $O(1)$ .

A second manifestation is the use of external workflows as the workflow sequence, i.e. the sequence in which a number of actions are performed, is a separate change driver from the way each of the actions are implemented. Viewing workflows as separate change drivers allows us to modify the workflow, e.g. changing the order in which we want to execute certain steps, without changing the way in which the actions are implemented.

A final manifestation is related to the software architectures. The use of a multi-layer architecture, as described in section 3.1.1.2, separates each change driver into a separate layer.

### 3.2.1.2. Data version transparency

The data version transparency theorem is concerned with how data entities are passed to processing functions. An entity has data version transparency if it can have multiple versions without affecting related processing functions: it should be possible to upgrade an entity without affecting related functions and methods that process it.

**Theorem 2** *A data structure that is passed through the interface of a processing function needs to exhibit version transparency in order to achieve stability.*

This feature is in fact present in nearly every environment as this corresponds to the notion of polymorphism. Other notable examples include web services or microservices as detailed in section 3.1.1.3 and the concept of information hiding and encapsulation.

### 3.2.1.3. Action version transparency

This theorem is concerned with how functions are called by other functions. It means that functions can have multiple versions without affecting other functions that call it during their own execution, i.e. it should be possible to upgrade functions without affecting the rest of the system.

**Theorem 3** *A processing function that is called by another processing function, needs to exhibit version transparency in order to achieve stability.*

In practice this theorem is present in concepts as polymorphism and the use of wrapper functions.

### 3.2.1.4. Separation of states

p283, 308

The Separation of States theorem is concerned with how calls between processing functions are handled.

### 3.2. NORMALISED SYSTEMS THEORY

What this theorem tells us is that, keeping in mind evolvability of the system, we need a separate data structure to store error states which may arise at runtime. For example, assume we have a function that performs a certain retrieval operation. If we change the implementation of this function its interface will remain the same since we need to comply with Data Version Transparency and Action Version Transparency. This new implementation may however lead to a new error state which requires certain handling. As the number of calls to this retrieval function may keep growing, so too will the number of error states grow. If we then have a separate data structure to handle error states we can simply trigger the corresponding action from the data structure irrespective of what the calling function was, i.e. we don't need to provide endless try-catch clauses. This leads us to the following formulation of the theorem

**Theorem 4** *Calling a processing function within another processing function, needs to exhibit state keeping in order to achieve stability.*

This theorem mainly manifests itself in asynchronous communication and processing. A more relevant manifestation are workflow systems as the principle of Separation of States corresponds to stateful workflow where intermediate states from the workflow are stored internally. Think of this as a state machine where the presence of a certain state will trigger the execution of the processing function of that state.

#### 3.2.2. Maintainability

In order to ensure maintainability of a software system, NST implements the notion of diagnosability in order to counter the inevitable entropy of the system. This diagnosability dictates that every observable system state is deterministically traceable to the software primitives causing the state.

This is vital as the number of internal states increases during each function invocation, leading to an increased complexity in the state space of the system. In this state space we can distinguish between internal states and external states.

Internal states describe the system during execution of an invocation, e.g. the steps you see when running a debugger in an IDE such as IntelliJ, and consist of the following items.

- The values of the global data structure, e.g. the values which become available when the software starts and can be used during the entire lifetime of the program. They only cease to exist when the program ends.
- The values of the local data structure which are only available during the lifespan of an individual invocation. They may contain intermediate results or store values obtain from external sources or other invocations.

External states on the other hand are far less detailed and contain only

- The external inputs the application may take during its lifetime.
- The external outputs the application may generate during its lifetime.

It goes without saying that the inputs will be taken into the internal state of the system and will reside there until the end of the application's lifetime, and that the outputs generated by the application will go to live on outside of the application and may remain there after the application has stopped.

If we want to use diagnostics to make statements about execution of software we need to define what is considered “properly executed”. In NST this has been defined as “the uncertainty related to the identification of the cause of the notion ‘has not been executed properly’ ” where the notion ‘has not been executed properly’ corresponds to the occurrence of an irregular execution of a task which caused improper execution or failure of the software.

### 3.2.3. Extendibility

To ensure extendibility in a system, NST focuses on modularization and how to keep modularity combinatorics under control. In this section we will briefly discuss the law of exponential variation gains and law of exponential ripple cost as well as their impact on extendibility.

In what follows we will use following notations and conventions

- $U_i$ : The unit of work in a system or module which cannot be used or activated separately.
- $M_i$ : Module in a system or artifact which can be used or activated separately.
- $k_i$ : Number of variants or versions of a specific unit of work or module  $i$ .
- $d_i$ : Number of other modules that are dependent on module  $i$  in a system.
- $c_i$ : Number of copies of a specific uniet of work  $i$  in other modules.

#### 3.2.3.1. Cohesion: Law of exponential variation gains

Consider a system which consists of  $n$  units of work  $U_i$ . In general the need often arises for different variants or versions of a unit of work. Suppose our system contains  $N$  different units of work and that every such unit  $U_i$  exists in  $k_i$  different versions or variants, then we have a few modular strategies to cope with these variations.

We could opt to go with one global monolithic system module in which all versions of all units of work would coexist. Having  $k_i$  variants, this single module would contain

$$\sum_{i=1}^N k_i$$

system part versions. This module would also contain all the logic glue to keep everything together and choose which version to use. If we then want to introduce a new version of any unit of work, we need at least

$$\sum_{i=1}^N (k_i - 1)$$

upgrades for the overall module. It is clear that this way of working would lead to a rather complicated monolithic module and should be avoided.

As alternative option we could choose to work with a different variant or version of the single global module for every specific combination of variants for the various units of work. In case all possible combinations of variants would be required and provided, this would lead to

$$\prod_{i=1}^N k_i$$

### 3.2. NORMALISED SYSTEMS THEORY

different variants of the aggregation module. If every possible combination of variants would be provided, every unit of work  $i$  would be duplicated

$$\prod_{j=1}^{i1} k_j \cdot \prod_{j=i+1}^N k_j$$

times. This also implies that any change on unit of work  $i$  would imply the duplication of that change as many times as there are duplications of that unit of work.

Of course in practice not all possible combinations will be required nor will they be provided, but it is clear that the efforts for development and maintenance for such a system would be exponentially dependent on the number of individual variants.

The more elegant - and efficient - solution would be to use separate and cohesive module variants. Suppose the overall system is decomposed into  $N$  independent modules where every unit of work  $U_i$  is embedded in a separate cohesive module. This would mean that there is a separate module variant for every variant of every unit of work and that every unit  $U_i$  has  $k_i$  module variants or versions corresponding to the  $k_i$  variants of that unit. The total amount of module versions to develop and maintain - the effort - would then become

$$\sum_{i=1}^N k_i$$

while the total amount of possible system aggregations - the yield - would equal

$$\prod_{i=1}^N k_i$$

This means that there is an exponential relationship between the effort for individual module variants and the total yield of available system variations

$$Nk \longrightarrow k^N$$

In other words, the development and maintenance of various variants of modules leads to an exponential gain in possible system variations. We call this relation the law of exponential variation gains.

#### 3.2.3.2. Coupling: Law of exponential ripple costs

We can distinguish three types of coupling.

We can have coupling between system modules. Suppose we have  $N$  different modules  $M_i$  which each contain a single unit of work  $U_i$  where every module  $M_i$  has  $k_i$  different variants and for each module  $M_i$  there are  $d_i$  other modules  $M_j$  that depend on it. Consider the case where we want to change a specific variant of a specific module. In case the module boundary or interface is not correctly encapsulated or some other dependency is not correctly shielded, it may be possible that these changes will ripple through to the related modules. The number of changes that might be needed would then equal

$$\sum_{j=1}^{d_i} k_j$$

where  $k_j$  is the number of variants of a dependent module  $M_j$ . It is even possible that changes will continue rippling to the dependent modules of the dependent modules, and when no guarantee would exist that the ripple effect would stop there it could lead to a series of changes

$$dk + d^2k + d^3k + \dots$$

We call this exponential increase in efforts or costs the law of exponential ripple costs.

We can also have coupling between various systems. Suppose a module  $M_i$  is reused in another system as module  $M_x$  with version  $v_x$  and  $d_x$  dependent modules. Within this system the number of first order dependencies would then be

$$\sum_{j=1}^{d_x} k_j$$

But as mentioned before we could also have second, third, ... order dependencies within this system and many more in still more other systems. It is obvious that this calls for standardization of basic modules in various domains.

A final sort of coupling we can distinguish is coupling within system modules. Suppose a part of a module  $M_i$  is duplicated or copied in the  $k_i$  versions of this module. Changing this part would then imply  $k_i$  changes and, as this part is combined with the rest of module within that module, it is quite likely this change may ripple through the rest of the module leading to  $k_i$  identical changes and  $k_i$  different changes. This means that a common part in all versions of a specific module will lead to

$$2k_i$$

changes from which  $k_i + 1$  are different. In case a unit of work is duplicated across  $C$  different modules, the number of changes could be even higher.

### 3.2.4. Elements and layers

Implementing the theorems as specified in section 3.2.1 yields an architecture in which 5 types of elements are implemented across 6 layers. For sake of clarity we provide an overview of the data elements in Table 3.2.4 and an overview of the layers in Table 3.2.4.

## 3.3. Refactoring: fixing what's broken

Refactoring is the art of improving the quality of existing code without changing the application's appearance to the outside world. The goal of refactoring is to make code easier to understand, maintain, and modify, while reducing the risk of introducing new bugs and tackling technical debt.

The need for refactoring can arise due to various reasons. For instance, the code may have become difficult to understand or modify due to its complexity or lack of proper documentation. It may also have become outdated or redundant due to changes in business requirements or technological advancements.

Some common refactoring techniques include identifying and isolating code smells, applying design patterns and principles, simplifying code and applying separation of concerns, and using tests to ensure correctness and consistency. Refactoring should be done incrementally and iteratively, focusing on small, manageable chunks of code that can be improved without causing major disruptions or breaking

Element	Description
Data	A Data Element is at the core of NS systems and represents information used in the application. Data Elements are often related to nouns used to describe the workings of your application.
Task	A Task Element corresponds to a 'verb'. It operates on a Data Element and will execute some actions. A Task will take the element it is working on and return a task result that can either be success or failure. It is important that tasks are implemented independent of the state of the Data Element it is working on.
Flow	A Flow Element describes a process in the application. It allows a user to define a number of states and transitions which describe the flow of this process.
Connector	Connector Elements provide I/O support and are provided as-is.
Trigger	The Trigger Element provides the functionality of triggering or periodically invoking the flow orchestrators and correspond to a system clock. It is almost completely confined to the Logic Layer.

Table 3.1.: Overview of the Elements in NST

Layer	Description
Data	The Data Layer contains all data and CRUD classes to provide data access.
Control	The Control Layer contains the server-side logic, it receives requests from the View Layer and uses both the Shared and Proxy Layer to handle them.
Logic	Contains Bean classes which provide transactions and implementations of tasks and commands.
Shared	The Shared Layer is used by all other layers (except the View Layer) and contains classes which are mainly used by the other layers to communicate.
Proxy	The Proxy Layer provides encapsulation for the remoting technology. Artifacts in the Logic Layer use EJB to implement transactions and scaling, but they can only be accessed by retrieving the objects through JNDI lookups.
View	The View Layer contains the visualization of the application and defines the pages the user will see and interact with.

Table 3.2.: Overview of the Layers in NST

functionality.

Refactoring can be a challenging and time-consuming process, particularly for large and complex codebases. However, it should not be seen as a one-time task or a last-minute solution to fix code problems. Instead, it should be an integral part of the development process, focusing on continuous improvement, collaboration, and code quality.

Refactoring should not be confused with reengineering. Where refactoring is focused on improving the quality of existing code, the aim of reengineering is to alter the fundamental structure of a system or application. Refactoring can thus be used as a tool in reengineering efforts.

### 3.4. SonarQube

SonarQube is a tool that provides a detailed overview of numerous metrics when run on a codebase. In this thesis we will use SonarQube to obtain results on following metrics:

- Complexity: this value denotes the cyclomatic complexity of the codebase. A higher value indicates more branches in the codebase.
- Duplications: number of duplicated blocks of lines, files and percentage of duplicated lines. The expectancy is that these numbers will be relatively high when run on codebases following NST due to the way  $\mu$ Radian generates its code.
- Code smells: the total count of code smell issues found in the codebase.
- Technical debt: provides an estimate for the effort to fix all code smells.
- Maintainability rating: rating given to the project relative to the value of the technical debt ratio. This takes into account the outstanding remediation cost, i.e. the sum of the estimated time to fix all open issues. If this time is
  - $\leq 5\%$  of the time that has already gone into the application, the rating is A
  - $6 \leq time \leq 10$  the rating is B
  - $11 \leq time \leq 20$  the rating is C
  - $21 \leq time \leq 50$  the rating is D
  - anything over 50% is an E
- Vulnerabilities: the number of vulnerability issues.
- Security rating:
  - A = 0 Vulnerabilities
  - B = at least 1 Minor Vulnerability
  - C = at least 1 Major Vulnerability
  - D = at least 1 Critical Vulnerability
  - E = at least 1 Blocker Vulnerability

### 3.4. SONARQUBE

- Security remediation effort: the effort to fix all vulnerability issues.

#### 3.4.1. Other tools

Another tool that might result in useful results would be CodeScene, which is a tool based on dynamic analysis unlike SonarQube which only performs static code analysis. When used in a professional setting, CodeScene allows for developers to track how much time has been spent on certain parts of the code, so called hotspots, as well as define certain goals on problematic areas.

When run on both the old and NST version, on first glance there didn't seem to be any significant results. For example, the code health shown in Figure3.4 reveals no issues with code health while the NST version is deemed to have some issues as shown in Figure3.5.

If we dive deeper into the different hotspots of our NST application, it turns out the the three hotspots

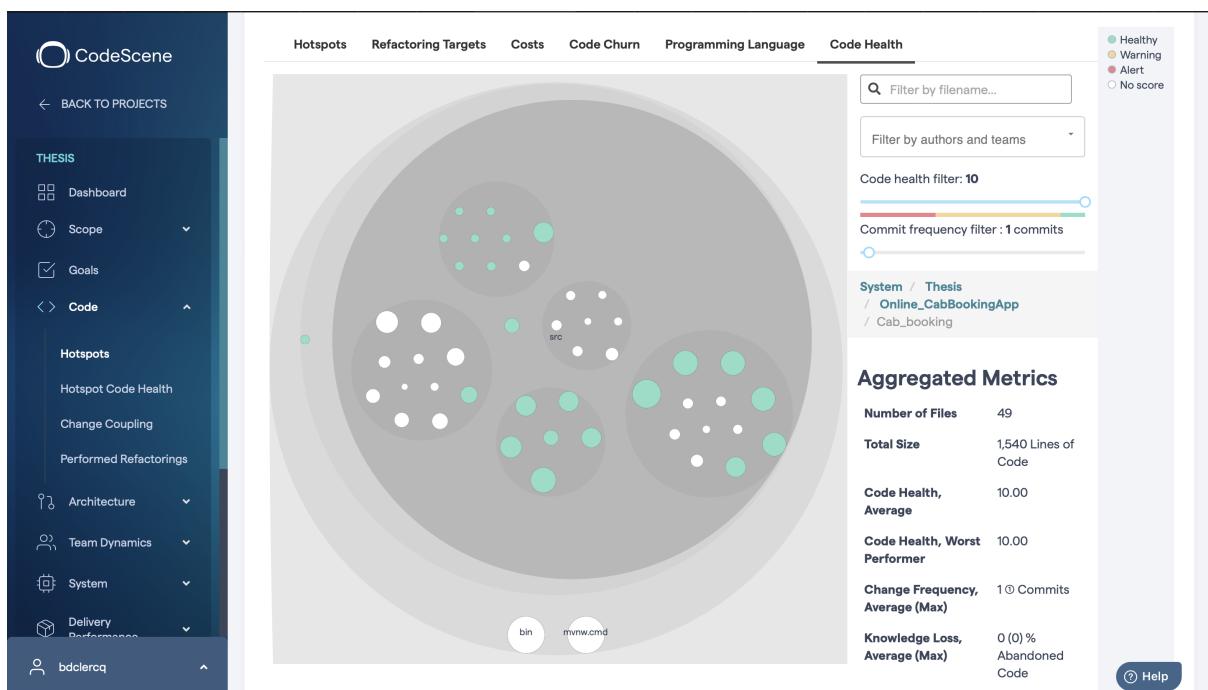


Figure 3.4.: Code health of original system according to Codescene

which contain red zones are actually part of JavaScript libraries and therefore fall outside the scope of NST. When we shift our focus to the part we generate using  $\mu$ Radiant we see that, while most of the code seems to be healthy, we do have some slight issues as revealed in Figure3.6. These issues in components/onlineCabBookingComp/ are as follows

- data layer: TripBookingCruds.java
    - Low cohesion in the module as measured with the LCOM4 metric (Lack of Cohesion Measure). With LCOM4, the functions inside a module are related if a) they access the same data members, or b) they call each other. High Cohesion is desirable as it means the Separation of Concerns theorem is followed as it should be. CodeScene has found that this module has at least 4 different responsibilities amongst its 47 functions.
- If we take a look at the X-ray of the module however, CodeScene tells us “There's no significant structural recommendations to make. Keep up the good work!”.

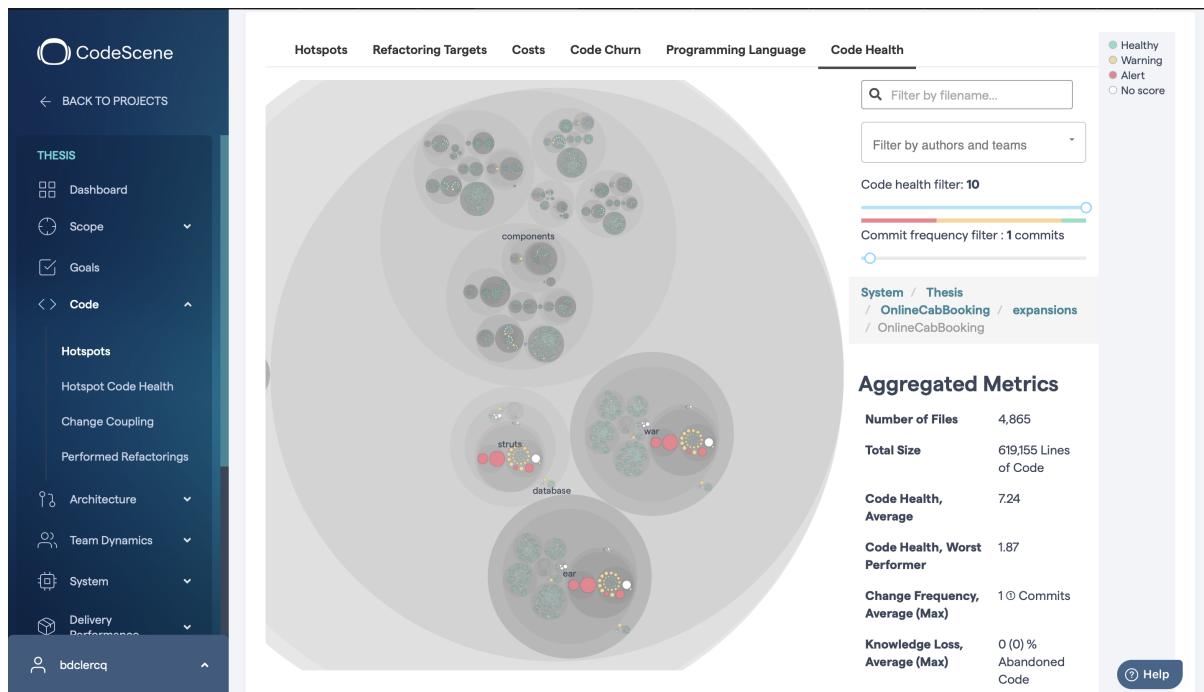


Figure 3.5.: Code health of NST system according to Codescene

- The create method has a cyclomatic complexity of 11, which is deemed to high.
- Duplicate code blocks have been detected. Seeing as these code blocks are generated by the expanders, without requiring any work from the developer, these duplicate code block issues can be ignored.
- shared layer
  - ComponentMetaData.java
    - \* Low cohesion in the module as measured with the LCOM4 metric (Lack of Cohesion Measure). With LCOM4, the functions inside a module are related if a) they access the same data members, or b) they call each other. High Cohesion is desirable as it means the Separation of Concerns theorem is followed as it should be. CodeScene has found that this module has at least 4 different responsibilities amongst its 13 functions.
    - \* Once again duplicate code blocks have been detected.
    - \* There are 4 methods which are too large and should be split up, these methods are getDataElementDef\_TripBooking(LoC = 89 lines), getDataElementDef\_Driver(LoC = 75 lines), getDataElementDef\_Person(LoC = 72 lines), getDataElementDef\_Address(LoC = 70 lines).
  - TripBookingMapper.java
    - \* This file has 2 bumpy roads, which are functions that contain multiple chunks of nested conditional logic inside the same body. The deeper the nesting and the more bumps, the lower the code health. CodeScene considers the following rules for the code health impact: 1) The deeper the nested conditional logic of each bump, the higher the tax on our working memory. 2) The more bumps inside a function, the more expensive it is to refactor as each bump represents a missing abstraction. 3) The larger each bump – that

### 3.5. SYSTEM UNDER OBSERVATION

is, the more lines of code it spans – the harder it is to build up a mental model of the function. The nesting depth for what is considered a bump is 2 levels of conditionals. The reported bumpy roads are convertToMap and convertToDetails.

- \* Complex methods: the same methods reported in the item above have a cyclomatic complexity of 19 and 20 respectively and should be simplified. As in other components these methods are also reported to have a too high cyclomatic complexity, it might be something to consider changing in the expanders.
- \* The convertToDetails method is reported to be too long and should be split up into smaller routines. As this is also a recurring remark in other components it might be beneficial to try and tackle this issue as well.
- view layer: tripBooking-finders.js, person-finders.js, driver-finders.js. In all these files the different finders which are generated are reported as being duplicate code blocks.

The other features of CodeScene, such as Costs and Code Churn, are currently useless as they depend on previous commits to assess changes.

If we thus take into account the limited amount of information we can get from this tool given the current systems, we decided not to use this tool as a basis for comparison.

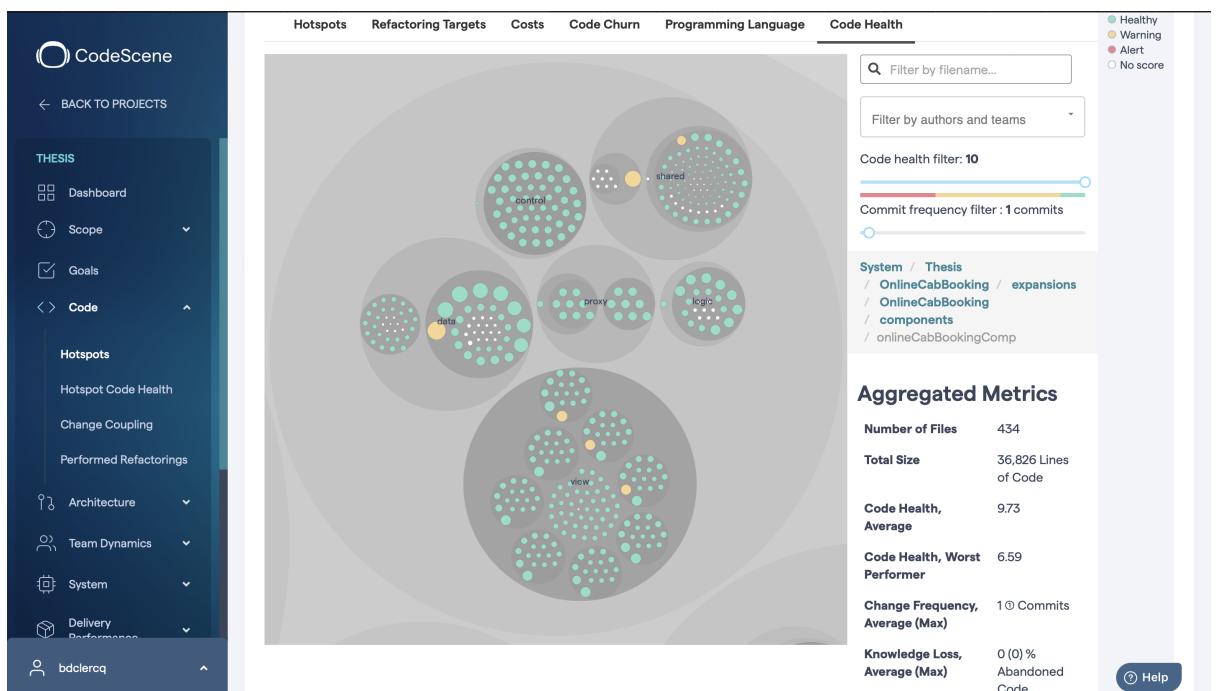


Figure 3.6.: Code health of the custom component according to Codescene

## 3.5. System under observation

In order to verify whether NST delivers on the promises it makes we decided to perform a small experiment.

We selected a small legacy system with as criteria that it should be written in Java and it should be

a CRUD based system, basically eliminating games for example and music processing software. The system we chose was found on GitHub at [https://github.com/VaibhavTyagi010/Online\\_CabBookingApp](https://github.com/VaibhavTyagi010/Online_CabBookingApp).

This Cab booking application was created using a specific kind of **multi-tier** architecture, namely Controller-Service-Repository. In this implementation the developers decided to extend this pattern with two additional layers, resulting in a total of five layers: Controller, Entity, Exception, Repository and Service.

### 3.5.1. Controller layer

Figure 3.7 provides a diagram of how the Control layer is implemented. As we can see the classes in the Control layer provide the CRUD, Create-Read-Update-Delete, operations needed for the application as well as the necessary getters and setters to allow access to the private members.

Classes in this layer define user-interface according to the REST principles. They relay requests they receive to the corresponding Service layer object.

### 3.5.2. Entity layer

Classes in the Entity layer define how data is stored in the database. Figure 3.8 provides an overview of the entities which are currently present in the application and how they are connected.

It is worth noting that in the current implementation, the use of inheritance breaks the Separation of Concerns principle. This can be circumvented by creating both the superclass and subclass as separate data elements and making the superclass entity a datamember of the subclass entity.

If we apply this to the current system, we would get data elements for Abstractuser, Customer, Admin and Driver where Customer, Admin and Driver would contain a pointer to an Abstractuser.

### 3.5.3. Exception layer

In Figure 3.9 we can see that some time has been put towards custom error handling.

### 3.5.4. Repository layer

The DAOs in this layer, as given in Figure 3.10, provide an API through which the application can access the database layer.

If we transform this system into NST, these finders and views can be generated simply by adding them to the corresponding data element in  $\mu$ Radiant.

### 3.5.5. Service layer

The service layer calls the API defined by the Repository layer to allow users to perform CRUD operations in the application. Figure 3.11 provides an overview of the services which are currently provided

### 3.5. SYSTEM UNDER OBSERVATION

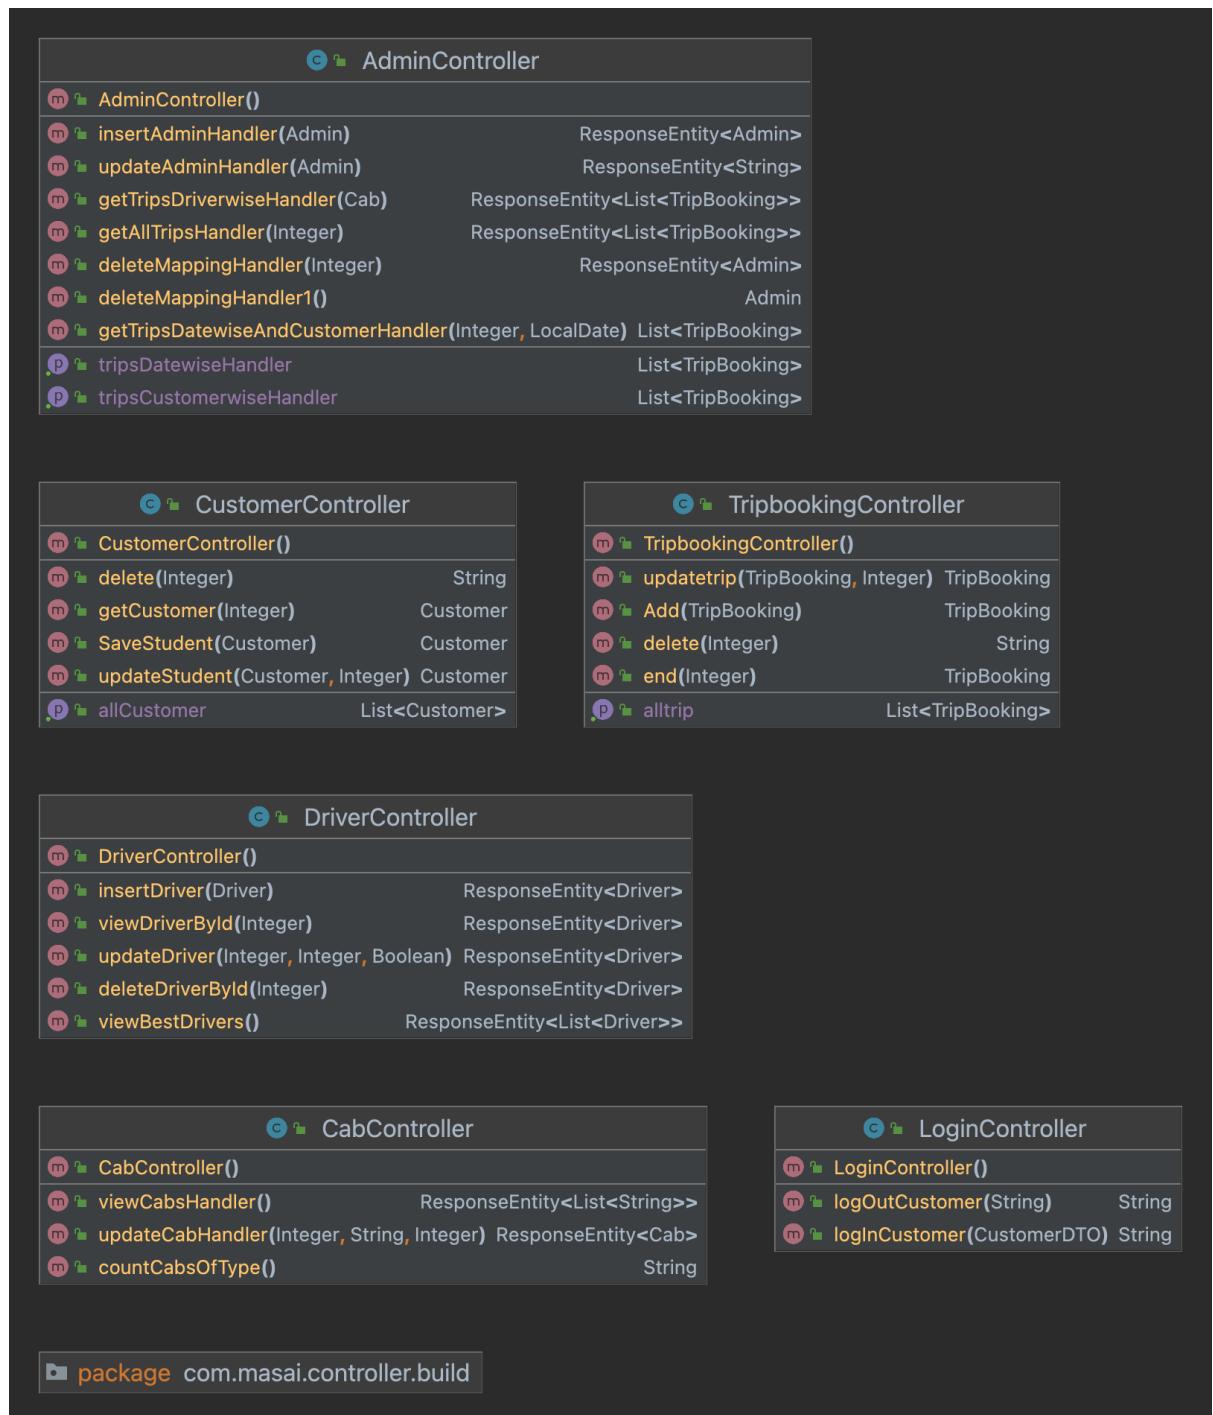


Figure 3.7.: Diagram of the control layer

to the user.

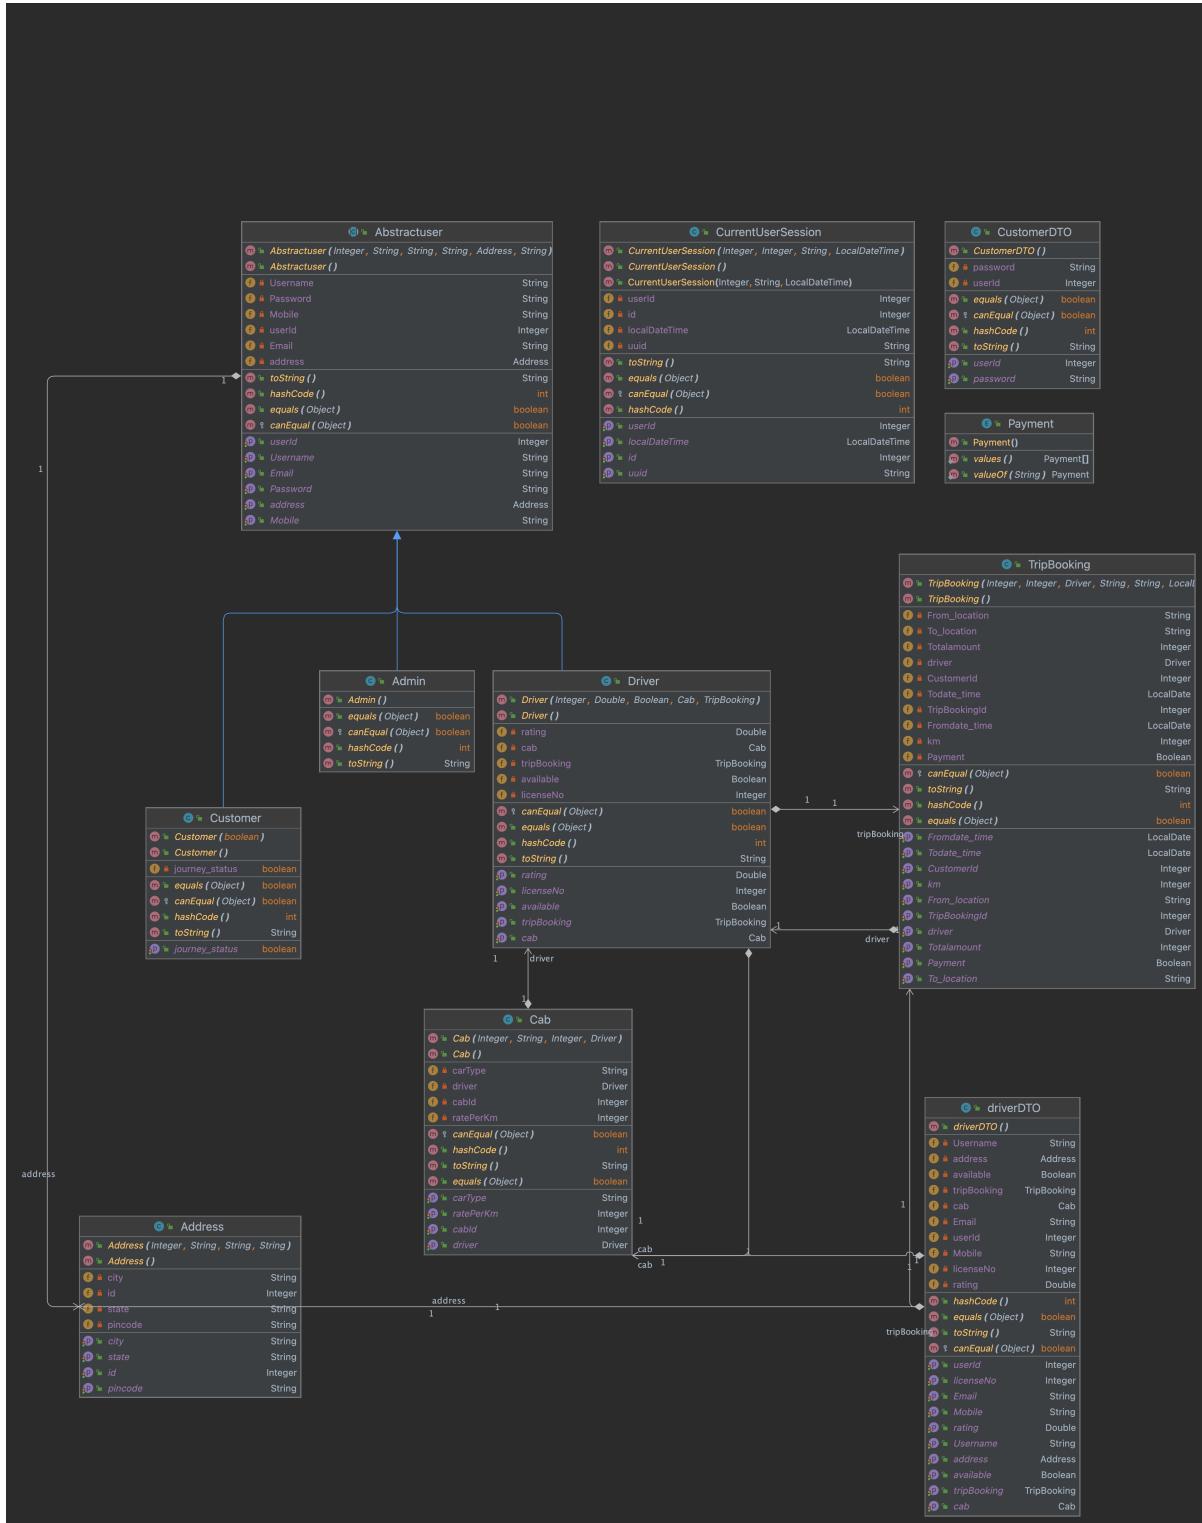


Figure 3.8.: Diagram of the entity layer

### 3.5. SYSTEM UNDER OBSERVATION

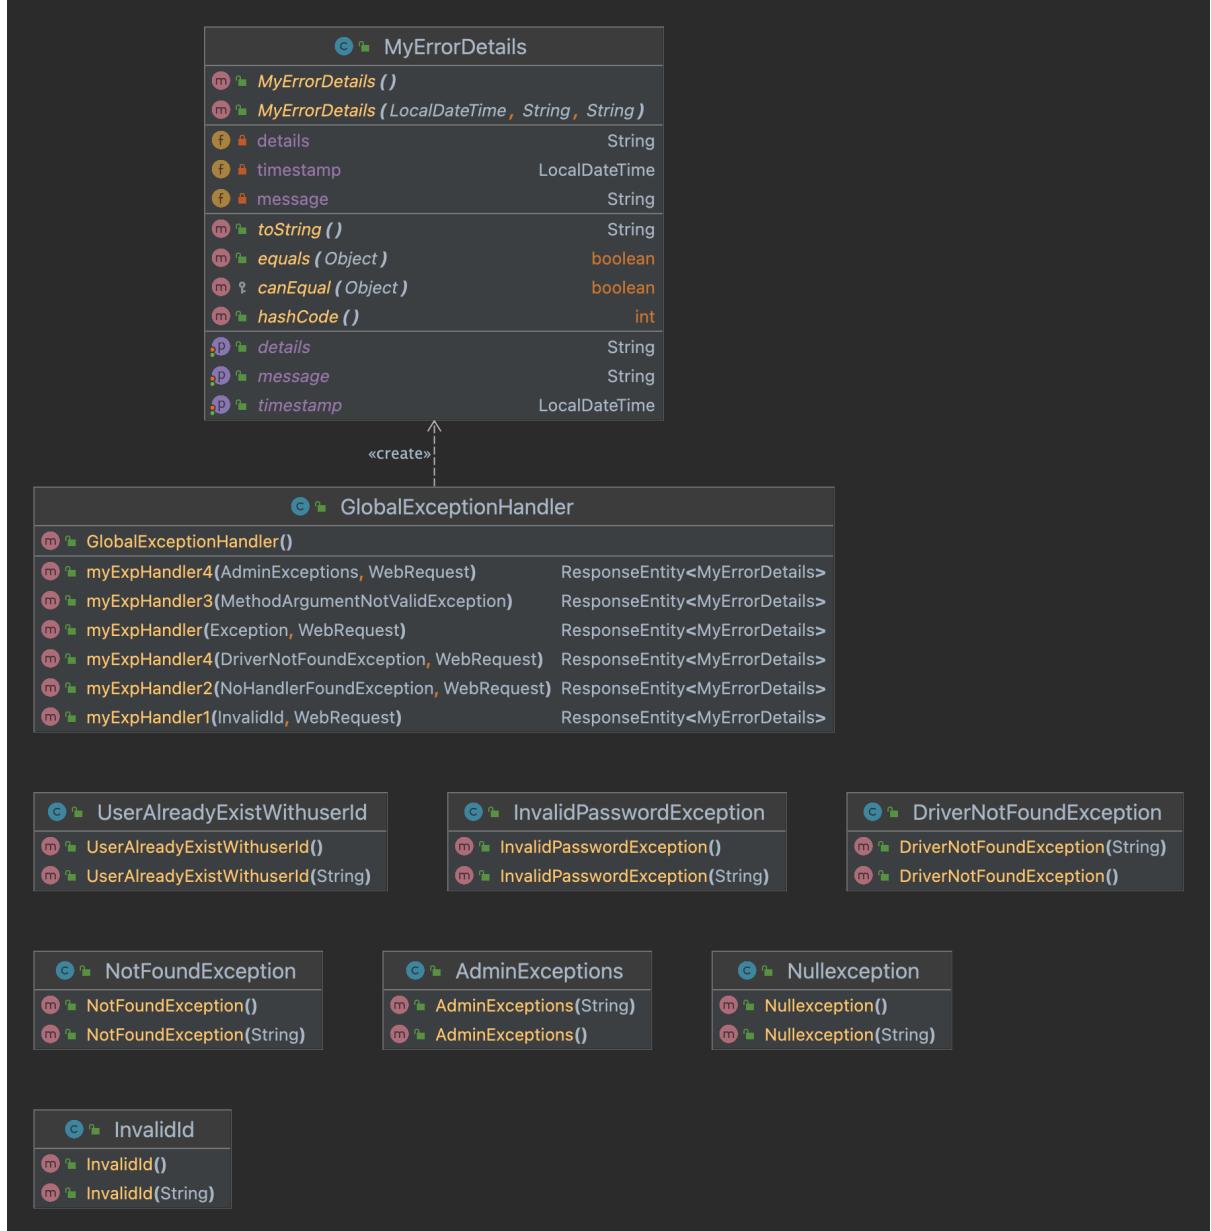


Figure 3.9.: Diagram of the exception layer

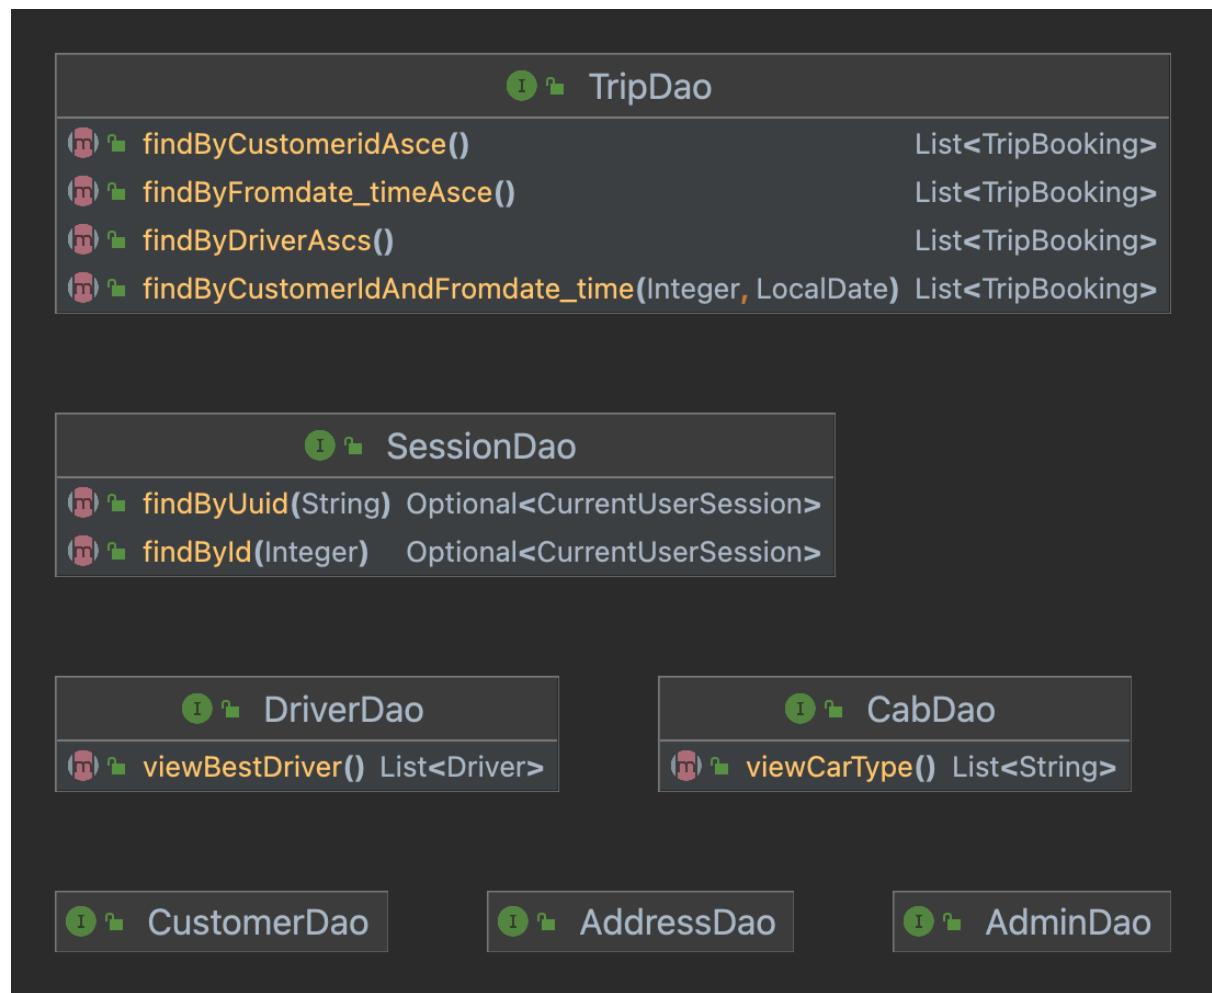


Figure 3.10.: Diagram of the repository layer

### 3.5. SYSTEM UNDER OBSERVATION

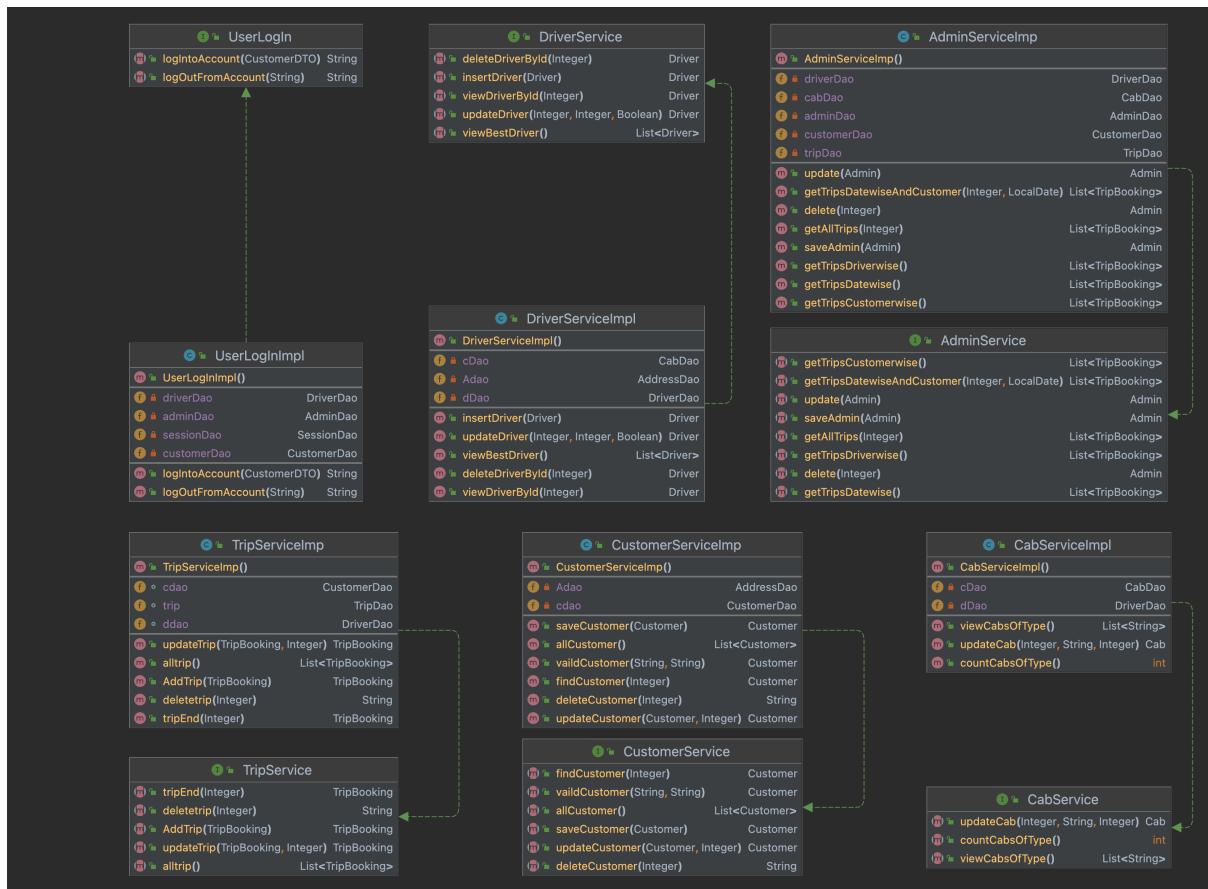


Figure 3.11.: Diagram of the service layer



## 4. Refactoring into NST

### 4.1. $\mu$ Radiant

$\mu$ Radiant is a tool which enables developers to model their NS application. As shown in Figure 4.1, an NS Application has the following structure

- The application is represented by an Application element.
- Each Application can have multiple ApplicationInstances. An ApplicationInstance is a combination of settings to expand the application with.
- Each Application can have multiple Components, and Components can be reused among multiple applications. Components group Data, Task, Flow Elements and ValueFieldTypes.
  - DataElements represent information in the system, which can be stored, queried and updated.
  - TaskElements represent executable pieces of logic. A task has a target DataElement. When executed, the task receives an instance of that DataElement as parameter.
  - Flow Elements are linked to a DataElement with a dedicated status field. The Flow Element describes a number of state transitions. Each transition has a begin-state and a TaskElement to execute if an instance of the DataElement reaches that state.
  - ValueFieldTypes allow the implementation of custom types for ValueFields in the DataElements of the Component.

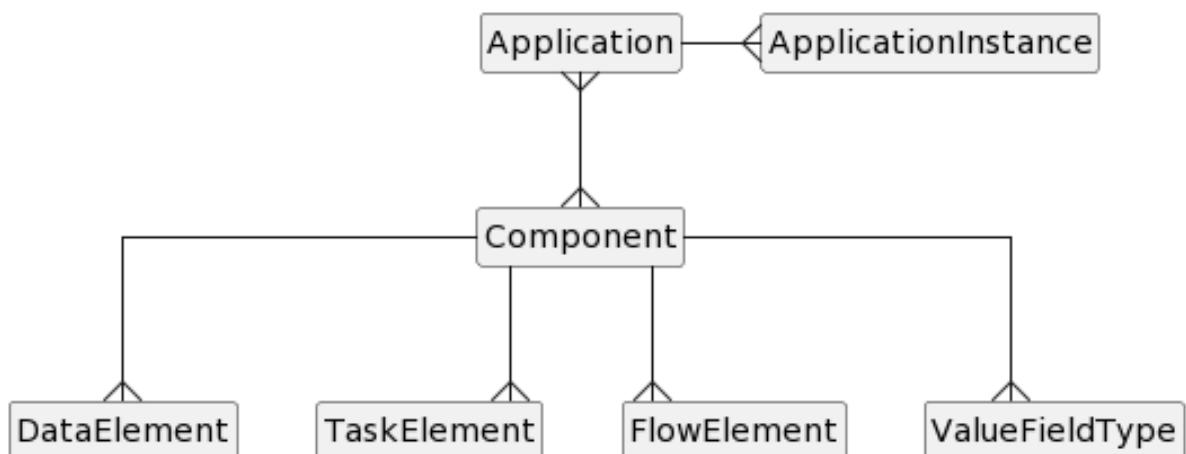


Figure 4.1.: Abstract structure of NS applications

## 4.2. The transformation process

In order to start our transformation process, we first need to collect a list of all data elements we will need in our application. As our goal is to recreate the system as closely as possible we decided to only take into consideration the data entities which were already present without adding new ones unless absolutely necessary.

### 4.2.1. Installation

In order to replicate what we will discuss further in this chapter, readers need to obtain a license from <https://foundation.stars-end.net> and proceed with the installation as described on <https://foundation.stars-end.net/docs/tools/micro-radiant/installation>.

After successful installation,  $\mu$ Radiant can be started by executing following script

```
1 #!/bin/bash
2 cd <path/to/installation/location>/micro-radiant-1/
3 ./start.sh
4 open http://localhost:9050/models
```

This will open the editor in the users default browser and will look similar to Figure 4.2. Before we

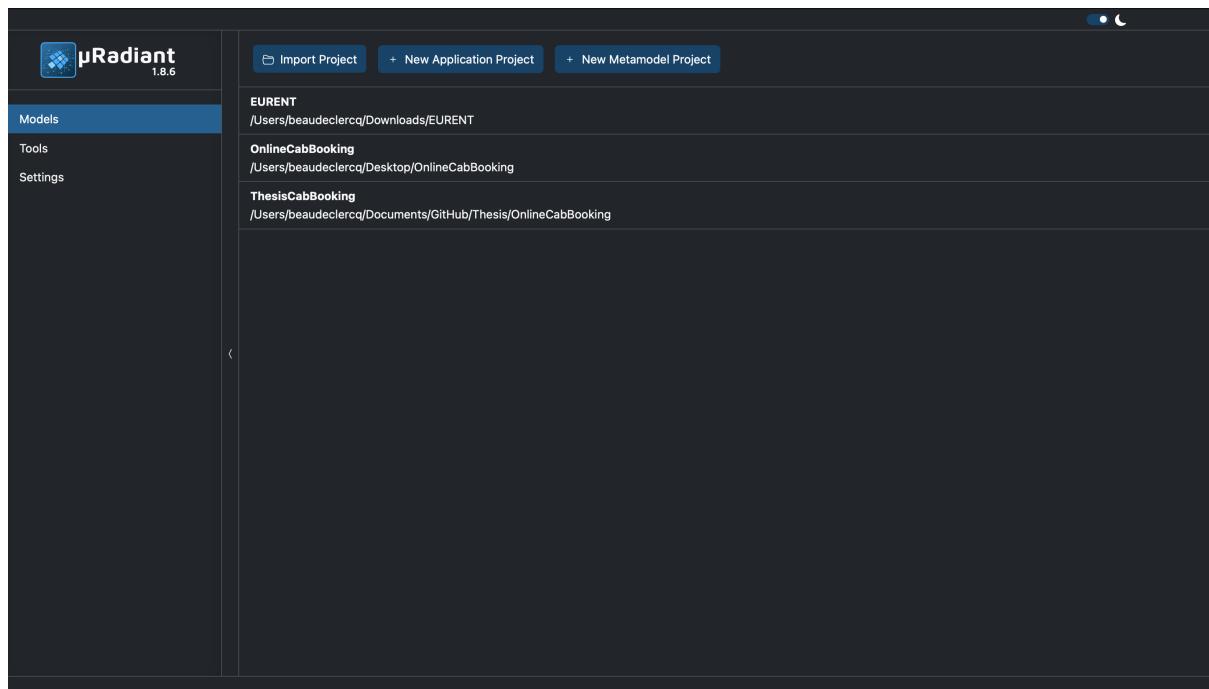


Figure 4.2.:  $\mu$ Radiant homepage

continue with creating the project, we first need to determine which data elements we will need. After all, good preparation is half of success.

### 4.2.2. The Data Elements

After going through the different layers of the original application we decided to keep the following list of data elements

- Address: this element represents the addresses which will be used in the application. An address consists of a state, city, zipcode, street and house number.
- Cab: in accordance to the initial application, a cab has a rate per km, a reference to its driver and a reference to its car type. As the viewCarType() method from the CabDao simply returns a list of all unique car types currently available as cab, and we opted to create a separate data element for car types to better adhere to the separation of concerns principle, this method can be replaced by using a generic or custom finder in the CarType data element.
- CarType: currently a car type is nothing more than a string. The reason we decided to create a separate data element for this is that even though currently it only consists of a name, in the future we might feel the need to add more generic information to cabs of a certain type such as a predefined rate per km, a maximum number of occupants etc. By already separating in this stage we create the possibility for easy alteration later on.
- Customer: as in the initial system, a customer has a journey status and a link to an AbstractUser or Person as we renamed it.
- Driver: unlike a customer, a driver has a fair number of attributes. In addition to a link to a Person, a driver also has a license number, a rating, a link to the cab which it owns, its current booking and a flag to signify whether or not the driver is currently available.
- Payment: we decided to take the price of the trip booking out of the TripBooking element and put it in the more logical Payment element. Along with this total fare the Payment element also has a flag to signify whether or not it has been payed.
- Person: the Person data element represents the AbstractUser from the initial system. A person has a username, password, email address, mobile number and a physical address.
- TripBooking: a trip booking represents the core of the application. A booking entails a customer who has requested a driver to transport him between two locations at two specified pickup times. Based on the total distance in km and the rate per km from the driver, a payment will be created for the customer to pay.

### 4.2.3. Creating the project

Now that the preliminary work is completed, it is time to start the fun part: creating the application.

We do this by selecting the “New Application Project” from the home screen as shown in Figure4.2, which will lead us to the project creation screen as depicted in Figure4.3. The first four fields are fairly straightforward. A user selects the folder where the project needs to be saved, provide a name for the application itself and optionally a name for the applicationInstance as well as a name for the first component.

The groupId is used in the POM.xml file which will be generated and should correspond to that of your company. The name for the dockerImage is once again self explanatory.

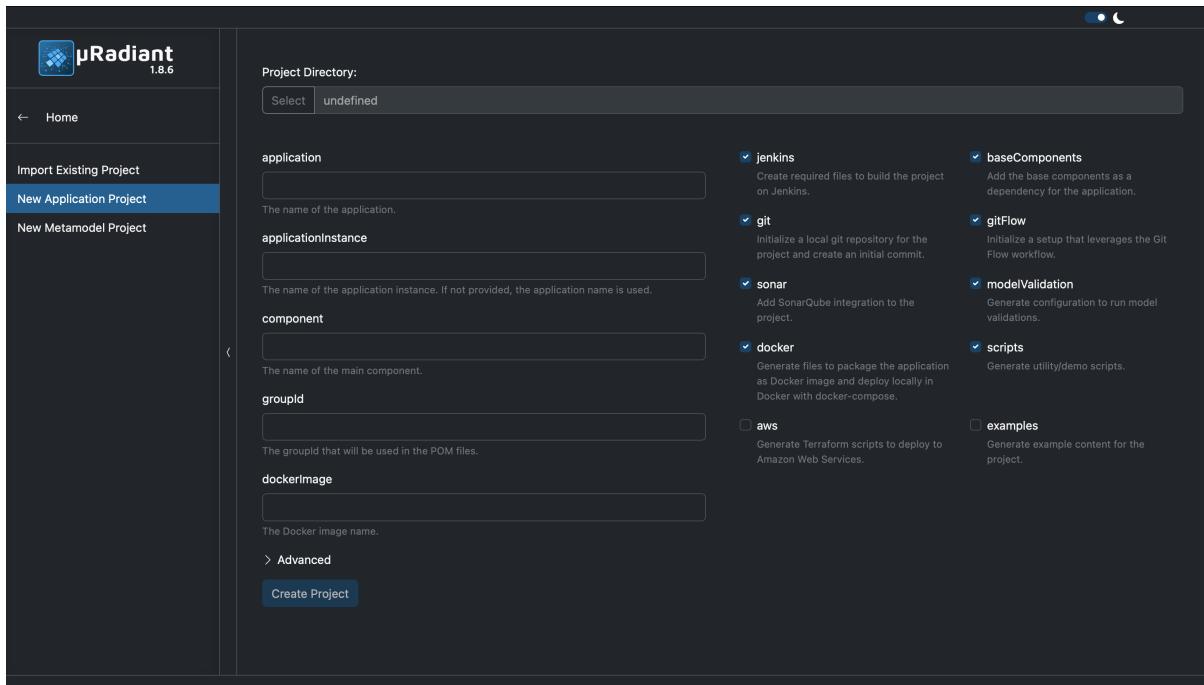


Figure 4.3.: Creating a new project

On the right hand side a number of boxes are checked by default and it is advised to keep them checked as they make life easier down the road. If wanted for demo purposes it might be usefull to check the ‘examples’ checkbox to generate dummy data.

Once everything is filled out correctly, simply click ‘Create Project’ and you’re on your almost ready to do business. If you want to you can go ahead and already take a look at your project by selecting it from the μRadianit homescreen. The project home page will look similar to Figure4.4.

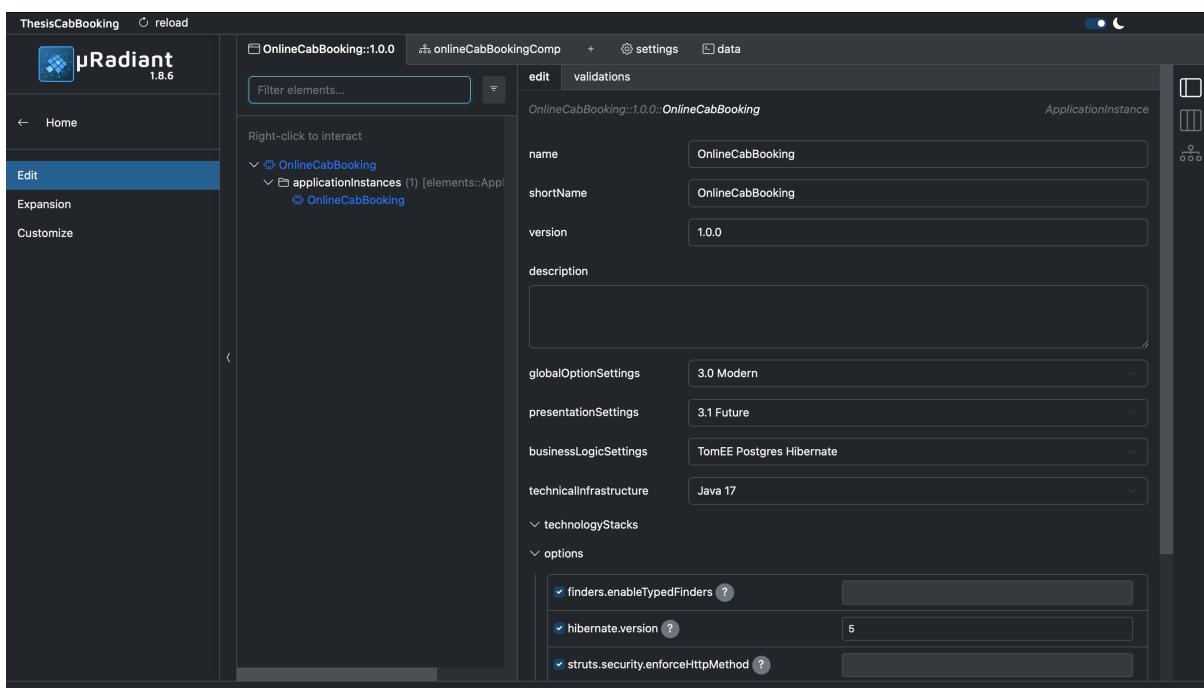


Figure 4.4.: Overview of an NS application

#### 4.2.4. Adding Data Elements

Once the project has been created, we can start adding the data elements we gathered during our preliminary work to our component. To do so we navigate to our component from the application overview, it will look like Figure 4.5. Adding a data element is easy: right-click on the component in the finder pane,

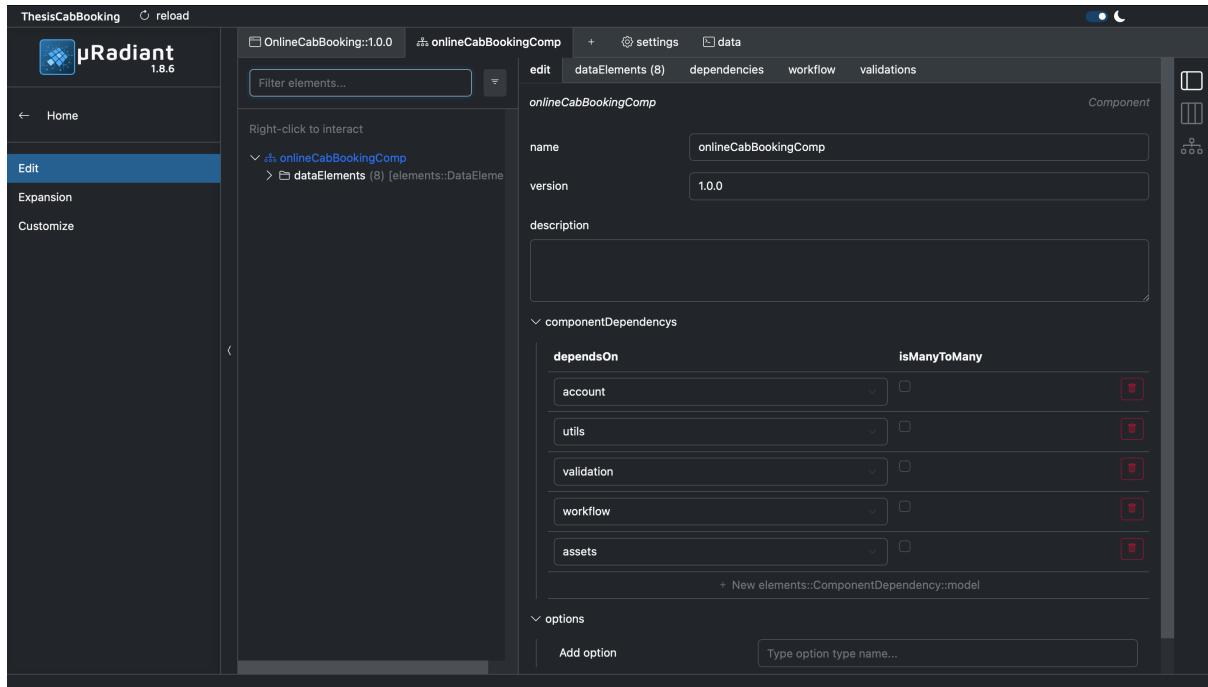


Figure 4.5.: Overview of a component in an application

select “New DataElement” and fill in the required fields as depicted in Figure 4.6.

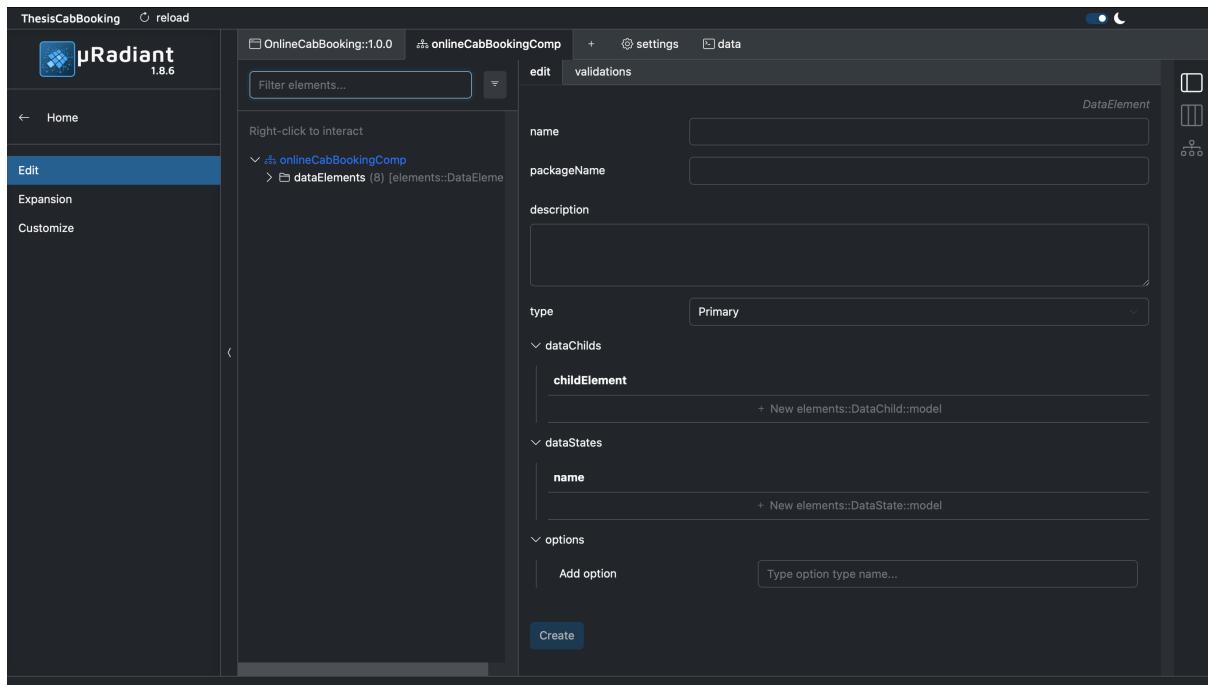


Figure 4.6.: Adding a data element to a component

### 4.2.5. Adding Flow and Task Elements

Of course it makes little sense to have an application where every booking needs to be manually approved and every invoice has to be created by an admin. Thankfully we can automate this by introducing two of the other elements introduced in section 3.2.4: the Flow and Task elements.

In NST, workflows are realized by keeping a status field with each instance. Instances are then moved from state to state by executing the corresponding StateTask and setting the appropriate EndState.

#### 4.2.5.1. Defining the Flow

We decided to start with defining the Flow element for the application. In Figure 4.7 we show the final result of what the element should look like in  $\mu$ Radiant. We suggest filling in the details in following

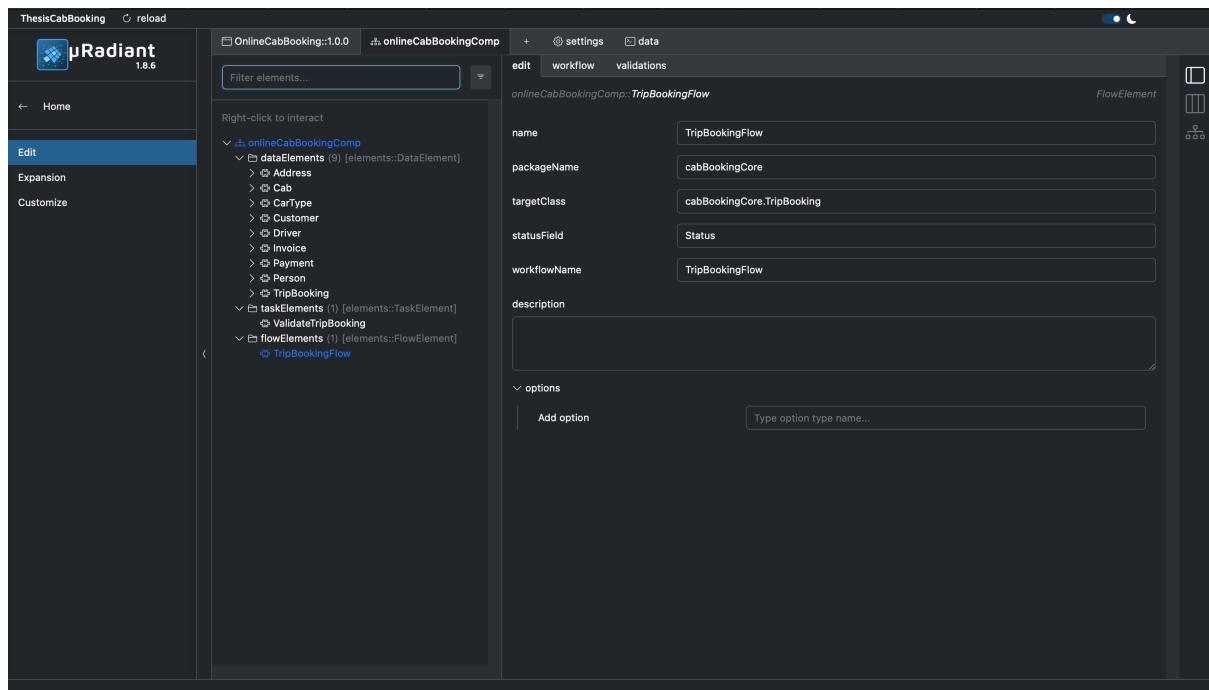


Figure 4.7.: Adding a flow element to a component

order

- **packageName:** this is the name of the package in which the data element on which we want to introduce our flow resides. In our application we only have one package, cabBookingCore, so this one is easy.
- **targetClass:** the class on which you want the Flow to work preceded by the package name, in our case cabBookingCore.TripBooking as we want to add a Flow to automate the booking of our cabs.
- **name:** the name of a flow should always be the name of the data element followed by Flow, which gives us TripBookingFlow.
- **statusField:** name of the field from the data element in which we can store the status, naturally you call this one ‘state’ in your element. In the flow element this is entered with a capital first

## 4.2. THE TRANSFORMATION PROCESS

letter. Don't forget to create a finder for the status field!

- workFlowName: the name you want to give to the actual workflow instance, the safest option is to just take the same thing as in the name field.

Once these fields have been filled in according to what kind of flow you are adding to your application, you simply click 'Create' and the flow will be added as element for you to use.

Each Flow Element must have a corresponding [Target Element]TaskStatus Data Element, an example of which has been provided in Figure4.8.

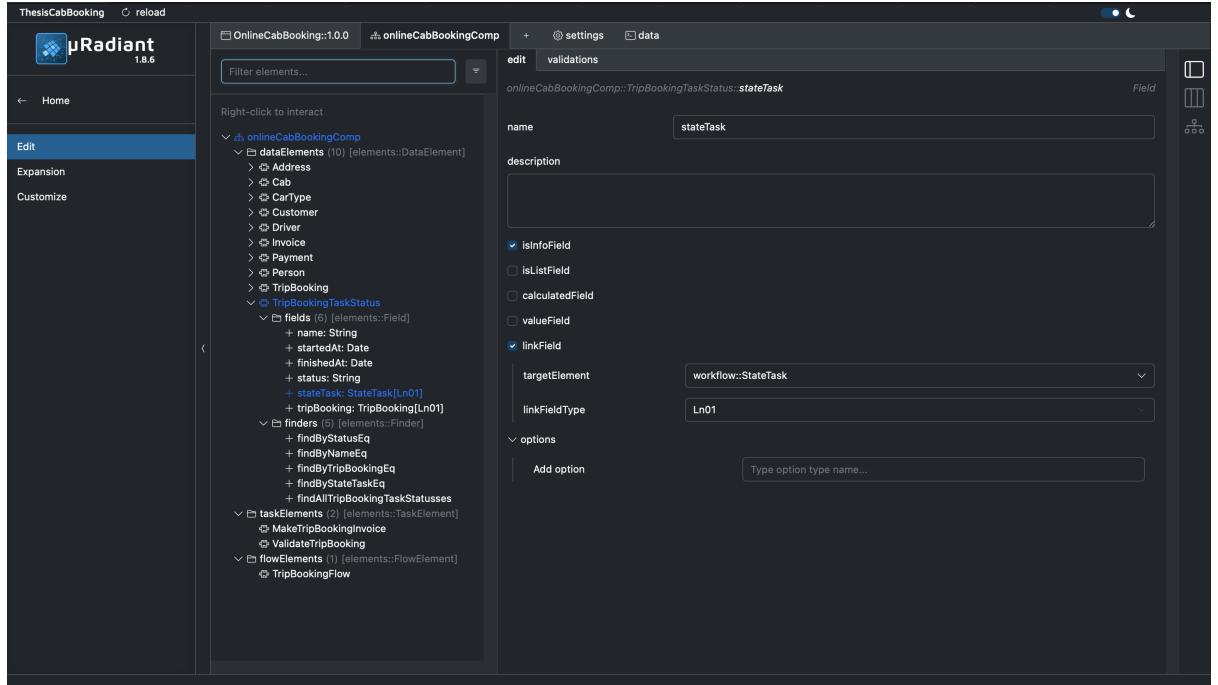


Figure 4.8.: Task status element for TripBooking

Before we continue with adding the Tasks we are going to add the states we want our TripBooking to have. To do so we go to our data element and fill in the dataStates as shown in Figure4.9. We finish by adding a status field to our data element which is of type String.

Once all this is done we can move on to the next step.

### 4.2.5.2. Defining the tasks

We want our flow to perform two different actions: firstly we want to validate our booking, secondly we want to create an invoice for our customer.

Adding a task is fairly easy: right-click on the component and select 'New TaskElement'. A blank version of Figure4.10 and Figure4.11 will appear.

The fields that need to be filled in are as follows

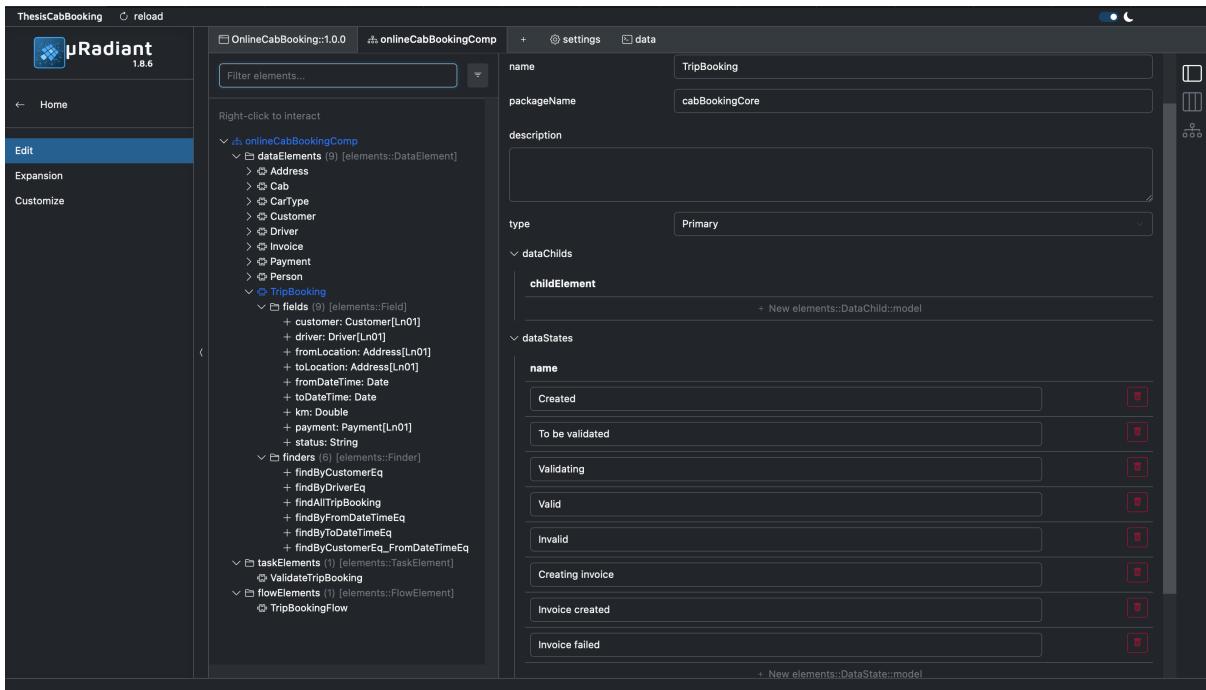


Figure 4.9.: Adding data states to a data element

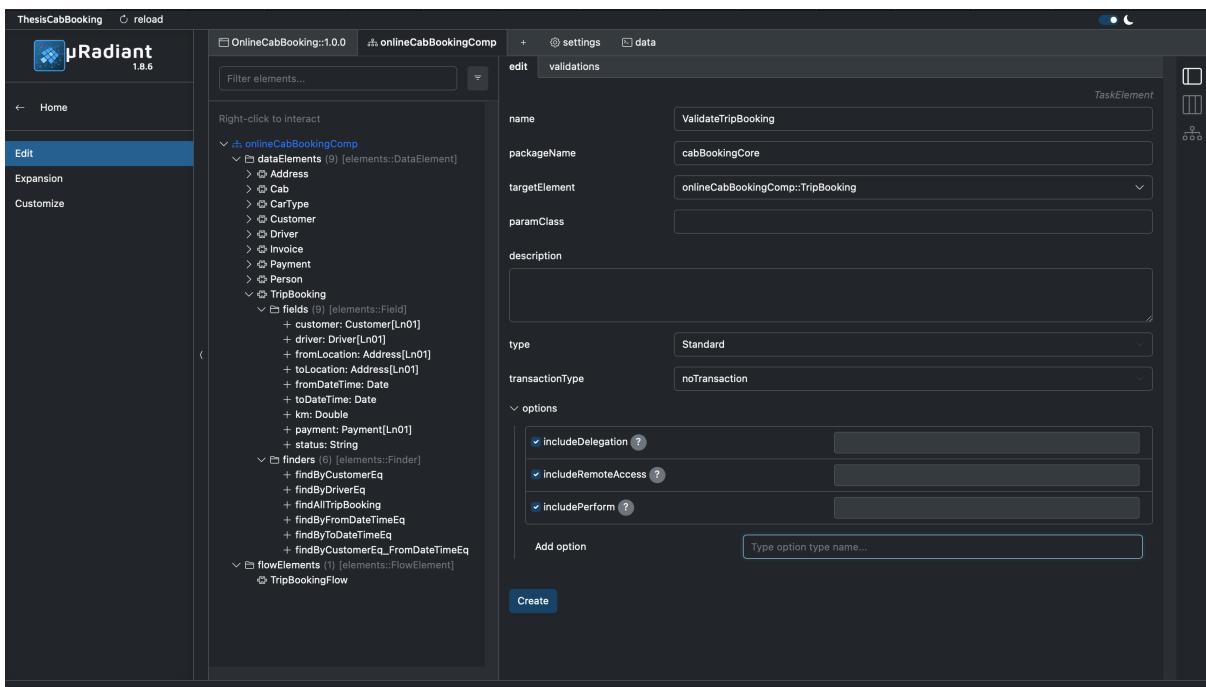


Figure 4.10.: Adding the validate booking task

- **targetElement:** name of the element the task will work on, comes in a dropdown menu.
- **packageName:** package where the task needs to be part of, i.e. the package mentioned in targetElement.
- **name:** name of the task element, e.g. ValidateBooking or MakeBookingInvoice.
- **paramClass:** no need to fill this in.

## 4.2. THE TRANSFORMATION PROCESS

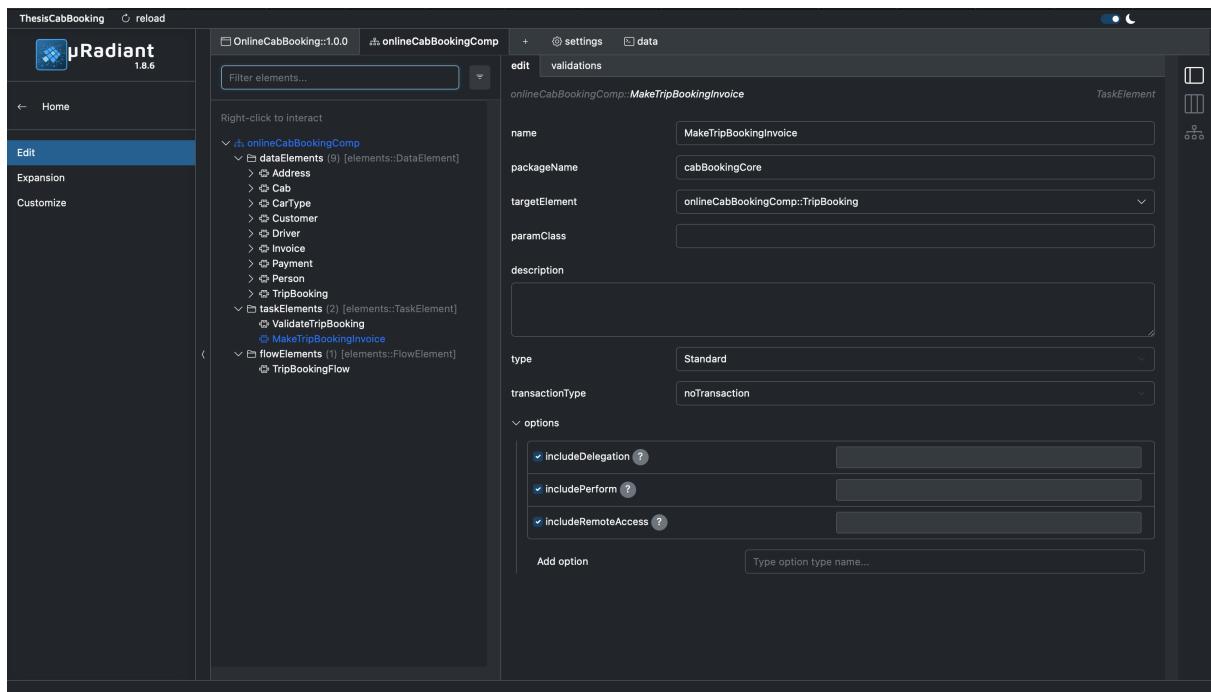


Figure 4.11.: Adding the make invoice task

- **description:** brief description of what the task will do, you can fill this in now or do it later as happens to all documentation.
- **type:** dropdown menu from which you can select how the task should be started, in most cases this will be ‘Standard’. For details on the other options we refer to the online documentation.
- **transactionType:** dropdown menu from which you need to pick the correct option for what you want to achieve.
- **options:** it’s not needed to add options here, we opted to follow given examples by adding includeDelegation, includePerform and includeRemoteAccess.

### 4.2.6. Creating the workflow

Once the necessary Elements are created it is time to build the application and define the workflow. Workflows are realized by keeping a status field with each instance. Instances are then moved from state to state by executing a task and setting the appropriate end state. Such a state transition is defined in a State Task. State Tasks have the following properties

- **name:** Name of task that is to be executed.
- **processor:** A combination of the component and task -name.
- **implementation:** The implementation class used for the task. The expanders generate 2 implementation classes by default with postfixes -Impl and -Impl2.
- **params:** A way to further parameterize the implementation class.
- **beginstate:** Instances that contain this state will be picked up by the workflow to perform this state task.

- **interimstate:** Instances picked up by the workflow receive this state while they are being processed.
- **failedstate:** If the task fails, the instance will receive this state.
- **endstate:** If the task is successful, the instance will receive this state.

#### 4.2.7. On the $\mu$ Radiant Flow

It is worth spending some time on how the cycle of generating an application with  $\mu$ Radiant works. In this cycle can distinguish six phases.

##### 4.2.7.1. Model/Coding phase

As we start a new application we first need to model our application just as we did in the section above. Or, if an application has already been generated, we have been happily coding some custom extras. Once we are happy with the model we created/altered or think our code is ready for testing, we can proceed to the next phase.

A small remark: it is not possible to harvest changes after adding models. For this reason, good practice would be to harvest changes made to the code before making changes with the model editor.

##### 4.2.7.2. Validation

After each remodel  $\mu$ Radiant will automatically validate the model. If any problems arise they can still be fixed and generation of non-working code is avoided.

##### 4.2.7.3. Harvest

If any custom code was written which needs to be kept, don't forget to Harvest it. Harvesting will extract the custom code from between the relevant anchors and store it in the harvest folder located under the components directory. Once we had a successful harvest we can continue to the next phase.

##### 4.2.7.4. Expand

The expansion phase is where the magic happens. During this phase two things happen. Firstly, all the existing code is deleted. Secondly, based on the defined models, settings and saved custom code, a new application is generated. For coding purposes it is advised to open the pom.xml file from the /expansions/{application} folder as a project in an IDE such as Intelij in stead of just opening the root folder as a project.

##### 4.2.7.5. Build

The previous phases can already be executed in an endless cycle, for example if you forgot something or suddenly a more elegant solution came to mind. If however the code seems mature enough to be

run, all we need to do before we can deploy is to build our application. This will trigger all necessary commands and scripts to produce a deployable application.

#### 4.2.7.6. Deploy

Deploying an application requires Docker to start and maintain containers. Once the application has been successfully deployed it is time to start experimenting with the newly added functionalities and layouts. Once the application has been deployed we will inevitably end up in the first phase again and the cycle will start again.

#### 4.2.8. On Anchors and Custom Code

In the code generated by the expanders a lot of anchors are present. These anchors denote places where things happen: where imports are located (on top of the file), where variables are located, the phases in which a method is divided, ...

Most of these anchors are there to provide clarity and while you can alter the code between them, any changes will be removed when the code is regenerated to ensure the four design principles are not violated. But, as it would be strange to not be able to write custom code at all, some anchors clearly indicate that between them you can safely write custom instructions. These anchors are denoted with “//anchor:custom-[something],[start—end]”.

Any custom code written between these anchors is extracted and saved to a safe place during the harvest phase.

#### 4.2.9. On Finders, Projections and Custom Finders

As no application would be complete without some finding functionality,  $\mu$ Radiant is able to provide most of these by default. A expansive list of operators is provided with which you can filter out certain field-values. It is even possible to combine multiple operators and field-values to obtain more complex finders. As can be seen in Figure 4.13 the name is automatically generated by adding field-operator pairs.

If however no combination of operators would satisfy the requirements, there is always the option of selection the ‘isCustom’ option. Doing so will add some small changes to the code which allow for the implementation of this custom finder. Implementing a custom finder happens according to following steps

- Add the finder in  $\mu$ Radiant and declare it as custom finder, otherwise the finder is put among the regular findBy-methods anchors and cannot be changed without manual overriding in the custom methods anchor, which would lead to code duplication.
- Add the necessary fields for your finder to the corresponding `|finder|Details` class along with any additional getters and setters.
- Modify the mapping functions `jsToViewModel` and `ViewModelTojs` found in `|finder|-ko-mapper.js` in the view layer.
- Create an appropriate form `|finder|-ko-form.html` in the view layer.

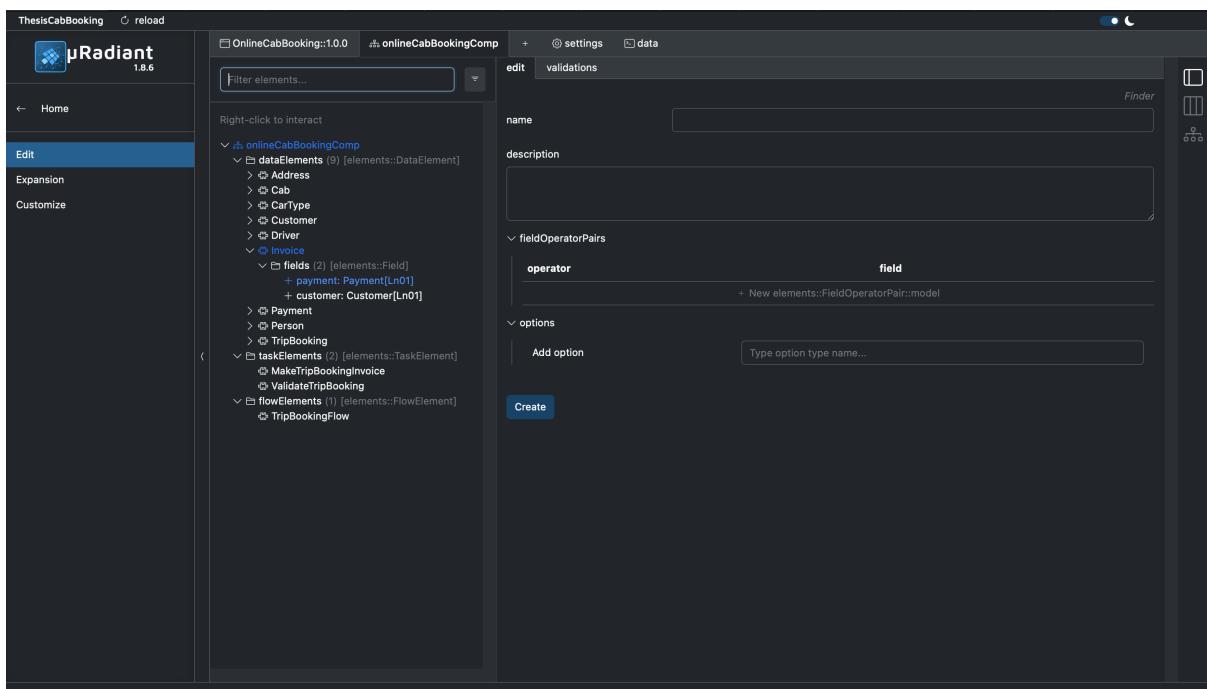


Figure 4.12.: Blank page for adding a new finder to a data element

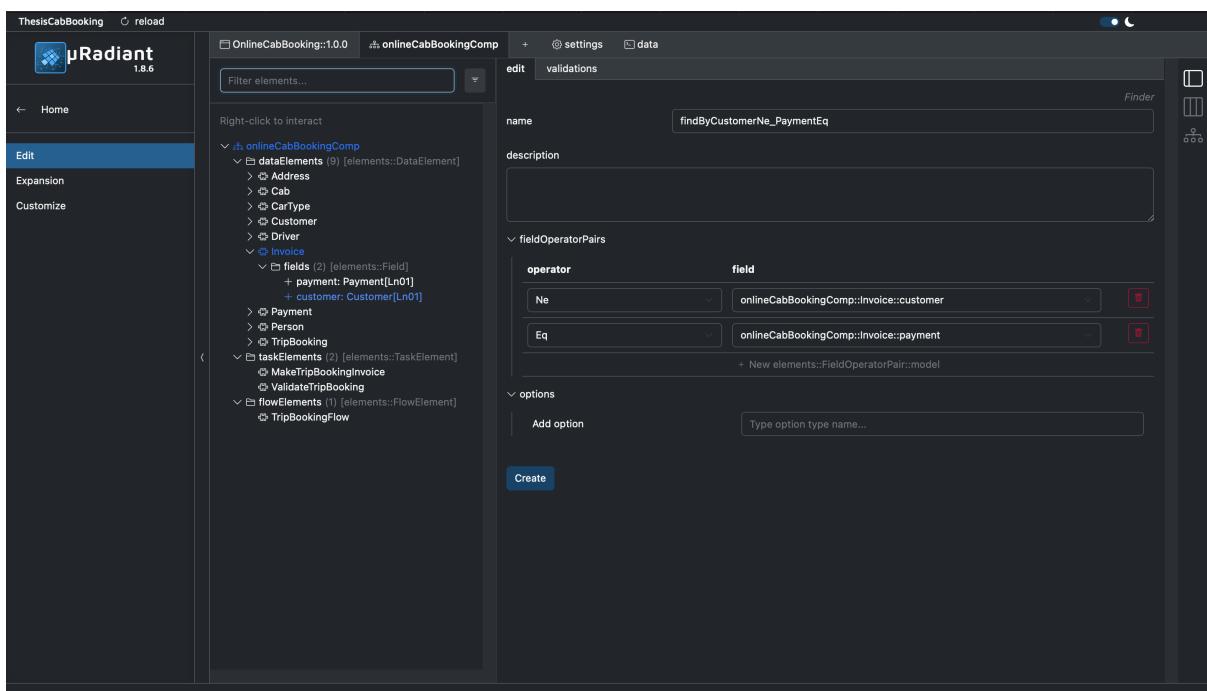


Figure 4.13.: Example finder for the invoice data element

- Add named queries to the `|DataElement|Data` class in the data layer.
- Use these named queries in the `|DataElement|FinderBean` to fetch the data for the finder.

A final item that needs some attention is the presence of data projections. DataProjections define a set of fields from their parent DataElement. When data is passed through the application, it is always represented in the form of a projection. There are 3 projections that are always present for each DataElement.

#### 4.3. FUTURE WORK

- Details Projection: The details projection serves as the default representation for most backend logic and is used to create and modify. The details projection is defined by the set of all fields in the DataElement excluding the calculated Fields.
- Info Projection: The info projection is a reduced set of fields from the DataElement. It defines the columns in the table in the user interface by showing the fields which have been marked as infoField.
- DataRef projection: The DataRef projection is used in several places to reference an instance of a DataElement without requiring the entire DataElement projection. By default, this projection only contains the id (and name if possible) of that instance. It is also possible to define your own Fields to add to the DataRef To do so, add the option functionalKey with the names of the fields separated with .. This functionalKey option generates a DataRef class which contains said Fields in addition to the name and id, so they can be used to look up instances. This new dataRef class will then be used when fetching data from the DataBase. Adding this option forces all functional fields to be contained in each projection.

Developers can also define their own projections. These projections define a set of reference fields, which point to fields of the DataElement.

### 4.3. Future work

- Testing: at time of writing there is no functionality to automatically generate tests. As unit tests and smoke tests can be generated automatically, it might be a useful extension in future development of the expanders.
- Automatic import: currently there is only functionality to import existing NS projects. It might be useful to provide functionality which allows for automatic importing of data elements as most software architectures already provide some separation of data entities, business logic and visual representation.



# 5. Conclusions

## 5.1. Short term findings

Metric	Original system	NST system
Bugs	3	0
Reliability rating	C	A
Reliability remediation effort	35m	0m
Cognitive Complexity	26	2104
Code smells	120	3124
Technical debt	6h26	35d
Debt ratio	1.0%	2.2%
Maintainability rating	A	A
Vulnerabilities	8	0
Security rating	D	A
Security remediation effort	1h20	0

### 5.1.1. Explanations

#### 5.1.1.1. Bugs

According to the SonarQube scan, the original version of the project contains 3 bugs which they classify as major. These bugs are

- In exception/GlobalExceptionHandler, the getFieldError() method could throw a NullPointerException and currently this behavior is not anticipated. Remedying this bug is estimated to take about 10 minutes.
- In service/UserLogInImpl 2 bugs are present, both in the logIntoAccount method. The first one refers to an issue on line 46 where the ID is being retrieved from an object obtained by using the get() method from the base Java library. As this method may throw a NoSuchElementException and this behavior is not anticipated, this is considered a major issue.

The second bug refers to the if-statement on line 50. This statement tries to resolve the before mentioned issue, but in the current implementation it makes little sense to have the check at that location. According to SonarQube remedying these bugs will take 10 and 15 minutes respectively. A possible solution to resolve both bugs would be to refactor the method as follows

```

1 public String logIntoAccount(CustomerDTO userDto) {
2     Optional<Customer> opt_customer = customerDao.findById(userDto.getUserId());
3     // Optional<Driver> opt_driver = driverDao.findById(userDto.getUserId());
4     // Optional<Admin> opt_admin = adminDao.findById(userDto.getUserId());
5     Integer userId = opt_customer.get().getUserId();
6     Optional<CurrentUserSession> currentUserOptional = sessionDao.findById(
7         userId);
8     if (!opt_customer.isPresent()) {
9         throw new AdminExceptions("user not found");
10    }
11    if (currentUserOptional.isPresent()) {
12        throw new UserAlreadyExistWithuserId("User already logged in with this
13        number");
14    }
15    if (opt_customer.get().getPassword().equals(userDto.getPassword())) {
16        String key = RandomString.make(6);
17        CurrentUserSession currentUserSession = new CurrentUserSession(
18            opt_customer.get().getUserId(), key,
19            LocalDateTime.now());
20        sessionDao.save(currentUserSession);
21
22        return currentUserSession.toString();
23    } else {
24        throw new InvalidPasswordException("Please Enter Valid Password");
25    }
}

```

```

1 public String logIntoAccount(CustomerDTO userDto) {
2     Optional<Customer> opt_customer = customerDao.findById(userDto.getUserId());
3     // Optional<Driver> opt_driver = driverDao.findById(userDto.getUserId());
4     // Optional<Admin> opt_admin = adminDao.findById(userDto.getUserId());
5     if (opt_customer.isPresent()) {
6         Integer userId = opt_customer.get().getUserId();
7         Optional<CurrentUserSession> currentUserOptional = sessionDao.findById(
8             userId);
9         if (currentUserOptional.isPresent()) {
10            throw new UserAlreadyExistWithuserId("User already logged in with this
11            number");
12        }
13        if (opt_customer.get().getPassword().equals(userDto.getPassword())) {
14            String key = RandomString.make(6);
15            CurrentUserSession currentUserSession = new CurrentUserSession(
16                opt_customer.get().getUserId(), key,
17                LocalDateTime.now());
18            sessionDao.save(currentUserSession);
19
20            return currentUserSession.toString();
21        } else {
22            throw new InvalidPasswordException("Please Enter Valid Password");
23        }
24    } else {
25        throw new AdminExceptions("user not found");
26    }
}

```

## 5.1. SHORT TERM FINDINGS

The NST version of the same project currently does not contain any bugs.

### 5.1.1.2. Reliability rating and remediation

Due to the bugs mentioned in the previous point, the original project has received a reliability rating of C. The total time to resolve all bugs and go from C to A is estimated to be 35 minutes, which is still an acceptable amount of time given the nature of the bugs.

The NST version of the project already has a rating of A as there are no bugs found.

**Conclusion:** when it comes down to bugs, using NST limits the possibility for introducing them due to the use of expanders during code generation/rejuvenation.

### 5.1.1.3. Cognitive complexity

The cognitive complexity of an application refers to how difficult it is to understand it. According to the white paper which can be retrieved at <https://www.sonarsource.com/resources/cognitive-complexity/>, the cognitive complexity is increased when encountering loop structures, conditionals, catches, nested structures, switches and recursion. For example, the use of a switch statement with multiple cases will result in a lower cognitive complexity than when using a if-else if structure.

The original system has a cognitive complexity of 26, all of which is located in the service layer of the application, while the NST version has a cognitive complexity of 2104 with following distribution

Control	468	Data	856
Logic	152	Proxy	416
Shared	212	View	0

This high complexity is due to the way in which code is currently generated by the expanders. For example, in the create method from the CabCruds class try-catch structures are used in combination with nested conditionals. This use of if-else structures can be justified by the need for anchor points between which developers can add custom code. Listing 5.1.1.3 shows the method mentioned above.

```
1 // anchor:create:start
2 public CrudsResult<DataRef> create(ParameterContext<CabDetails> detailsParameter)
3 {
4     if (sessionContext.getRollbackOnly()) {
5         return getDiagnosticHelper().createCrudsError(CREATE_ERROR_MSG_KEY);
6     }
7     Context context = detailsParameter.getContext();
8     CabDetails details = detailsParameter.getValue();
9     // @anchor:preCreate:start
10    // @anchor:preCreate:end
11    // anchor:custom-preCreate:start
12    // anchor:custom-preCreate:end
13    // @anchor:preCreate-validation:start
14    // @anchor:preCreate-validation:end
15    CabData cabData = new CabData();
16    try {
17        // @anchor:create-beforeProjection:start
18        // @anchor:create-beforeProjection:end
19        // anchor:custom-create-beforeProjection:start
        // anchor:custom-create-beforeProjection:end
```

```

20     detailsProjector.toData(cabData, detailsParameter);
21     // @anchor:create-afterProjection:start
22     // @anchor:create-afterProjection:end
23     // anchor:custom-create-afterProjection:start
24     // anchor:custom-create-afterProjection:end
25 } catch (Exception e) {
26     sessionContext.setRollbackOnly();
27     if (logger.isErrorEnabled()) {
28         logger.error(
29             "Cannot fill data object", e
30         );
31     }
32     return getDiagnosticHelper().createCrudsError(CREATE_ERROR_MSG_KEY);
33 }
34 // @anchor:implicitNameFieldOnly-beforePersist:start
35 String identity = details.getName();
36 if (identity != null) {
37     identity = identity.trim();
38     identity = identity.length() > 0 ? identity : null;
39 }
39 // @anchor:implicitNameFieldOnly-beforePersist:end
40 try {
41     entityManager.persist(cabData);
42     // @anchor:implicitNameFieldOnly-afterPersist:start
43     if (identity == null) {
44         identity = "C-" + cabData.getId().toString();
45     }
46     cabData.setName(identity);
47     entityManager.flush();
48     // @anchor:implicitNameFieldOnly-afterPersist:end
49     if (cabData != null) {
50         if (logger.isDebugEnabled()) {
51             logger.debug(
52                 "Created CabCruds with { id : " + cabData.getId() + ", name: " +
53 cabData.getName() + " }"
54             );
55         }
56     }
57 } catch (Exception e) {
58     sessionContext.setRollbackOnly();
59     if (logger.isErrorEnabled()) {
60         logger.error(
61             "Cannot perform entry creation", e
62         );
63     }
64     return getDiagnosticHelper().createCrudsError(CREATE_ERROR_MSG_KEY);
65 }
66 CrudsResult<DataRef> result = getDataRefFromData(detailsParameter.construct(
67 cabData));
68 // @anchor:postCreate:start
69 // @anchor:postCreate:end
70 // anchor:custom-postCreate:start
71 // anchor:custom-postCreate:end
72     return result;
73 }
73 // anchor:create:end

```

Listing 5.1: Create method with high cognitive complexity

### 5.1.1.4. Code smells

Our scan of the original project revealed the presence of 120 code smells in the code base. Some of the revealed tags are

Code smell	Original system	NST system
Unused: these code smells refer to deletion of commented code blocks and unused imports.	38	2k
Clumsy: clumsy code smells indicate situations that could have been implemented in more efficient ways, such as specifying types in constructors via diamond constructors, immediately returning expressions instead of storing them in temporary variables and using appropriate methods to check for emptiness.	38	146
Convention: reports violations of coding conventions such as formatting and naming.	29	283
Bad practice: these issues will work as intended, but the way in which they are implemented are acknowledged to be bad practice.	3	101
CERT: reports violations of CERT standard rules as can be found at <a href="https://wiki.sei.cmu.edu/">https://wiki.sei.cmu.edu/</a> .	5	226
CWE: relates to rules in the Common Weakness Enumeration	16	258
Design: reveals questionable design choices, such as having duplicates of literals instead of using a variable for it.	4	80

Fixing all code smells in the original system would take an estimated 6h 26 minutes while the NST system would require 35 days of work.

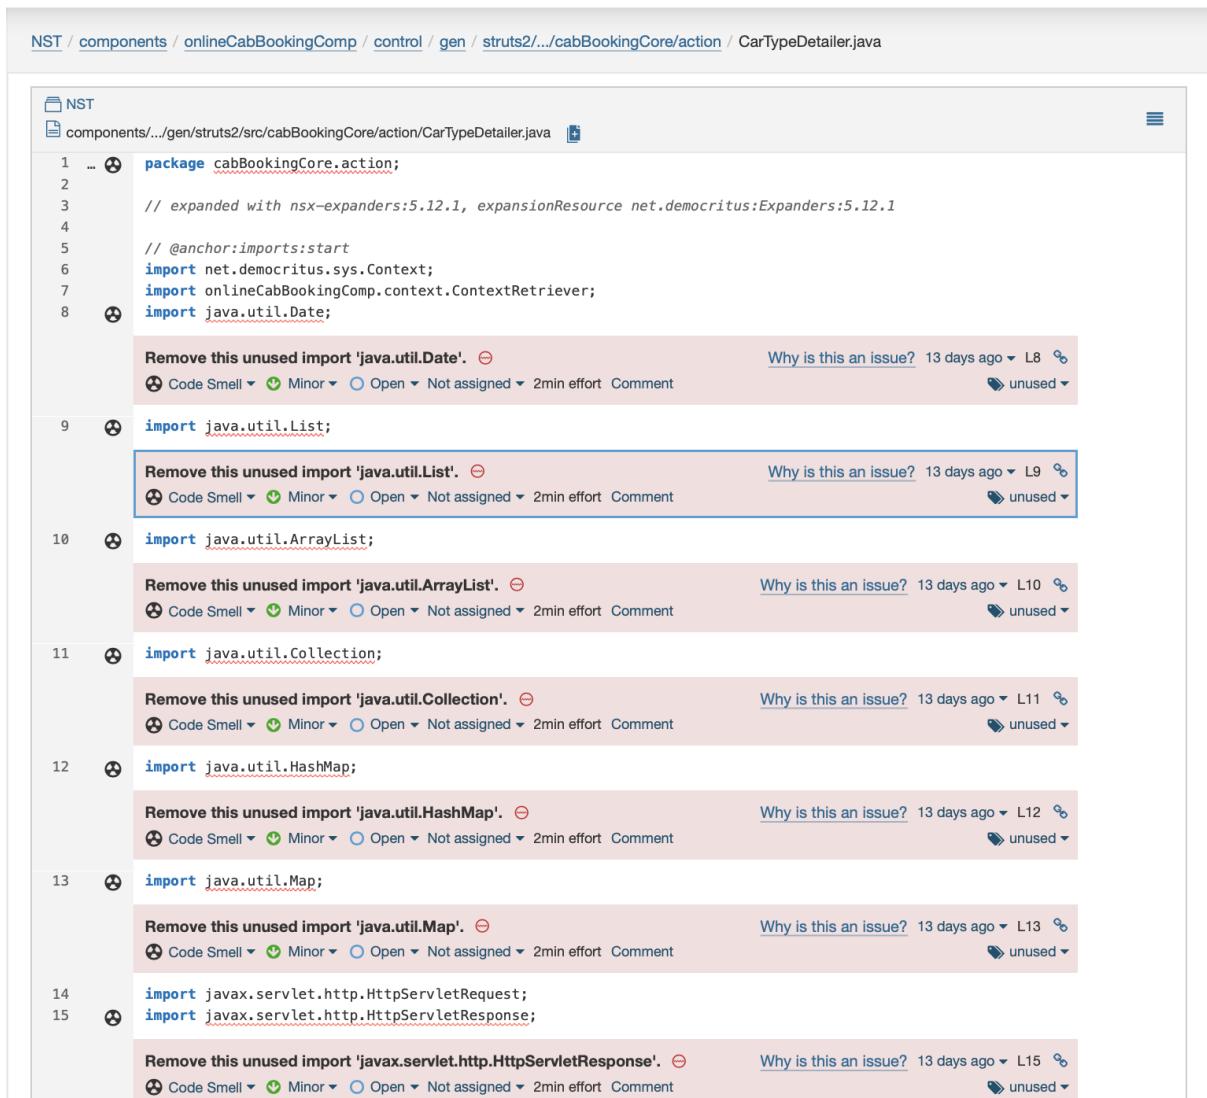
It is actually quite surprising that the system following the NST principles has such a high number of code smells, so let's take a closer look at some of the categories.

- Convention: almost all issues with this tag refer to the renaming of the package from cabBooking-Core to cabbookingcore. This can easily be done in the  $\mu$ Radiant editor and can be fixed in less than 5 minutes. The estimated time to fix all these issues is estimated to be 5 days and 5 hours.
- Unused: SonarQube reports over 2k code smells of unused code. The main issue here is that, while these are indeed correctly identified as currently unused, not all of them can be removed due to the way in which the code is generated by the expanders. For example, imports such as shown in Figure5.1 and Figure5.2 are included on a “these are most commonly used in this context” basis. Therefor it's not possible to resolve these smells as they are created by design. This means that of the estimated 35 days needed to resolve the code smells, 12 will always be there.
- CWE: the majority of these issues fall in one of three categories. A first category is one where the different agents created to provide access to the different objects should be declared as either transient or serializable. A next category would be one removal of deprecated methods. A final category is avoiding the use of generic exceptions such as RuntimeException. Code smells that fall in either of the first two categories can be resolved by updating the expanders. Code smells that fall in the third category are harder to resolve as they would require the

automatic creation of per-class exceptions which can then be further customized by the developer. This would lead to an exponential increase in number of classes which is both impractical and undesirable.

In total it would take 9 days and 7 hours to resolve all smells with the CWE tag. If we take into account that about a third of the workload cannot be resolved by design, there will always be an outstanding cost of at least 3 days.

- Bad practice: all code smells with this tag can be resolved by either adding a @deprecated tag or @Override annotation in the expanders. SonarQube gives an estimated cost of 1 day.
- Design: all except one of the design code smells refer to the use of duplicate string literals. The estimated cost for fixing these smells is 1 day 6 hours.



The screenshot shows a SonarQube interface for the file `CarTypeDetailer.java`. The code contains several unused import statements, each highlighted with a red underline. Below each underlined import is a tooltip providing details about the smell:

- Line 1: `Remove this unused import 'java.util.Date'.` (Code Smell, Minor, Open, Not assigned, 2min effort, Comment)
- Line 9: `Remove this unused import 'java.util.List'.` (Code Smell, Minor, Open, Not assigned, 2min effort, Comment)
- Line 10: `Remove this unused import 'java.util.ArrayList'.` (Code Smell, Minor, Open, Not assigned, 2min effort, Comment)
- Line 11: `Remove this unused import 'java.util.Collection'.` (Code Smell, Minor, Open, Not assigned, 2min effort, Comment)
- Line 12: `Remove this unused import 'java.util.HashMap'.` (Code Smell, Minor, Open, Not assigned, 2min effort, Comment)
- Line 13: `Remove this unused import 'java.util.Map'.` (Code Smell, Minor, Open, Not assigned, 2min effort, Comment)
- Line 15: `Remove this unused import 'javax.servlet.http.HttpServletRequest'.` (Code Smell, Minor, Open, Not assigned, 2min effort, Comment)

Figure 5.1.: Example of the unused imports generated by  $\mu$ Radiant

### 5.1.1.5. Technical debt and Maintainability Rating

As stated in the previous paragraph, it would take 35 days to fix all code smells from our NST system. We can solve almost 6 days of work in less than 5 minutes using  $\mu$ Radiant and an additional 9 days of

## 5.1. SHORT TERM FINDINGS

<input type="checkbox"/> components/.../src/cabBookingCore/action/AddressDeleter.java		
<input type="checkbox"/> Remove this unused import 'java.util.Date'.	13 days ago	L8 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'java.util.List'.	13 days ago	L9 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'java.util.ArrayList'.	13 days ago	L10 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'java.util.Collection'.	13 days ago	L11 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'java.util.HashMap'.	13 days ago	L12 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'java.util.Map'.	13 days ago	L13 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'javax.servlet.http.HttpServletRequest'.	13 days ago	L14 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'javax.servlet.http.HttpServletResponse'.	13 days ago	L15 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'javax.servlet.http.HttpSession'.	13 days ago	L16 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'org.apache.struts2.ServletActionContext'.	13 days ago	L17 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'net.democritus.sys.DataRef'.	13 days ago	L21 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'net.democritus.sys.PageRef'.	13 days ago	L22 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'net.democritus.sys.SearchDataRef'.	13 days ago	L24 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'net.democritus.sys.SearchResult'.	13 days ago	L25 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		
<input type="checkbox"/> Remove this unused import 'com.opensymphony.xwork2.ActionContext'.	13 days ago	L28 %  unused
<input checked="" type="checkbox"/> Code Smell  Minor  Open  Not assigned  2min effort		

Figure 5.2.: Example of the unused imports generated by  $\mu$ Radiant

work can be taken care of by tweaking the expanders. Of the 35 days needed to fix all code smells at least 15 days can not be resolved due to the design of NST.

In the previous discussion we have not mentioned CERT, CWE and clumsy code smells. The reason for this is that both CERT and CWE smells are also tagged with other tags which have already been mentioned and their remediation cost has thus already been counted. The clumsy smells require a case-by-case inspection: some could be fixed quite easily in the expanders, such as merging nested conditional statements, while others cannot be resolved due to the presence of anchors as is the case in for instance the getCustomDisplayName method from the CarTypeCruds class. As these smells only account for a workload of 7 hours to remedy and overall they are classified as minor issues, it might not be feasible to spend time at tweaking the expanders just for the sake of reducing code smells. The remaining 5 days of work are distributed among various other tags which we will not discuss here for

sake of interest.

Our initial system on the other hand can be released of its code smells in less than 7 hours.

We already know that the technical debt of software will increase over time if a developer does not consciously tries and avoid it. The main benefit of NST is that, since the majority of the code is generated/rejuvenated by  $\mu$ Radiant, the technical debt will almost never increase as fast as otherwise.

The maintainability rating is closely related to the technical debt ratio. As both versions have a debt ratio below 5%, both versions get a maintainability rating of A.

#### 5.1.1.6. Vulnerabilities

In our initial system, SonarQube found 8 vulnerability issues, all of which entail replacing a persistent entity with a POJO or DTO object. Our NST version comes with zero vulnerability issues.

#### 5.1.1.7. Security rating and remediation

As all of the vulnerabilities found in the initial system are deemed to be critical issues, this version has received a security rating of D which is the second last lowest rating. Luckily the estimated remediation effort is 1 hour 20 minutes, so we can easily improve this rating from D to A.

Since our NST version has no detected vulnerabilities, it received a rating of A.

#### 5.1.1.8. Conclusion

While at first glance the NST version of the system looks to be way worse than our initial version - according to SonarQube we would need well over a month to remedy all issues! - things aren't as bad at all.

Indeed, if we take a closer look at the two areas in which NST seems to perform poorly - namely cognitive complexity and code smells - it becomes clear that many of the issues can be resolved by tweaking the expanders by for instance switching from try-catch statements to switch statements in order to lower the cognitive complexity or resolve 8 days of remediation cost from code smells, thereby providing cleaner code.

The main question here would be the following: is it worth recreating a system using the  $\mu$ Radiant tool? The answer is: it depends.

If an architectural blueprint is present - which really should be the case in a professional setting - it should not take long to transform an application using  $\mu$ Radiant. The advantage obtained from doing so, i.e. obtaining a system build according to the four design principles given in section 3.2.1 which can be mathematically proven [VM03] to deliver software with the desired properties, outweigh the initial high cognitive complexity and high amount of code smells.

## 5.2. In the long run

As the long term benefits of transforming a system according to NST versus traditional refactoring cannot be measured during this thesis, this part is not covered at time of writing.

## 6. Previous research projects

### 6.1. Research project 1

Title/subject: Sketch recognition

Promotor: Hans Vangheluwe

Related to this thesis: No.

Abstract: Sketch recognition is a problem that has been around for decades, and over the years several solutions have been proposed to resolve it. In this report we will first take a look at some of the earlier work that has been performed on this problem. The second part of will go into more detail about a prototype build based on one of these previous works, describing how we can implement it as well as providing some background information about the different aspects involved. To conclude we will go over the current limitations and provide possible improvements.

### 6.2. Research project 2

Title/subject: Visualizing large TCR networks

Promotor: Pieter Meysman and Sofie Giellis

Related to this thesis: No.

Abstract: In this research project, the aim was to find ways to efficiently visualise large TCR networks. To this end some of the available tools and frameworks were explored before selecting two viable options, Igraph and NetworkX. After some experimentation Igraph remained as the final option.



## A. Bibliography



# Bibliography

- [Bab21] Shibu Babuchandran. Software release life cycle (srlc): Understand the 6 main stages, 2021. [Online; publication date Dec 6 2021]. 4
- [DSN03] Serge Demeyer, Ducasse Stephane, and Oscar Nierstrasz. *Object-oriented reengineering patterns*. Morgan Kaufman Pub., 2003. 11
- [ECPB12] Thomas Erl, Benjamin Carlyle, Cesare Pautasso, and Raj Balasubramanian. *SOA with REST Principles, Patterns Constraints for Building Enterprise Solutions with REST*. Pearson, 2012. 10
- [Ext21] Extensibility. Extensibility — Wikipedia, the free encyclopedia, 2021. [Online]. 4
- [GHJ<sup>+</sup>94] Erich Gamma Gamma, Richard Helm, Ralph Johnson, John Vlissides, and Grady Booch. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley Professional, 1994. 11
- [NGI] NGINX Glossary. What is load balancing? [Online]. 8
- [Pie] Michael Pierce. Software release life cycle (srlc): Understand the 6 main stages. [Online; no publication date]. 3
- [Ric15] Mark Richards. *Software Architecture Patterns Understanding Common Architecture Patterns and When to Use Them*. O'Reilly, 2015. 8
- [sof] What is software rot, and why does it happen? [Online; no publication date or author]. 4
- [Sof23] Software rot. Software rot — Wikipedia, the free encyclopedia, 2023. [Online]. 3
- [sta21] The importance of stability and reliability testing in software development, 2021. [Online]. 3
- [tec] What is technical debt and what should you do about it? [Online; no publication date or author]. 4
- [VM03] Jan Verelst and Herwig and Mannaert. *Normalised Systems Theory*. Morgan Kaufman Pub., 2003. 54